Red Hat Enterprise Linux 8

Deploying different types of servers

A guide to deploying different types of servers in Red Hat Enterprise Linux 8

Red Hat Enterprise Linux 8 Deploying different types of servers

A guide to deploying different types of servers in Red Hat Enterprise Linux 8

Legal Notice

Abstract

This document describes how to configure and run different types of servers on Red Hat Enterprise Linux 8, including Apache HTTP web server, Samba server, NFS server, available database servers, and the CUPS server.

Table of Contents

MAKING OPEN SOURCE MORE INCLUSIVE

Red Hat is committed to replacing problematic language in our code, documentation, and web properties. We are beginning with these four terms: master, slave, blacklist, and whitelist. Because of the enormity of this endeavor, these changes will be implemented gradually over several upcoming releases. For more details, see our CTO Chris Wright's message .

PROVIDING FEEDBACK ON RED HAT DOCUMENTATION

We appreciate your feedback on our documentation. Let us know how we can improve it.

Submitting comments on specific passages

1. View the documentation in the **Multi-page HTML** format and ensure that you see the **Feedback** button in the upper right corner after the page fully loads.

2. Use your cursor to highlight the part of the text that you want to comment on.

3. Click the **Add Feedback** button that appears near the highlighted text.

4. Add your feedback and click **Submit**.

Submitting feedback through Bugzilla (account required)

1. Log in to the Bugzilla website.

2. Select the correct version from the **Version** menu.

3. Enter a descriptive title in the **Summary** field.

4. Enter your suggestion for improvement in the **Description** field. Include links to the relevant parts of the documentation.

5. Click **Submit Bug**.

CHAPTER 1. SETTING UP THE APACHE HTTP WEB SERVER

1.1. INTRODUCTION TO THE APACHE HTTP WEB SERVER

A *web server* is a network service that serves content to a client over the web. This typically means web pages, but any other documents can be served as well. Web servers are also known as HTTP servers, as they use the *hypertext transport protocol* (**HTTP**).

The **Apache HTTP Server**, **httpd**, is an open source web server developed by the Apache Software Foundation.

If you are upgrading from a previous release of Red Hat Enterprise Linux, you have to update the **httpd** service configuration accordingly. This section reviews some of the newly added features, and guides you through the update of prior configuration files.

1.2. NOTABLE CHANGES IN THE APACHE HTTP SERVER

The **Apache HTTP Server** has been updated from version 2.4.6 in RHEL 7 to version 2.4.37 in RHEL 8. This updated version includes several new features, but maintains backwards compatibility with the RHEL 7 version at the level of configuration and Application Binary Interface (ABI) of external modules.

New features include:

- **HTTP/2** support is now provided by the **mod_http2** package, which is a part of the **httpd** module.

- systemd socket activation is supported. See **httpd.socket(8)** man page for more details.

- Multiple new modules have been added:

  - **mod_proxy_hcheck** - a proxy health-check module

  - **mod_proxy_uwsgi** - a Web Server Gateway Interface (WSGI) proxy

  - **mod_proxy_fdpass** - provides support for the passing the socket of the client to another process

  - **mod_cache_socache** - an HTTP cache using, for example, memcache backend

  - **mod_md** - an ACME protocol SSL/TLS certificate service

- The following modules now load by default:

  - **mod_request**

  - **mod_macro**

  - **mod_watchdog**

- A new subpackage, **httpd-filesystem**, has been added, which contains the basic directory layout for the **Apache HTTP Server** including the correct permissions for the directories.

- Instantiated service support, **httpd@.service** has been introduced. See the **httpd.service** man page for more information.

- A new **httpd-init.service** replaces the **%post script** to create a self-signed **mod_ssl** key pair.

- Automated TLS certificate provisioning and renewal using the Automatic Certificate Management Environment (ACME) protocol is now supported with the **mod_md** package (for use with certificate providers such as **Let's Encrypt**).

- The **Apache HTTP Server** now supports loading TLS certificates and private keys from hardware security tokens directly from **PKCS#11** modules. As a result, a **mod_ssl** configuration can now use **PKCS#11** URLs to identify the TLS private key, and, optionally, the TLS certificate in the **SSLCertificateKeyFile** and **SSLCertificateFile** directives.

- A new **ListenFree** directive in the **/etc/httpd/conf/httpd.conf** file is now supported. Similarly to the **Listen** directive, **ListenFree** provides information about IP addresses, ports, or IP address-and-port combinations that the server listens to. However, with **ListenFree**, the **IP_FREEBIND** socket option is enabled by default. Hence, **httpd** is allowed to bind to a nonlocal IP address or to an IP address that does not exist yet. This allows **httpd** to listen on a socket without requiring the underlying network interface or the specified dynamic IP address to be up at the time when **httpd** is trying to bind to it.

  Note that the **ListenFree** directive is currently available only in RHEL 8.

  For more details on **ListenFree**, see the following table:

  Table 1.1. ListenFree directive's syntax, status, and modules

Syntax	Status	Modules
ListenFree [IP-address:]portnumber [protocol]	MPM	event, worker, prefork, mpm_winnt, mpm_netware, mpmt_os2

Other notable changes include:

- The following modules have been removed:

  - **mod_file_cache**

  - **mod_nss**
    Use **mod_ssl** as a replacement. For details about migrating from **mod_nss**, see Section 1.14, "Exporting a private key and certificates from an NSS database to use them in an Apache web server configuration".

  - **mod_perl**

- The default type of the DBM authentication database used by the **Apache HTTP Server** in RHEL 8 has been changed from **SDBM** to **db5**.

- The **mod_wsgi** module for the **Apache HTTP Server** has been updated to Python 3. WSGI applications are now supported only with Python 3, and must be migrated from Python 2.

- The multi-processing module (MPM) configured by default with the **Apache HTTP Server** has changed from a multi-process, forked model (known as **prefork**) to a high-performance multi-threaded model, **event**.
  Any third-party modules that are not thread-safe need to be replaced or removed. To change the configured MPM, edit the **/etc/httpd/conf.modules.d/00-mpm.conf** file. See the **httpd.service(8)** man page for more information.

- The minimum UID and GID allowed for users by suEXEC are now 1000 and 500, respectively (previously 100 and 100).

- The **/etc/sysconfig/httpd** file is no longer a supported interface for setting environment variables for the **httpd** service. The **httpd.service(8)** man page has been added for the systemd service.

- Stopping the **httpd** service now uses a "graceful stop" by default.

- The **mod_auth_kerb** module has been replaced by the **mod_auth_gssapi** module.

1.3. UPDATING THE CONFIGURATION

To update the configuration files from the **Apache HTTP Server** version used in Red Hat Enterprise Linux 7, choose one of the following options:

- If **/etc/sysconfig/httpd** is used to set environment variables, create a systemd drop-in file instead.

- If any third-party modules are used, ensure they are compatible with a threaded MPM.

- If suexec is used, ensure user and group IDs meet the new minimums.

You can check the configuration for possible errors by using the following command:

```
# apachectl configtest
Syntax OK
```

1.4. THE APACHE CONFIGURATION FILES

The **httpd**, by default, reads the configuration files after start. You can see the list of the locations of configuration files in the table below.

Table 1.2. The httpd service configuration files

Path	Description
/etc/httpd/conf/httpd.conf	The main configuration file.
/etc/httpd/conf.d/	An auxiliary directory for configuration files that are included in the main configuration file.
/etc/httpd/conf.modules.d/	An auxiliary directory for configuration files which load installed dynamic modules packaged in Red Hat Enterprise Linux. In the default configuration, these configuration files are processed first.

Although the default configuration is suitable for most situations, you can use also other configuration options. For any changes to take effect, restart the web server first.

To check the configuration for possible errors, type the following at a shell prompt:

```
# apachectl configtest
Syntax OK
```

To make the recovery from mistakes easier, make a copy of the original file before editing it.

1.5. MANAGING THE HTTPD SERVICE

This section describes how to start, stop, and restart the **httpd** service.

Prerequisites

- The Apache HTTP Server is installed.

Procedure

- To start the **httpd** service, enter:

  ```
  # systemctl start httpd
  ```

- To stop the **httpd** service, enter:

  ```
  # systemctl stop httpd
  ```

- To restart the **httpd** service, enter:

  ```
  # systemctl restart httpd
  ```

1.6. SETTING UP A SINGLE-INSTANCE APACHE HTTP SERVER

This section describes how to set up a single-instance Apache HTTP Server to serve static HTML content.

Follow the procedure in this section if the web server should provide the same content for all domains associated with the server. If you want to provide different content for different domains, set up name-based virtual hosts. For details, see Configuring Apache name-based virtual hosts .

Procedure

1. Install the **httpd** package:

   ```
   # yum install httpd
   ```

2. Open the TCP port **80** in the local firewall:

   ```
   # firewall-cmd --permanent --add-port=80/tcp
   # firewall-cmd --reload
   ```

3. Enable and start the **httpd** service:

   ```
   # systemctl enable --now httpd
   ```

4. Optional: Add HTML files to the **/var/www/html/** directory.

> **NOTE**
>
> When adding content to **/var/www/html/**, files and directories must be readable by the user under which **httpd** runs by default. The content owner can be the either the **root** user and **root** user group, or another user or group of the administrator's choice. If the content owner is the **root** user and **root** user group, the files must be readable by other users. The SELinux context for all the files and directories must be **httpd_sys_content_t**, which is applied by default to all content within the **/var/www** directory.

Verification steps

- Connect with a web browser to **http://** *server_IP_or_host_name* /.
  If the **/var/www/html/** directory is empty or does not contain an **index.html** or **index.htm** file, Apache displays the **Red Hat Enterprise Linux Test Page**. If **/var/www/html/** contains HTML files with a different name, you can load them by entering the URL to that file, such as **http://** *server_IP_or_host_name* / *example.html*.

Additional resources

- Apache manual: Installing the Apache HTTP server manual .

- See the **httpd.service(8)** man page.

1.7. CONFIGURING APACHE NAME-BASED VIRTUAL HOSTS

Name-based virtual hosts enable Apache to serve different content for different domains that resolve to the IP address of the server.

The procedure in this section describes setting up a virtual host for both the **example.com** and **example.net** domain with separate document root directories. Both virtual hosts serve static HTML content.

Prerequisites

- Clients and the web server resolve the **example.com** and **example.net** domain to the IP address of the web server.
  Note that you must manually add these entries to your DNS server.

Procedure

1. Install the **httpd** package:

   ```
   # yum install httpd
   ```

2. Edit the **/etc/httpd/conf/httpd.conf** file:

   a. Append the following virtual host configuration for the **example.com** domain:

   ```
   <VirtualHost *:80>
       DocumentRoot "/var/www/example.com/"
       ServerName example.com
   ```

```
        CustomLog /var/log/httpd/example.com_access.log combined
        ErrorLog /var/log/httpd/example.com_error.log
</VirtualHost>
```

These settings configure the following:

- All settings in the **<VirtualHost *:80>** directive are specific for this virtual host.

- **DocumentRoot** sets the path to the web content of the virtual host.

- **ServerName** sets the domains for which this virtual host serves content.
  To set multiple domains, add the **ServerAlias** parameter to the configuration and specify the additional domains separated with a space in this parameter.

- **CustomLog** sets the path to the access log of the virtual host.

- **ErrorLog** sets the path to the error log of the virtual host.

> **NOTE**
>
> Apache uses the first virtual host found in the configuration also for requests that do not match any domain set in the **ServerName** and **ServerAlias** parameters. This also includes requests sent to the IP address of the server.

3. Append a similar virtual host configuration for the **example.net** domain:

```
<VirtualHost *:80>
    DocumentRoot "/var/www/example.net/"
    ServerName example.net
    CustomLog /var/log/httpd/example.net_access.log combined
    ErrorLog /var/log/httpd/example.net_error.log
</VirtualHost>
```

4. Create the document roots for both virtual hosts:

```
# mkdir /var/www/example.com/
# mkdir /var/www/example.net/
```

5. If you set paths in the **DocumentRoot** parameters that are not within **/var/www/**, set the **httpd_sys_content_t** context on both document roots:

```
# semanage fcontext -a -t httpd_sys_content_t "/srv/example.com(/.*)?"
# restorecon -Rv /srv/example.com/
# semanage fcontext -a -t httpd_sys_content_t "/srv/example.net(/.\*)?"
# restorecon -Rv /srv/example.net/
```

These commands set the **httpd_sys_content_t** context on the /**srv**/**example.com**/ and /**srv**/**example.net**/ directory.

Note that you must install the **policycoreutils-python-utils** package to run the **restorecon** command.

6. Open port **80** in the local firewall:

```
# firewall-cmd --permanent --add-port=80/tcp
# firewall-cmd --reload
```

7. Enable and start the **httpd** service:

```
# systemctl enable --now httpd
```

Verification steps

1. Create a different example file in each virtual host's document root:

```
# echo "vHost example.com" > /var/www/example.com/index.html
# echo "vHost example.net" > /var/www/example.net/index.html
```

2. Use a browser and connect to **http://example.com**. The web server shows the example file from the **example.com** virtual host.

3. Use a browser and connect to **http://example.net**. The web server shows the example file from the **example.net** virtual host.

Additional resources

- Installing the Apache HTTP Server manual – Virtual Hosts

1.8. CONFIGURING KERBEROS AUTHENTICATION FOR THE APACHE HTTP WEB SERVER

To perform Kerberos authentication in the Apache HTTP web server, RHEL 8 uses the **mod_auth_gssapi** Apache module. The Generic Security Services API (**GSSAPI**) is an interface for applications that make requests to use security libraries, such as Kerberos. The **gssproxy** service allows to implement privilege separation for the **httpd** server, which optimizes this process from the security point of view.
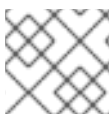
> **NOTE**
>
> The **mod_auth_gssapi** module replaces the removed **mod_auth_kerb** module.

Prerequisites

- The **httpd** and **gssproxy** packages are installed.

- The Apache web server is set up and the **httpd** service is running.

1.8.1. Setting up GSS-Proxy in an IdM environment

This procedure describes how to set up **GSS-Proxy** to perform Kerberos authentication in the Apache HTTP web server.

Procedure

1. Enable access to the **keytab** file of HTTP/<SERVER_NAME>@realm principal by creating the service principal:

```
# ipa service-add HTTP/<SERVER_NAME>
```

2. Retrieve the **keytab** for the principal stored in the **/etc/gssproxy/http.keytab** file:

```
# ipa-getkeytab -s $(awk '/^server =/ {print $3}' /etc/ipa/default.conf) -k
/etc/gssproxy/http.keytab -p HTTP/$(hostname -f)
```

This step sets permissions to 400, thus only the **root** user has access to the **keytab** file. The **apache** user does not.

3. Create the **/etc/gssproxy/80-httpd.conf** file with the following content:

```
[service/HTTP]
  mechs = krb5
  cred_store = keytab:/etc/gssproxy/http.keytab
  cred_store = ccache:/var/lib/gssproxy/clients/krb5cc_%U
  euid = apache
```

4. Restart and enable the **gssproxy** service:

```
# systemctl restart gssproxy.service
# systemctl enable gssproxy.service
```

Additional resources

- **gssproxy(8)** man pages

- **gssproxy-mech(8)** man pages

- **gssproxy.conf(5)** man pages

1.8.2. Configuring Kerberos authentication for a directory shared by the Apache HTTP web server

This procedure describes how to configure Kerberos authentication for the **/var/www/html/private/** directory.

Prerequisites

- The **gssproxy** service is configured and running.

Procedure

1. Configure the **mod_auth_gssapi** module to protect the **/var/www/html/private/** directory:

```
<Location /var/www/html/private>
  AuthType GSSAPI
  AuthName "GSSAPI Login"
  Require valid-user
</Location>
```

2. Create the **/etc/systemd/system/httpd.service** file with the following content:

```
.include /lib/systemd/system/httpd.service
[Service]
Environment=GSS_USE_PROXY=1
```

3. Reload the **systemd** configuration:

   ```
   # systemctl daemon-reload
   ```

4. Restart the **httpd** service:

   ```
   # systemctl restart httpd.service
   ```

Verification steps

1. Obtain a Kerberos ticket:

   ```
   # kinit
   ```

2. Open the URL to the protected directory in a browser.

1.9. CONFIGURING TLS ENCRYPTION ON AN APACHE HTTP SERVER

By default, Apache provides content to clients using an unencrypted HTTP connection. This section describes how to enable TLS encryption and configure frequently used encryption-related settings on an Apache HTTP Server.

Prerequisites

- The Apache HTTP Server is installed and running.

1.9.1. Adding TLS encryption to an Apache HTTP Server

This section describes how to enable TLS encryption on an Apache HTTP Server for the **example.com** domain.

Prerequisites

- The Apache HTTP Server is installed and running.

- The private key is stored in the **/etc/pki/tls/private/example.com.key** file.
  For details about creating a private key and certificate signing request (CSR), as well as how to request a certificate from a certificate authority (CA), see your CA's documentation. Alternatively, if your CA supports the ACME protocol, you can use the **mod_md** module to automate retrieving and provisioning TLS certificates.

- The TLS certificate is stored in the **/etc/pki/tls/certs/example.com.crt** file. If you use a different path, adapt the corresponding steps of the procedure.

- The CA certificate is stored in the **/etc/pki/tls/certs/ca.crt** file. If you use a different path, adapt the corresponding steps of the procedure.

- Clients and the web server resolve the host name of the server to the IP address of the web server.

Procedure

1. Install the **mod_ssl** package:

   > # **yum install mod_ssl**

2. Edit the **/etc/httpd/conf.d/ssl.conf** file and add the following settings to the **<VirtualHost _default_:443>** directive:

   a. Set the server name:

      > **ServerName** *example.com*

      > IMPORTANT
      >
      > The server name must match the entry set in the **Common Name** field of the certificate.

   b. Optional: If the certificate contains additional host names in the **Subject Alt Names** (SAN) field, you can configure **mod_ssl** to provide TLS encryption also for these host names. To configure this, add the **ServerAliases** parameter with corresponding names:

      > **ServerAlias** *www.example.com server.example.com*

   c. Set the paths to the private key, the server certificate, and the CA certificate:

      > **SSLCertificateKeyFile "/etc/pki/tls/private/example.com.key"**
      > **SSLCertificateFile "/etc/pki/tls/certs/example.com.crt"**
      > **SSLCACertificateFile "/etc/pki/tls/certs/ca.crt"**

3. For security reasons, configure that only the **root** user can access the private key file:

   > # **chown root:root /etc/pki/tls/private/example.com.key**
   > # **chmod 600 /etc/pki/tls/private/example.com.key**

   > WARNING
   >
   > If the private key was accessed by unauthorized users, revoke the certificate, create a new private key, and request a new certificate. Otherwise, the TLS connection is no longer secure.

4. Open port **443** in the local firewall:

   > # **firewall-cmd --permanent --add-port=443/tcp**
   > # **firewall-cmd --reload**

5. Restart the **httpd** service:

> # **systemctl restart httpd**
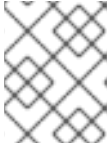
> **NOTE**
>
> If you protected the private key file with a password, you must enter this password each time when the **httpd** service starts.

Verification steps

- Use a browser and connect to **https://** *example.com*.

Additional resources

- SSL/TLS Encryption

- Security considerations for TLS in RHEL 8

1.9.2. Setting the supported TLS protocol versions on an Apache HTTP Server

By default, the Apache HTTP Server on RHEL uses the system-wide crypto policy that defines safe default values, which are also compatible with recent browsers. For example, the **DEFAULT** policy defines that only the **TLSv1.2** and **TLSv1.3** protocol versions are enabled in apache.

This section describes how to manually configure which TLS protocol versions your Apache HTTP Server supports. Follow the procedure if your environment requires to enable only specific TLS protocol versions, for example:

- If your environment requires that clients can also use the weak **TLS1** (TLSv1.0) or **TLS1.1** protocol.

- If you want to configure that Apache only supports the **TLSv1.2** or **TLSv1.3** protocol.

Prerequisites

- TLS encryption is enabled on the server as described in Adding TLS encryption to an Apache HTTP server.

Procedure

1. Edit the **/etc/httpd/conf/httpd.conf** file, and add the following setting to the **<VirtualHost>** directive for which you want to set the TLS protocol version. For example, to enable only the **TLSv1.3** protocol:

   > **SSLProtocol -All TLSv1.3**

2. Restart the **httpd** service:

   > # **systemctl restart httpd**

Verification steps

1. Use the following command to verify that the server supports **TLSv1.3**:

> # **openssl s_client -connect *example.com*:443 -tls1_3**

2. Use the following command to verify that the server does not support **TLSv1.2**:

> # **openssl s_client -connect *example.com*:443 -tls1_2**

If the server does not support the protocol, the command returns an error:

> 140111600609088:error:1409442E:SSL routines:ssl3_read_bytes:tlsv1 alert protocol
> version:ssl/record/rec_layer_s3.c:1543:SSL alert number 70

3. Optional: Repeat the command for other TLS protocol versions.

Additional resources

- **update-crypto-policies(8)** man page

- Using system-wide cryptographic policies .

- For further details about the **SSLProtocol** parameter, refer to the **mod_ssl** documentation in the Apache manual: Installing the Apache HTTP server manual .

1.9.3. Setting the supported ciphers on an Apache HTTP Server

By default, the Apache HTTP Server uses the system-wide crypto policy that defines safe default values, which are also compatible with recent browsers. For the list of ciphers the system-wide crypto allows, see the **/etc/crypto-policies/back-ends/openssl.config** file.

This section describes how to manually configure which ciphers your Apache HTTP Server supports. Follow the procedure if your environment requires specific ciphers.

Prerequisites

- TLS encryption is enabled on the server as described in Adding TLS encryption to an Apache HTTP server.

Procedure

1. Edit the **/etc/httpd/conf/httpd.conf** file, and add the **SSLCipherSuite** parameter to the **<VirtualHost>** directive for which you want to set the TLS ciphers:

> **SSLCipherSuite**
> **"EECDH+AESGCM:EDH+AESGCM:AES256+EECDH:AES256+EDH:!SHA1:!SHA256"**

This example enables only the **EECDH+AESGCM**, **EDH+AESGCM**, **AES256+EECDH**, and **AES256+EDH** ciphers and disables all ciphers which use the **SHA1** and **SHA256** message authentication code (MAC).

2. Restart the **httpd** service:

> # **systemctl restart httpd**

Verification steps

1. To display the list of ciphers the Apache HTTP Server supports:

   a. Install the **nmap** package:

      ```
      # yum install nmap
      ```

   b. Use the **nmap** utility to display the supported ciphers:

      ```
      # nmap --script ssl-enum-ciphers -p 443 example.com
      ...
      PORT    STATE SERVICE
      443/tcp open  https
      | ssl-enum-ciphers:
      |   TLSv1.2:
      |     ciphers:
      |       TLS_ECDHE_RSA_WITH_AES_256_GCM_SHA384 (ecdh_x25519) - A
      |       TLS_DHE_RSA_WITH_AES_256_GCM_SHA384 (dh 2048) - A
      |       TLS_ECDHE_RSA_WITH_CHACHA20_POLY1305_SHA256 (ecdh_x25519) - A
      ...
      ```

Additional resources

- **update-crypto-policies(8)** man page

- Using system-wide cryptographic policies .

- SSLCipherSuite

1.10. CONFIGURING TLS CLIENT CERTIFICATE AUTHENTICATION

Client certificate authentication enables administrators to allow only users who authenticate using a certificate to access resources on the web server. This section describes how to configure client certificate authentication for the **/var/www/html/Example/** directory.

If the Apache HTTP Server uses the TLS 1.3 protocol, certain clients require additional configuration. For example, in Firefox, set the **security.tls.enable_post_handshake_auth** parameter in the **about:config** menu to **true**. For further details, see Transport Layer Security version 1.3 in Red Hat Enterprise Linux 8.

Prerequisites

- TLS encryption is enabled on the server as described in Adding TLS encryption to an Apache HTTP server.

Procedure

1. Edit the **/etc/httpd/conf/httpd.conf** file and add the following settings to the **<VirtualHost>** directive for which you want to configure client authentication:

   ```
   <Directory "/var/www/html/Example/">
    SSLVerifyClient require
   </Directory>
   ```

The **SSLVerifyClient require** setting defines that the server must successfully validate the client certificate before the client can access the content in the **/var/www/html/Example/** directory.

2. Restart the **httpd** service:

   ```
   # systemctl restart httpd
   ```

Verification steps

1. Use the **curl** utility to access the **https://example.com/Example/** URL without client authentication:

   ```
   $ curl https://example.com/Example/
   curl: (56) OpenSSL SSL_read: error:1409445C:SSL routines:ssl3_read_bytes:tlsv13 **alert
   certificate required**, errno 0
   ```

   The error indicates that the web server requires a client certificate authentication.

2. Pass the client private key and certificate, as well as the CA certificate to **curl** to access the same URL with client authentication:

   ```
   $ curl --cacert ca.crt --key client.key --cert client.crt https://example.com/Example/
   ```

   If the request succeeds, **curl** displays the **index.html** file stored in the **/var/www/html/Example/** directory.

Additional resources

- mod_ssl configuration

1.11. SECURING WEB APPLICATIONS ON A WEB SERVER USING MODSECURITY

ModSecurity is an open source web application firewall (WAF) supported by various web servers such as Apache, Nginx, and IIS, which reduces security risks in web applications. ModSecurity provides customizable rule sets for configuring your server.

The **mod_security-crs** package contains the core rule set (CRS) with rules against cross-website scripting, bad user agents, SQL injection, Trojans, session hijacking, and other exploits.

1.11.1. Deploying the ModSecurity web-based application firewall for Apache

To reduce risks related to running web-based applications on your web server by deploying ModSecurity, install the **mod_security** and **mod_security_crs** packages for the Apache HTTP server. The **mod_security_crs** package provides the core rule set (CRS) for the ModSecurity web-based application firewall (WAF) module.

Procedure

1. Install the **mod_security**, **mod_security_crs**, and **httpd** packages:

   ```
   # yum install -y mod_security mod_security_crs httpd
   ```

2. Start the **httpd** server:

```
# systemctl restart httpd
```

Verification

1. Verify that the ModSecurity web-based application firewall is enabled on your Apache HTTP server:

```
# httpd -M | grep security
 security2_module (shared)
```

2. Check that the **/etc/httpd/modsecurity.d/activated_rules/** directory contains rules provided by **mod_security_crs**:

```
# ls /etc/httpd/modsecurity.d/activated_rules/
...
REQUEST-921-PROTOCOL-ATTACK.conf
REQUEST-930-APPLICATION-ATTACK-LFI.conf
...
```

Additional resources

- Red Hat JBoss Core Services ModSecurity Guide

- An introduction to web application firewalls for Linux sysadmins

1.11.2. Adding a custom rule to ModSecurity

If the rules contained in the ModSecurity core rule set (CRS) do not fit your scenario and if you want to prevent additional possible attacks, you can add your custom rules to the rule set used by the ModSecurity web-based application firewall. The following example demonstrates the addition of a simple rule. For creating more complex rules, see the reference manual on the ModSecurity Wiki website.

Prerequisites

- ModSecurity for Apache is installed and enabled.

Procedure

1. Open the **/etc/httpd/conf.d/mod_security.conf** file in a text editor of your choice, for example:

```
# vi /etc/httpd/conf.d/mod_security.conf
```

2. Add the following example rule after the line starting with **SecRuleEngine On**:

```
SecRule ARGS:data "@contains evil" "deny,status:403,msg:'param data contains evil
data',id:1"
```

The previous rule forbids the use of resources to the user if the **data** parameter contains the **evil** string.

3. Save the changes, and quit the editor.

4. Restart the **httpd** server:

```
# systemctl restart httpd
```

Verification

1. Create a *test*.**html** page:

```
# echo "mod_security test" > /var/www/html/test.html
```

2. Restart the **httpd** server:

```
# systemctl restart httpd
```

3. Request *test.html* without malicious data in the **GET** variable of the HTTP request:

```
$ curl http://localhost/test.html?data=good

mod_security test
```

4. Request *test.html* with malicious data in the **GET** variable of the HTTP request:

```
$ curl localhost/test.html?data=xxxevilxxx

<!DOCTYPE HTML PUBLIC "-//IETF//DTD HTML 2.0//EN">
<html><head>
<title>403 Forbidden</title>
</head><body>
<h1>Forbidden</h1>
<p>You don't have permission to access this resource.</p>
</body></html>
```

5. Check the **/var/log/httpd/error_log** file, and locate the log entry about denying access with the **param data containing an evil data** message:

```
[Wed May 25 08:01:31.036297 2022] [:error] [pid 5839:tid 139874434791168] [client
::1:45658] [client ::1] ModSecurity: Access denied with code 403 (phase 2). String match
"evil" at ARGS:data. [file "/etc/httpd/conf.d/mod_security.conf"] [line "4"] [id "1"] [msg "param
data contains evil data"] [hostname "localhost"] [uri "/test.html"] [unique_id
"Yo4amwIdsBG3yZqSzh2GuwAAAIY"]
```

Additional resources

- ModSecurity Wiki

1.12. INSTALLING THE APACHE HTTP SERVER MANUAL

This section describes how to install the Apache HTTP Server manual. This manual provides a detailed documentation of, for example:

- Configuration parameters and directives

- Performance tuning

- Authentication settings

- Modules

- Content caching

- Security tips

- Configuring TLS encryption

After installing the manual, you can display it using a web browser.

Prerequisites

- The Apache HTTP Server is installed and running.

Procedure

1. Install the **httpd-manual** package:

   ```
   # yum install httpd-manual
   ```

2. Optional: By default, all clients connecting to the Apache HTTP Server can display the manual. To restrict access to a specific IP range, such as the **192.0.2.0/24** subnet, edit the **/etc/httpd/conf.d/manual.conf** file and add the **Require ip 192.0.2.0/24** setting to the **<Directory "/usr/share/httpd/manual">** directive:

   ```
   <Directory "/usr/share/httpd/manual">
   ...
       **Require ip 192.0.2.0/24**
   ...
   </Directory>
   ```

3. Restart the **httpd** service:

   ```
   # systemctl restart httpd
   ```

Verification steps

1. To display the Apache HTTP Server manual, connect with a web browser to **http://***host_name_or_IP_address***/manual/**

1.13. WORKING WITH APACHE MODULES

The **httpd** service is a modular application, and you can extend it with a number of *Dynamic Shared Objects* (**DSO**s). *Dynamic Shared Objects* are modules that you can dynamically load or unload at runtime as necessary. You can find these modules in the **/usr/lib64/httpd/modules/** directory.

1.13.1. Loading a DSO module

As an administrator, you can choose the functionality to include in the server by configuring which modules the server should load. To load a particular DSO module, use the **LoadModule** directive. Note

that modules provided by a separate package often have their own configuration file in the **/etc/httpd/conf.modules.d/** directory.

Prerequisites

- You have installed the **httpd** package.

Procedure

1. Search for the module name in the configuration files in the **/etc/httpd/conf.modules.d/** directory:

   > # **grep mod_ssl.so /etc/httpd/conf.modules.d/***

2. Edit the configuration file in which the module name was found, and uncomment the **LoadModule** directive of the module:

   > **LoadModule ssl_module modules/mod_ssl.so**

3. If the module was not found, for example, because a RHEL package does not provide the module, create a configuration file, such as **/etc/httpd/conf.modules.d/30-example.conf** with the following directive:

   > **LoadModule ssl_module modules/<custom_module>.so**

4. Restart the **httpd** service:

   > # **systemctl restart httpd**

1.13.2. Compiling a custom Apache module

You can create your own module and build it with the help of the **httpd-devel** package, which contains the include files, the header files, and the **APache eXtenSion** (**apxs**) utility required to compile a module.

Prerequisites

- You have the **httpd-devel** package installed.

Procedure

- Build a custom module with the following command:

  > # **apxs -i -a -c module_name.c**

Verification steps

- Load the module the same way as described in Loading a DSO module .

1.14. EXPORTING A PRIVATE KEY AND CERTIFICATES FROM AN NSS DATABASE TO USE THEM IN AN APACHE WEB SERVER CONFIGURATION

RHEL 8 no longer provides the **mod_nss** module for the Apache web server, and Red Hat recommends using the **mod_ssl** module. If you store your private key and certificates in a Network Security Services (NSS) database, for example, because you migrated the web server from RHEL 7 to RHEL 8, follow this procedure to extract the key and certificates in Privacy Enhanced Mail (PEM) format. You can then use the files in the **mod_ssl** configuration as described in Configuring TLS encryption on an Apache HTTP server.

This procedure assumes that the NSS database is stored in **/etc/httpd/alias/** and that you store the exported private key and certificates in the **/etc/pki/tls/** directory.

Prerequisites

- The private key, the certificate, and the certificate authority (CA) certificate are stored in an NSS database.

Procedure

1. List the certificates in the NSS database:

   ```
   # certutil -d /etc/httpd/alias/ -L
   Certificate Nickname         Trust Attributes
                                SSL,S/MIME,JAR/XPI

   Example CA                  C,,
   Example Server Certificate     u,u,u
   ```

   You need the nicknames of the certificates in the next steps.

2. To extract the private key, you must temporarily export the key to a PKCS #12 file:

   a. Use the nickname of the certificate associated with the private key, to export the key to a PKCS #12 file:

   ```
   # pk12util -o /etc/pki/tls/private/export.p12 -d /etc/httpd/alias/ -n "Example Server Certificate"
   Enter password for PKCS12 file: password
   Re-enter password: password
   pk12util: PKCS12 EXPORT SUCCESSFUL
   ```

   Note that you must set a password on the PKCS #12 file. You need this password in the next step.

   b. Export the private key from the PKCS #12 file:

   ```
   # openssl pkcs12 -in /etc/pki/tls/private/export.p12 -out /etc/pki/tls/private/server.key -nocerts -nodes
   Enter Import Password: password
   MAC verified OK
   ```

   c. Delete the temporary PKCS #12 file:

> # rm /etc/pki/tls/private/export.p12

3. Set the permissions on **/etc/pki/tls/private/server.key** to ensure that only the **root** user can access this file:

   > # chown root:root /etc/pki/tls/private/server.key
   > # chmod 0600 /etc/pki/tls/private/server.key

4. Use the nickname of the server certificate in the NSS database to export the CA certificate:

   > # certutil -d /etc/httpd/alias/ -L -n "Example Server Certificate" -a -o
   > /etc/pki/tls/certs/server.crt

5. Set the permissions on **/etc/pki/tls/certs/server.crt** to ensure that only the **root** user can access this file:

   > # chown root:root /etc/pki/tls/certs/server.crt
   > # chmod 0600 /etc/pki/tls/certs/server.crt

6. Use the nickname of the CA certificate in the NSS database to export the CA certificate:

   > # certutil -d /etc/httpd/alias/ -L -n "*Example CA*" -a -o /etc/pki/tls/certs/ca.crt

7. Follow Configuring TLS encryption on an Apache HTTP server to configure the Apache web server, and:

   - Set the **SSLCertificateKeyFile** parameter to **/etc/pki/tls/private/server.key**.

   - Set the **SSLCertificateFile** parameter to **/etc/pki/tls/certs/server.crt**.

   - Set the **SSLCACertificateFile** parameter to **/etc/pki/tls/certs/ca.crt**.

Additional resources

- **certutil(1)** man page

- **pk12util(1)** man page

- **pkcs12(1ssl)** man page

1.15. ADDITIONAL RESOURCES

- **httpd(8)**

- **httpd.service(8)**

- **httpd.conf(5)**

- **apachectl(8)**

- Using GSS-Proxy for Apache httpd operation.

- Configuring applications to use cryptographic hardware through PKCS #11.

CHAPTER 2. SETTING UP AND CONFIGURING NGINX

NGINX is a high performance and modular server that you can use, for example, as a:

- Web server

- Reverse proxy

- Load balancer

This section describes how to NGINX in these scenarios.

2.1. INSTALLING AND PREPARING NGINX

Red Hat uses Application Streams to provide different versions of NGINX. This section describes how to:

- Select a stream and install NGINX

- Open the required ports in the firewall

- Enable and start the **nginx** service

Using the default configuration, NGINX runs as a web server on port **80** and provides content from the **/usr/share/nginx/html/** directory.

Prerequisites

- RHEL 8 is installed.

- The host is subscribed to the Red Hat Customer Portal.

- The **firewalld** service is enabled and started.

Procedure

1. Display the available NGINX module streams:

   ```
   # yum module list nginx
   Red Hat Enterprise Linux 8 for x86_64 - AppStream (RPMs)
   Name        Stream      Profiles       Summary
   nginx       1.14 [d]     common [d]     nginx webserver
   nginx       1.16         common [d]     nginx webserver
   ...

   Hint: [d]efault, [e]nabled, [x]disabled, [i]nstalled
   ```

2. If you want to install a different stream than the default, select the stream:

   ```
   # yum module enable nginx:stream_version
   ```

3. Install the **nginx** package:

   ```
   # yum install nginx
   ```

4. Open the ports on which NGINX should provide its service in the firewall. For example, to open the default ports for HTTP (port 80) and HTTPS (port 443) in **firewalld**, enter:

```
# firewall-cmd --permanent --add-port={80/tcp,443/tcp}
# firewall-cmd --reload
```

5. Enable the **nginx** service to start automatically when the system boots:

```
# systemctl enable nginx
```

6. Optionally, start the **nginx** service:

```
# systemctl start nginx
```

If you do not want to use the default configuration, skip this step, and configure NGINX accordingly before you start the service.

Verification steps

1. Use the **yum** utility to verify that the **nginx** package is installed:

```
# yum list installed nginx
Installed Packages
nginx.x86_64    1:1.14.1-9.module+el8.0.0+4108+af250afe    @rhel-8-for-x86_64-appstream-
rpms
```

2. Ensure that the ports on which NGINX should provide its service are opened in the firewalld:

```
# firewall-cmd --list-ports
80/tcp 443/tcp
```

3. Verify that the **nginx** service is enabled:

```
# systemctl is-enabled nginx
enabled
```

Additional resources

- For details about Subscription Manager, see the Using and Configuring Subscription Manager guide.

- For further details about Application Streams, modules, and installing packages, see the Installing, managing, and removing user-space components guide.

- For details about configuring firewalls, see the Securing networks guide.

2.2. CONFIGURING NGINX AS A WEB SERVER THAT PROVIDES DIFFERENT CONTENT FOR DIFFERENT DOMAINS

By default, NGINX acts as a web server that provides the same content to clients for all domain names associated with the IP addresses of the server. This procedure explains how to configure NGINX:

- To serve requests to the **example.com** domain with content from the **/var/www/example.com/** directory

- To serve requests to the **example.net** domain with content from the **/var/www/example.net/** directory

- To serve all other requests, for example, to the IP address of the server or to other domains associated with the IP address of the server, with content from the **/usr/share/nginx/html/** directory

Prerequisites

- NGINX is installed

- Clients and the web server resolve the **example.com** and **example.net** domain to the IP address of the web server.
  Note that you must manually add these entries to your DNS server.

Procedure

1. Edit the **/etc/nginx/nginx.conf** file:

   a. By default, the **/etc/nginx/nginx.conf** file already contains a catch-all configuration. If you have deleted this part from the configuration, re-add the following **server** block to the **http** block in the **/etc/nginx/nginx.conf** file:

   ```
   server {
       listen      80 default_server;
       listen      [::]:80 default_server;
       server_name  _;
       root        /usr/share/nginx/html;
   }
   ```

   These settings configure the following:

   - The **listen** directive define which IP address and ports the service listens. In this case, NGINX listens on port **80** on both all IPv4 and IPv6 addresses. The **default_server** parameter indicates that NGINX uses this **server** block as the default for requests matching the IP addresses and ports.

   - The **server_name** parameter defines the host names for which this **server** block is responsible. Setting **server_name** to _ configures NGINX to accept any host name for this **server** block.

   - The **root** directive sets the path to the web content for this **server** block.

   b. Append a similar **server** block for the **example.com** domain to the **http** block:

   ```
   server {
       server_name  example.com;
       root        /var/www/example.com/;
       access_log   /var/log/nginx/example.com/access.log;
       error_log    /var/log/nginx/example.com/error.log;
   }
   ```

   - The **access_log** directive defines a separate access log file for this domain.

- The **error_log** directive defines a separate error log file for this domain.

c. Append a similar **server** block for the **example.net** domain to the **http** block:

```
server {
    server_name  example.net;
    root         /var/www/example.net/;
    access_log   /var/log/nginx/example.net/access.log;
    error_log    /var/log/nginx/example.net/error.log;
}
```

2. Create the root directories for both domains:

```
# mkdir -p /var/www/example.com/
# mkdir -p /var/www/example.net/
```

3. Set the **httpd_sys_content_t** context on both root directories:

```
# semanage fcontext -a -t httpd_sys_content_t "/var/www/example.com(/.*)?"
# restorecon -Rv /var/www/example.com/
# semanage fcontext -a -t httpd_sys_content_t "/var/www/example.net(/.\*)?"
# restorecon -Rv /var/www/example.net/
```

These commands set the **httpd_sys_content_t** context on the **/var/www/example.com/** and **/var/www/example.net/** directories.

Note that you must install the **policycoreutils-python-utils** package to run the **restorecon** commands.

4. Create the log directories for both domains:

```
# mkdir /var/log/nginx/example.com/
# mkdir /var/log/nginx/example.net/
```

5. Restart the **nginx** service:

```
# systemctl restart nginx
```

Verification steps

1. Create a different example file in each virtual host's document root:

```
# echo "Content for example.com" > /var/www/example.com/index.html
# echo "Content for example.net" > /var/www/example.net/index.html
# echo "Catch All content" > /usr/share/nginx/html/index.html
```

2. Use a browser and connect to **http://example.com**. The web server shows the example content from the **/var/www/example.com/index.html** file.

3. Use a browser and connect to **http://example.net**. The web server shows the example content from the **/var/www/example.net/index.html** file.

4. Use a browser and connect to **http://***IP_address_of_the_server*. The web server shows the example content from the **/usr/share/nginx/html/index.html** file.

2.3. ADDING TLS ENCRYPTION TO AN NGINX WEB SERVER

This section describes how to enable TLS encryption on an NGINX web server for the **example.com** domain.

Prerequisites

- NGINX is installed.

- The private key is stored in the **/etc/pki/tls/private/example.com.key** file.
  For details about creating a private key and certificate signing request (CSR), as well as how to request a certificate from a certificate authority (CA), see your CA's documentation.

- The TLS certificate is stored in the **/etc/pki/tls/certs/example.com.crt** file. If you use a different path, adapt the corresponding steps of the procedure.

- The CA certificate has been appended to the TLS certificate file of the server.

- Clients and the web server resolve the host name of the server to the IP address of the web server.

- Port **443** is open in the local firewall.

Procedure

1. Edit the **/etc/nginx/nginx.conf** file, and add the following **server** block to the **http** block in the configuration:

   ```
   server {
       listen              443 ssl;
       server_name         example.com;
       root                /usr/share/nginx/html;
       ssl_certificate     /etc/pki/tls/certs/example.com.crt;
       ssl_certificate_key /etc/pki/tls/private/example.com.key;
   }
   ```

2. For security reasons, configure that only the **root** user can access the private key file:

   ```
   # chown root:root /etc/pki/tls/private/example.com.key
   # chmod 600 /etc/pki/tls/private/example.com.key
   ```

   > **WARNING**
   >
   > If the private key was accessed by unauthorized users, revoke the certificate, create a new private key, and request a new certificate. Otherwise, the TLS connection is no longer secure.

3. Restart the **nginx** service:

> # **systemctl restart nginx**

Verification steps

- Use a browser and connect to **https://example.com**

Additional resources

- [Security considerations for TLS in RHEL](#)

2.4. CONFIGURING NGINX AS A REVERSE PROXY FOR THE HTTP TRAFFIC

You can configure the NGINX web server to act as a reverse proxy for HTTP traffic. For example, you can use this functionality to forward requests to a specific subdirectory on a remote server. From the client perspective, the client loads the content from the host it accesses. However, NGINX loads the actual content from the remote server and forwards it to the client.

This procedure explains how to forward traffic to the **/example** directory on the web server to the URL **https://example.com**.

Prerequisites

- NGINX is installed

- Optional: TLS encryption is enabled on the reverse proxy.

Procedure

1. Edit the **/etc/nginx/nginx.conf** file and add the following settings to the **server** block that should provide the reverse proxy:

   > location /example {
   >   proxy_pass https://example.com;
   > }

   The **location** block defines that NGINX passes all requests in the **/example** directory to **https://example.com**.

2. Set the **httpd_can_network_connect** SELinux boolean parameter to **1** to configure that SELinux allows NGINX to forward traffic:

   > # **setsebool -P httpd_can_network_connect 1**

3. Restart the **nginx** service:

   > # **systemctl restart nginx**

Verification steps

- Use a browser and connect to **http://***host_name***/example** and the content of **https://example.com** is shown.

2.5. CONFIGURING NGINX AS AN HTTP LOAD BALANCER

You can use the NGINX reverse proxy feature to load-balance traffic. This procedure describes how to configure NGINX as an HTTP load balancer that sends requests to different servers, based on which of them has the least number of active connections. If both servers are not available, the procedure also defines a third host for fallback reasons.

Prerequisites

- NGINX is installed

Procedure

1. Edit the **/etc/nginx/nginx.conf** file and add the following settings:

   ```
   http {
       upstream backend {
           least_conn;
           server server1.example.com;
           server server2.example.com;
           server server3.example.com backup;
       }

       server {
           location / {
               proxy_pass http://backend;
           }
       }
   }
   ```

   The **least_conn** directive in the host group named **backend** defines that NGINX sends requests to **server1.example.com** or **server2.example.com**, depending on which host has the least number of active connections. NGINX uses **server3.example.com** only as a backup in case that the other two hosts are not available.

   With the **proxy_pass** directive set to **http://backend**, NGINX acts as a reverse proxy and uses the **backend** host group to distribute requests based on the settings of this group.

   Instead of the **least_conn** load balancing method, you can specify:

   - No method to use round robin and distribute requests evenly across servers.

   - **ip_hash** to send requests from one client address to the same server based on a hash calculated from the first three octets of the IPv4 address or the whole IPv6 address of the client.

   - **hash** to determine the server based on a user-defined key, which can be a string, a variable, or a combination of both. The **consistent** parameter configures that NGINX distributes requests across all servers based on the user-defined hashed key value.

   - **random** to send requests to a randomly selected server.

2. Restart the **nginx** service:

   ```
   # systemctl restart nginx
   ```

2.6. ADDITIONAL RESOURCES

- For the official NGINX documentation see https://nginx.org/en/docs/. Note that Red Hat does not maintain this documentation and that it might not work with the NGINX version you have installed.

- Configuring applications to use cryptographic hardware through PKCS #11 .

CHAPTER 3. USING SAMBA AS A SERVER

Samba implements the Server Message Block (SMB) protocol in Red Hat Enterprise Linux. The SMB protocol is used to access resources on a server, such as file shares and shared printers. Additionally, Samba implements the Distributed Computing Environment Remote Procedure Call (DCE RPC) protocol used by Microsoft Windows.

You can run Samba as:

- An Active Directory (AD) or NT4 domain member

- A standalone server

- An NT4 Primary Domain Controller (PDC) or Backup Domain Controller (BDC)

> **NOTE**
>
> Red Hat supports the PDC and BDC modes only in existing installations with Windows versions which support NT4 domains. Red Hat recommends not setting up a new Samba NT4 domain, because Microsoft operating systems later than Windows 7 and Windows Server 2008 R2 do not support NT4 domains.
>
> Red Hat does not support running Samba as an AD domain controller (DC).

Independently of the installation mode, you can optionally share directories and printers. This enables Samba to act as a file and print server.

3.1. UNDERSTANDING THE DIFFERENT SAMBA SERVICES AND MODES

This section describes the different services included in Samba and the different modes you can configure.

3.1.1. The Samba services

Samba provides the following services:

smbd

This service provides file sharing and printing services using the SMB protocol. Additionally, the service is responsible for resource locking and for authenticating connecting users. For authenticating domain members, **smbd** requires **winbindd**. The **smb systemd** service starts and stops the **smbd** daemon.
To use the **smbd** service, install the **samba** package.

nmbd

This service provides host name and IP resolution using the NetBIOS over IPv4 protocol. Additionally to the name resolution, the **nmbd** service enables browsing the SMB network to locate domains, work groups, hosts, file shares, and printers. For this, the service either reports this information directly to the broadcasting client or forwards it to a local or master browser. The **nmb systemd** service starts and stops the **nmbd** daemon.
Note that modern SMB networks use DNS to resolve clients and IP addresses. For Kerberos a working DNS setup is required.

To use the **nmbd** service, install the **samba** package.

winbindd

This service provides an interface for the Name Service Switch (NSS) to use AD or NT4 domain users and groups on the local system. This enables, for example, domain users to authenticate to services hosted on a Samba server or to other local services. The **winbind systemd** service starts and stops the **winbindd** daemon.

If you set up Samba as a domain member, **winbindd** must be started before the **smbd** service. Otherwise, domain users and groups are not available to the local system..

To use the **winbindd** service, install the **samba-winbind** package.
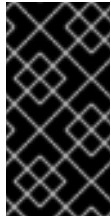
> **IMPORTANT**
>
> Red Hat only supports running Samba as a server with the **winbindd** service to provide domain users and groups to the local system. Due to certain limitations, such as missing Windows access control list (ACL) support and NT LAN Manager (NTLM) fallback, SSSD is not supported.

3.1.2. The Samba security services

The **security** parameter in the **[global]** section in the **/etc/samba/smb.conf** file manages how Samba authenticates users that are connecting to the service. Depending on the mode you install Samba in, the parameter must be set to different values:

On an AD domain member, set **security = ads**

In this mode, Samba uses Kerberos to authenticate AD users.
For details about setting up Samba as a domain member, see Setting up Samba as an AD domain member server.

On a standalone server, set **security = user**

In this mode, Samba uses a local database to authenticate connecting users.
For details about setting up Samba as a standalone server, see Setting up Samba as a standalone server.

On an NT4 PDC or BDC, set **security = user**

In this mode, Samba authenticates users to a local or LDAP database.

On an NT4 domain member, set **security = domain**

In this mode, Samba authenticates connecting users to an NT4 PDC or BDC. You cannot use this mode on AD domain members.
For details about setting up Samba as a domain member, see Setting up Samba as an AD domain member server.

Additional resources

- **security** parameter in the **smb.conf(5)** man page

3.1.3. Scenarios when Samba services and Samba client utilities load and reload their configuration

The following describes when Samba services and utilities load and reload their configuration:

- Samba services reload their configuration:

  - Automatically every 3 minutes

  - On manual request, for example, when you run the **smbcontrol all reload-config** command.

- Samba client utilities read their configuration only when you start them.

Note that certain parameters, such as **security** require a restart of the **smb** service to take effect and a reload is not sufficient.

Additional resources

- The **How configuration changes are applied** section in the **smb.conf(5)** man page

- The **smbd(8)**, **nmbd(8)**, and **winbindd(8)** man pages

3.1.4. Editing the Samba configuration in a safe way

Samba services automatically reload their configuration every 3 minutes. This procedure describes how to edit the Samba configuration in a way that prevents the services reload the changes before you have verified the configuration using the **testparm** utility.

Prerequisites

- Samba is installed.

Procedure

1. Create a copy of the **/etc/samba/smb.conf** file:

   # **cp** /etc/samba/smb.conf /etc/samba/samba.conf.*copy*

2. Edit the copied file and make the desired changes.

3. Verify the configuration in the **/etc/samba/samba.conf.*copy*** file:

   # **testparm -s** /etc/samba/samba.conf.*copy*

   If **testparm** reports errors, fix them and run the command again.

4. Override the **/etc/samba/smb.conf** file with the new configuration:

   # **mv** /etc/samba/samba.conf.*copy* /etc/samba/smb.conf

5. Wait until the Samba services automatically reload their configuration or manually reload the configuration:

   # **smbcontrol all reload-config**

Additional resources

- Scenarios when Samba services and Samba client utilities load and reload their configuration

3.2. VERIFYING THE SMB.CONF FILE BY USING THE TESTPARM UTILITY

The **testparm** utility verifies that the Samba configuration in the **/etc/samba/smb.conf** file is correct. The utility detects invalid parameters and values, but also incorrect settings, such as for ID mapping. If **testparm** reports no problem, the Samba services will successfully load the **/etc/samba/smb.conf** file. Note that **testparm** cannot verify that the configured services will be available or work as expected.

> **IMPORTANT**
>
> Red Hat recommends that you verify the **/etc/samba/smb.conf** file by using **testparm** after each modification of this file.

Prerequisites

- You installed Samba.

- The **/etc/samba/smb.conf** file exits.

Procedure

1. Run the **testparm** utility as the **root** user:

   > **# testparm**
   > Load smb config files from /etc/samba/smb.conf
   > rlimit_max: increasing rlimit_max (1024) to minimum Windows limit (16384)
   > **Unknown parameter encountered: "log levell"**
   > Processing section "[example_share]"
   > Loaded services file OK.
   > **ERROR: The idmap range for the domain * (tdb) overlaps with the range of DOMAIN (ad)!**
   >
   > Server role: ROLE_DOMAIN_MEMBER
   >
   > Press enter to see a dump of your service definitions
   >
   > # Global parameters
   > [global]
   >   ...
   >
   > [example_share]
   >   ...

   The previous example output reports a non-existent parameter and an incorrect ID mapping configuration.

2. If **testparm** reports incorrect parameters, values, or other errors in the configuration, fix the problem and run the utility again.

3.3. SETTING UP SAMBA AS A STANDALONE SERVER

You can set up Samba as a server that is not a member of a domain. In this installation mode, Samba authenticates users to a local database instead of to a central DC. Additionally, you can enable guest access to allow users to connect to one or multiple services without authentication.

3.3.1. Setting up the server configuration for the standalone server

This section describes how to set up the server configuration for a Samba standalone server.

Procedure

1. Install the **samba** package:

   ```
   # yum install samba
   ```

2. Edit the **/etc/samba/smb.conf** file and set the following parameters:

   ```
   [global]
    workgroup = Example-WG
    netbios name = Server
    security = user

    log file = /var/log/samba/%m.log
    log level = 1
   ```

   This configuration defines a standalone server named **Server** within the **Example-WG** work group. Additionally, this configuration enables logging on a minimal level (**1**) and log files will be stored in the **/var/log/samba/** directory. Samba will expand the **%m** macro in the **log file** parameter to the NetBIOS name of connecting clients. This enables individual log files for each client.

3. Optionally, configure file or printer sharing. See:

   - Setting up a share that uses POSIX ACLs

   - Setting up a share that uses Windows ACLs

   - Setting up Samba as a Print Server

4. Verify the **/etc/samba/smb.conf** file:

   ```
   # testparm
   ```

5. If you set up shares that require authentication, create the user accounts.
   For details, see Creating and enabling local user accounts .

6. Open the required ports and reload the firewall configuration by using the **firewall-cmd** utility:

   ```
   # firewall-cmd --permanent --add-service=samba
   # firewall-cmd --reload
   ```

7. Enable and start the **smb** service:

   ```
   # systemctl enable --now smb
   ```

Additional resources

- **smb.conf(5)** man page

3.3.2. Creating and enabling local user accounts

To enable users to authenticate when they connect to a share, you must create the accounts on the Samba host both in the operating system and in the Samba database. Samba requires the operating system account to validate the Access Control Lists (ACL) on file system objects and the Samba account to authenticate connecting users.

If you use the **passdb backend = tdbsam** default setting, Samba stores user accounts in the /**var**/**lib**/**samba**/**private**/**passdb.tdb** database.

The procedure in this section describes how to create a local Samba user named **example**.

Prerequisites

- Samba is installed and configured as a standalone server.

Procedure

1. Create the operating system account:

   > # **useradd -M -s** /**sbin**/**nologin example**

   This command adds the **example** account without creating a home directory. If the account is only used to authenticate to Samba, assign the /**sbin**/**nologin** command as shell to prevent the account from logging in locally.

2. Set a password to the operating system account to enable it:

   > # **passwd example**
   > Enter new UNIX password: *password*
   > Retype new UNIX password: *password*
   > passwd: password updated successfully

   Samba does not use the password set on the operating system account to authenticate. However, you need to set a password to enable the account. If an account is disabled, Samba denies access if this user connects.

3. Add the user to the Samba database and set a password to the account:

   > # **smbpasswd -a** *example*
   > New SMB password: *password*
   > Retype new SMB password: *password*
   > Added user example.

   Use this password to authenticate when using this account to connect to a Samba share.

4. Enable the Samba account:

   > # **smbpasswd -e** *example*
   > Enabled user example.

3.4. UNDERSTANDING AND CONFIGURING SAMBA ID MAPPING

Windows domains distinguish users and groups by unique Security Identifiers (SID). However, Linux requires unique UIDs and GIDs for each user and group. If you run Samba as a domain member, the **winbindd** service is responsible for providing information about domain users and groups to the operating system.

To enable the **winbindd** service to provide unique IDs for users and groups to Linux, you must configure ID mapping in the **/etc/samba/smb.conf** file for:

- The local database (default domain)

- The AD or NT4 domain the Samba server is a member of

- Each trusted domain from which users must be able to access resources on this Samba server

Samba provides different ID mapping back ends for specific configurations. The most frequently used back ends are:

Back end	Use case
tdb	The * default domain only
ad	AD domains only
rid	AD and NT4 domains
autorid	AD, NT4, and the * default domain

3.4.1. Planning Samba ID ranges

Regardless of whether you store the Linux UIDs and GIDs in AD or if you configure Samba to generate them, each domain configuration requires a unique ID range that must not overlap with any of the other domains.

> **WARNING**
>
> If you set overlapping ID ranges, Samba fails to work correctly.

Example 3.1. Unique ID Ranges

The following shows non-overlapping ID mapping ranges for the default (*), **AD-DOM**, and the **TRUST-DOM** domains.

```
[global]
...
idmap config * : backend = tdb
```

```
idmap config * : range = 10000-999999

idmap config AD-DOM:backend = rid
idmap config AD-DOM:range = 2000000-2999999

idmap config TRUST-DOM:backend = rid
idmap config TRUST-DOM:range = 4000000-4999999
```

> **IMPORTANT**
>
> You can only assign one range per domain. Therefore, leave enough space between the domains ranges. This enables you to extend the range later if your domain grows.
>
> If you later assign a different range to a domain, the ownership of files and directories previously created by these users and groups will be lost.

3.4.2. The * default domain

In a domain environment, you add one ID mapping configuration for each of the following:

- The domain the Samba server is a member of

- Each trusted domain that should be able to access the Samba server

However, for all other objects, Samba assigns IDs from the default domain. This includes:

- Local Samba users and groups

- Samba built-in accounts and groups, such as **BUILTIN\Administrators**

> **IMPORTANT**
>
> You must configure the default domain as described in this section to enable Samba to operate correctly.

The default domain back end must be writable to permanently store the assigned IDs.

For the default domain, you can use one of the following back ends:

tdb

When you configure the default domain to use the **tdb** back end, set an ID range that is big enough to include objects that will be created in the future and that are not part of a defined domain ID mapping configuration.
For example, set the following in the **[global]** section in the **/etc/samba/smb.conf** file:

```
idmap config * : backend = tdb
idmap config * : range = 10000-999999
```

For further details, see Using the TDB ID mapping back end .

autorid

When you configure the default domain to use the **autorid** back end, adding additional ID mapping configurations for domains is optional.

For example, set the following in the **[global]** section in the /**etc**/**samba**/**smb.conf** file:

```
idmap config * : backend = autorid
idmap config * : range = 10000-999999
```

For further details, see Using the autorid ID mapping back end .

3.4.3. Using the tdb ID mapping back end

The **winbindd** service uses the writable **tdb** ID mapping back end by default to store Security Identifier (SID), UID, and GID mapping tables. This includes local users, groups, and built-in principals.

Use this back end only for the ***** default domain. For example:

```
idmap config * : backend = tdb
idmap config * : range = 10000-999999
```

Additional resources

- The * default domain .

3.4.4. Using the ad ID mapping back end

This section describes how to configure a Samba AD member to use the **ad** ID mapping back end.

The **ad** ID mapping back end implements a read-only API to read account and group information from AD. This provides the following benefits:

- All user and group settings are stored centrally in AD.

- User and group IDs are consistent on all Samba servers that use this back end.

- The IDs are not stored in a local database which can corrupt, and therefore file ownerships cannot be lost.

> **NOTE**
>
> The **ad** ID mapping back end does not support Active Directory domains with one-way trusts. If you configure a domain member in an Active Directory with one-way trusts, use instead one of the following ID mapping back ends: **tdb**, **rid**, or **autorid**.

The ad back end reads the following attributes from AD:

AD attribute name	Object type	Mapped to
sAMAccountName	User and group	User or group name, depending on the object
uidNumber	User	User ID (UID)

AD attribute name	Object type	Mapped to
gidNumber	Group	Group ID (GID)
loginShell [a]	User	Path to the shell of the user
unixHomeDirectory [a]	User	Path to the home directory of the user
primaryGroupID [b]	User	Primary group ID

[a] Samba only reads this attribute if you set **idmap config *DOMAIN*:unix_nss_info = yes**.

[b] Samba only reads this attribute if you set **idmap config *DOMAIN*:unix_primary_group = yes**.

Prerequisites

- Both users and groups must have unique IDs set in AD, and the IDs must be within the range configured in the **/etc/samba/smb.conf** file. Objects whose IDs are outside of the range will not be available on the Samba server.

- Users and groups must have all required attributes set in AD. If required attributes are missing, the user or group will not be available on the Samba server. The required attributes depend on your configuration. .Prerequisites

- You installed Samba.

- The Samba configuration, except ID mapping, exists in the **/etc/samba/smb.conf** file.

Procedure

1. Edit the **[global]** section in the **/etc/samba/smb.conf** file:

   a. Add an ID mapping configuration for the default domain (*****) if it does not exist. For example:

      ```
      idmap config * : backend = tdb
      idmap config * : range = 10000-999999
      ```

   b. Enable the **ad** ID mapping back end for the AD domain:

      ```
      idmap config DOMAIN : backend = ad
      ```

   c. Set the range of IDs that is assigned to users and groups in the AD domain. For example:

      ```
      idmap config DOMAIN : range = 2000000-2999999
      ```

> **IMPORTANT**
>
> The range must not overlap with any other domain configuration on this server. Additionally, the range must be set big enough to include all IDs assigned in the future. For further details, see Planning Samba ID ranges.

d. Set that Samba uses the RFC 2307 schema when reading attributes from AD:

```
idmap config DOMAIN : schema_mode = rfc2307
```

e. To enable Samba to read the login shell and the path to the users home directory from the corresponding AD attribute, set:

```
idmap config DOMAIN : unix_nss_info = yes
```

Alternatively, you can set a uniform domain-wide home directory path and login shell that is applied to all users. For example:

```
template shell = /bin/bash
template homedir = /home/%U
```

f. By default, Samba uses the **primaryGroupID** attribute of a user object as the user's primary group on Linux. Alternatively, you can configure Samba to use the value set in the **gidNumber** attribute instead:

```
idmap config DOMAIN : unix_primary_group = yes
```

2. Verify the **/etc/samba/smb.conf** file:

```
# testparm
```

3. Reload the Samba configuration:

```
# smbcontrol all reload-config
```

Additional resources

- The * default domain

- **smb.conf(5)** and **idmap_ad(8)** man pages

- **VARIABLE SUBSTITUTIONS** section in the **smb.conf(5)** man page

3.4.5. Using the rid ID mapping back end

This section describes how to configure a Samba domain member to use the **rid** ID mapping back end.

Samba can use the relative identifier (RID) of a Windows SID to generate an ID on Red Hat Enterprise Linux.

NOTE

The RID is the last part of a SID. For example, if the SID of a user is **S-1-5-21-5421822485-1151247151-421485315-30014**, then **30014** is the corresponding RID.

The **rid** ID mapping back end implements a read-only API to calculate account and group information based on an algorithmic mapping scheme for AD and NT4 domains. When you configure the back end, you must set the lowest and highest RID in the **idmap config *DOMAIN* : *range*** parameter. Samba will not map users or groups with a lower or higher RID than set in this parameter.
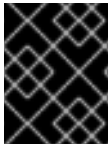
IMPORTANT

As a read-only back end, **rid** cannot assign new IDs, such as for **BUILTIN** groups. Therefore, do not use this back end for the ***** default domain.

Benefits of using the rid back end

- All domain users and groups that have an RID within the configured range are automatically available on the domain member.

- You do not need to manually assign IDs, home directories, and login shells.

Drawbacks of using the rid back end

- All domain users get the same login shell and home directory assigned. However, you can use variables.

- User and group IDs are only the same across Samba domain members if all use the **rid** back end with the same ID range settings.

- You cannot exclude individual users or groups from being available on the domain member. Only users and groups outside of the configured range are excluded.

- Based on the formulas the **winbindd** service uses to calculate the IDs, duplicate IDs can occur in multi-domain environments if objects in different domains have the same RID.

Prerequisites

- You installed Samba.

- The Samba configuration, except ID mapping, exists in the **/etc/samba/smb.conf** file.

Procedure

1. Edit the **[global]** section in the **/etc/samba/smb.conf** file:

   a. Add an ID mapping configuration for the default domain (*****) if it does not exist. For example:

   ```
   idmap config * : backend = tdb
   idmap config * : range = 10000-999999
   ```

   b. Enable the **rid** ID mapping back end for the domain:

   ```
   idmap config DOMAIN : backend = rid
   ```

c. Set a range that is big enough to include all RIDs that will be assigned in the future. For example:

```
idmap config DOMAIN : range = 2000000-2999999
```

Samba ignores users and groups whose RIDs in this domain are not within the range.

> **IMPORTANT**
>
> The range must not overlap with any other domain configuration on this server. Additionally, the range must be set big enough to include all IDs assigned in the future. For further details, see Planning Samba ID ranges.

d. Set a shell and home directory path that will be assigned to all mapped users. For example:

```
template shell = /bin/bash
template homedir = /home/%U
```

2. Verify the **/etc/samba/smb.conf** file:

```
# testparm
```

3. Reload the Samba configuration:

```
# smbcontrol all reload-config
```

Additional resources

- The * default domain

- **VARIABLE SUBSTITUTIONS** section in the **smb.conf(5)** man page

- Calculation of the local ID from a RID, see the **idmap_rid(8)** man page

3.4.6. Using the autorid ID mapping back end

This section describes how to configure a Samba domain member to use the **autorid** ID mapping back end.

The **autorid** back end works similar to the **rid** ID mapping back end, but can automatically assign IDs for different domains. This enables you to use the **autorid** back end in the following situations:

- Only for the * default domain

- For the * default domain and additional domains, without the need to create ID mapping configurations for each of the additional domains

- Only for specific domains
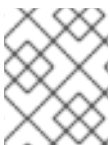
> **NOTE**
>
> If you use **autorid** for the default domain, adding additional ID mapping configuration for domains is optional.

Parts of this section were adopted from the idmap config autorid documentation published in the Samba Wiki. License: CC BY 4.0. Authors and contributors: See the history tab on the Wiki page.

Benefits of using the autorid back end

- All domain users and groups whose calculated UID and GID is within the configured range are automatically available on the domain member.

- You do not need to manually assign IDs, home directories, and login shells.

- No duplicate IDs, even if multiple objects in a multi-domain environment have the same RID.

Drawbacks

- User and group IDs are not the same across Samba domain members.

- All domain users get the same login shell and home directory assigned. However, you can use variables.

- You cannot exclude individual users or groups from being available on the domain member. Only users and groups whose calculated UID or GID is outside of the configured range are excluded.

Prerequisites

- You installed Samba.

- The Samba configuration, except ID mapping, exists in the **/etc/samba/smb.conf** file.

Procedure

1. Edit the **[global]** section in the **/etc/samba/smb.conf** file:

   a. Enable the **autorid** ID mapping back end for the ***** default domain:

      ```
      idmap config * : backend = autorid
      ```

   b. Set a range that is big enough to assign IDs for all existing and future objects. For example:

      ```
      idmap config * : range = 10000-999999
      ```

      Samba ignores users and groups whose calculated IDs in this domain are not within the range.

      > **WARNING**
      >
      > After you set the range and Samba starts using it, you can only increase the upper limit of the range. Any other change to the range can result in new ID assignments, and thus in losing file ownerships.

c. Optionally, set a range size. For example:

```
idmap config * : rangesize = 200000
```

Samba assigns this number of continuous IDs for each domain's object until all IDs from the range set in the **idmap config * : range** parameter are taken.

> **NOTE**
>
> If you set a rangesize, you need to adapt the range accordingly. The range needs to be a multiple of the rangesize.

d. Set a shell and home directory path that will be assigned to all mapped users. For example:

```
template shell = /bin/bash
template homedir = /home/%U
```

e. Optionally, add additional ID mapping configuration for domains. If no configuration for an individual domain is available, Samba calculates the ID using the **autorid** back end settings in the previously configured * default domain.

> **IMPORTANT**
>
> The range must not overlap with any other domain configuration on this server. Additionally, the range must be set big enough to include all IDs assigned in the future. For further details, see Planning Samba ID ranges.

2. Verify the **/etc/samba/smb.conf** file:

```
# testparm
```

3. Reload the Samba configuration:

```
# smbcontrol all reload-config
```

Additional resources

- **THE MAPPING FORMULAS** section in the **idmap_autorid(8)** man page

- **rangesize** parameter description in the **idmap_autorid(8)** man page

- **VARIABLE SUBSTITUTIONS** section in the **smb.conf(5)** man page

3.5. SETTING UP SAMBA AS AN AD DOMAIN MEMBER SERVER

If you are running an AD or NT4 domain, use Samba to add your Red Hat Enterprise Linux server as a member to the domain to gain the following:

- Access domain resources on other domain members

- Authenticate domain users to local services, such as **sshd**

- Share directories and printers hosted on the server to act as a file and print server

3.5.1. Joining a RHEL system to an AD domain

Samba Winbind is an alternative to the System Security Services Daemon (SSSD) for connecting a Red Hat Enterprise Linux (RHEL) system with Active Directory (AD). This section describes how to join a RHEL system to an AD domain by using **realmd** to configure Samba Winbind.

Procedure

1. If your AD requires the deprecated RC4 encryption type for Kerberos authentication, enable support for these ciphers in RHEL:

   ```
   # update-crypto-policies --set DEFAULT:AD-SUPPORT
   ```

2. Install the following packages:

   ```
   # yum install realmd oddjob-mkhomedir oddjob samba-winbind-clients \
       samba-winbind samba-common-tools samba-winbind-krb5-locator
   ```

3. To share directories or printers on the domain member, install the **samba** package:

   ```
   # yum install samba
   ```

4. Backup the existing **/etc/samba/smb.conf** Samba configuration file:

   ```
   # mv /etc/samba/smb.conf /etc/samba/smb.conf.bak
   ```

5. Join the domain. For example, to join a domain named **ad.example.com**:

   ```
   # realm join --membership-software=samba --client-software=winbind ad.example.com
   ```

   Using the previous command, the **realm** utility automatically:

   - Creates a **/etc/samba/smb.conf** file for a membership in the **ad.example.com** domain

   - Adds the **winbind** module for user and group lookups to the **/etc/nsswitch.conf** file

   - Updates the Pluggable Authentication Module (PAM) configuration files in the **/etc/pam.d/** directory

   - Starts the **winbind** service and enables the service to start when the system boots

6. Optionally, set an alternative ID mapping back end or customized ID mapping settings in the **/etc/samba/smb.conf** file. For details, see Understanding and configuring Samba ID mapping .

7. Verify that the **winbind** service is running:

   ```
   # systemctl status winbind
   ...
       Active: active (running) since Tue 2018-11-06 19:10:40 CET; 15s ago
   ```

> **IMPORTANT**
>
> To enable Samba to query domain user and group information, the **winbind** service must be running before you start **smb**.

8. If you installed the **samba** package to share directories and printers, enable and start the **smb** service:

   ```
   # systemctl enable --now smb
   ```

9. Optionally, if you are authenticating local logins to Active Directory, enable the **winbind_krb5_localauth** plug-in. See Using the local authorization plug-in for MIT Kerberos .

Verification steps

1. Display an AD user's details, such as the AD administrator account in the AD domain:

   ```
   # getent passwd "AD\administrator"
   AD\administrator:*:10000:10000::/home/administrator@AD:/bin/bash
   ```

2. Query the members of the domain users group in the AD domain:

   ```
   # getent group "AD\Domain Users"
      AD\domain users:x:10000:user1,user2
   ```

3. Optionally, verify that you can use domain users and groups when you set permissions on files and directories. For example, to set the owner of the **/srv/samba/example.txt** file to **AD\administrator** and the group to **AD\Domain Users**:

   ```
   # chown "AD\administrator":"AD\Domain Users" /srv/samba/example.txt
   ```

4. Verify that Kerberos authentication works as expected:

   a. On the AD domain member, obtain a ticket for the **administrator@AD.EXAMPLE.COM** principal:

      ```
      # kinit administrator@AD.EXAMPLE.COM
      ```

   b. Display the cached Kerberos ticket:

      ```
      # klist
      Ticket cache: KCM:0
      Default principal: administrator@AD.EXAMPLE.COM

      Valid starting       Expires            Service principal
      01.11.2018 10:00:00  01.11.2018 20:00:00
      krbtgt/AD.EXAMPLE.COM@AD.EXAMPLE.COM
            renew until 08.11.2018 05:00:00
      ```

5. Display the available domains:

   ```
   # wbinfo --all-domains
   BUILTIN
   ```

> SAMBA-SERVER
> AD

Additional resources

- If you do not want to use the deprecated RC4 ciphers, you can enable the AES encryption type in AD. See

- [Enabling the AES encryption type in Active Directory using a GPO](#) . Note that this can have an impact on other services in your AD.

- **realm(8)** man page

3.5.2. Using the local authorization plug-in for MIT Kerberos

The **winbind** service provides Active Directory users to the domain member. In certain situations, administrators want to enable domain users to authenticate to local services, such as an SSH server, which are running on the domain member. When using Kerberos to authenticate the domain users, enable the **winbind_krb5_localauth** plug-in to correctly map Kerberos principals to Active Directory accounts through the **winbind** service.

For example, if the **sAMAccountName** attribute of an Active Directory user is set to **EXAMPLE** and the user tries to log with the user name lowercase, Kerberos returns the user name in upper case. As a consequence, the entries do not match and authentication fails.

Using the **winbind_krb5_localauth** plug-in, the account names are mapped correctly. Note that this only applies to GSSAPI authentication and not for getting the initial ticket granting ticket (TGT).

Prerequisites

- Samba is configured as a member of an Active Directory.

- Red Hat Enterprise Linux authenticates log in attempts against Active Directory.

- The **winbind** service is running.

Procedure

Edit the **/etc/krb5.conf** file and add the following section:

```
[plugins]
localauth = {
    module = winbind:/usr/lib64/samba/krb5/winbind_krb5_localauth.so
    enable_only = winbind
}
```

Additional resources

- **winbind_krb5_localauth(8)** man page.

3.6. SETTING UP SAMBA ON AN IDM DOMAIN MEMBER

This section describes how to set up Samba on a host that is joined to a Red Hat Identity Management (IdM) domain. Users from IdM and also, if available, from trusted Active Directory (AD) domains, can access shares and printer services provided by Samba.

> **IMPORTANT**
>
> Using Samba on an IdM domain member is an unsupported Technology Preview feature and contains certain limitations. For example, IdM trust controllers do not support the Active Directory Global Catalog service, and they do not support resolving IdM groups using the Distributed Computing Environment / Remote Procedure Calls (DCE/RPC) protocols. As a consequence, AD users can only access Samba shares and printers hosted on IdM clients when logged in to other IdM clients; AD users logged into a Windows machine can not access Samba shares hosted on an IdM domain member.
>
> Customers deploying Samba on IdM domain members are encouraged to provide feedback to Red Hat.

Prerequisites

- The host is joined as a client to the IdM domain.

- Both the IdM servers and the client must run on RHEL 8.1 or later.

3.6.1. Preparing the IdM domain for installing Samba on domain members

Before you can set up Samba on an IdM client, you must prepare the IdM domain using the **ipa-adtrust-install** utility on an IdM server.

> **NOTE**
>
> Any system where you run the **ipa-adtrust-install** command automatically becomes an AD trust controller. However, you must run **ipa-adtrust-install** only once on an IdM server.

Prerequisites

- IdM server is installed.

- You need root privileges to install packages and restart IdM services.

Procedure

1. Install the required packages:

   [root@ipaserver ~]# **yum install ipa-server-trust-ad samba-client**

2. Authenticate as the IdM administrative user:

   [root@ipaserver ~]# **kinit admin**

3. Run the **ipa-adtrust-install** utility:

   [root@ipaserver ~]# **ipa-adtrust-install**

The DNS service records are created automatically if IdM was installed with an integrated DNS server.

If you installed IdM without an integrated DNS server, **ipa-adtrust-install** prints a list of service records that you must manually add to DNS before you can continue.

4. The script prompts you that the **/etc/samba/smb.conf** already exists and will be rewritten:

> WARNING: The smb.conf already exists. Running ipa-adtrust-install will break your existing Samba configuration.
>
> Do you wish to continue? [no]: **yes**

5. The script prompts you to configure the **slapi-nis** plug-in, a compatibility plug-in that allows older Linux clients to work with trusted users:

> Do you want to enable support for trusted domains in Schema Compatibility plugin?
> This will allow clients older than SSSD 1.9 and non-Linux clients to work with trusted users.
>
> Enable trusted domains support in slapi-nis? [no]: **yes**

6. When prompted, enter the NetBIOS name for the IdM domain or press **Enter** to accept the name suggested:

> Trust is configured but no NetBIOS domain name found, setting it now.
> Enter the NetBIOS name for the IPA domain.
> Only up to 15 uppercase ASCII letters, digits and dashes are allowed.
> Example: EXAMPLE.
>
> NetBIOS domain name [IDM]:

7. You are prompted to run the SID generation task to create a SID for any existing users:

> Do you want to run the ipa-sidgen task? [no]: **yes**

This is a resource-intensive task, so if you have a high number of users, you can run this at another time.

8. **(Optional)** By default, the Dynamic RPC port range is defined as **49152-65535** for Windows Server 2008 and later. If you need to define a different Dynamic RPC port range for your environment, configure Samba to use different ports and open those ports in your firewall settings. The following example sets the port range to **55000-65000**.

> [root@ipaserver ~]# **net conf setparm global 'rpc server dynamic port range' 55000-65000**
> [root@ipaserver ~]# **firewall-cmd --add-port=55000-65000/tcp**
> [root@ipaserver ~]# **firewall-cmd --runtime-to-permanent**

9. Restart the **ipa** service:

> [root@ipaserver ~]# **ipactl restart**

10. Use the **smbclient** utility to verify that Samba responds to Kerberos authentication from the IdM side:

```
[root@ipaserver ~]# smbclient -L server.idm.example.com -U user_name --use-
kerberos=required
lp_load_ex: changing to config backend registry
  Sharename     Type     Comment
  ---------     ----     -------
  IPC$          IPC      IPC Service (Samba 4.15.2)
...
```

3.6.2. Enabling the AES encryption type in Active Directory using a GPO

This section describes how to enable the AES encryption type in Active Directory (AD) using a group policy object (GPO). Certain features on RHEL, such as running a Samba server on an IdM client, require this encryption type.

Note that RHEL no longer supports the weak DES and RC4 encryption types.

Prerequisites

- You are logged into AD as a user who can edit group policies.

- The **Group Policy Management Console** is installed on the computer.

Procedure

1. Open the **Group Policy Management Console**.

2. Right-click **Default Domain Policy**, and select **Edit**. The **Group Policy Management Editor** opens.

3. Navigate to **Computer Configuration → Policies → Windows Settings → Security Settings → Local Policies → Security Options**.

4. Double-click the **Network security: Configure encryption types allowed for Kerberos** policy.

5. Select **AES256_HMAC_SHA1** and, optionally, **Future encryption types**.

6. Click **OK**.

7. Close the **Group Policy Management Editor**.

8. Repeat the steps for the **Default Domain Controller Policy**.

9. Wait until the Windows domain controllers (DC) applied the group policy automatically. Alternatively, to apply the GPO manually on a DC, enter the following command using an account that has administrator permissions:

   ```
   C:\> gpupdate /force /target:computer
   ```

3.6.3. Installing and configuring a Samba server on an IdM client

This section describes how to install and configure Samba on a client enrolled in an IdM domain.

Prerequisites

- Both the IdM servers and the client must run on RHEL 8.1 or later.

- The IdM domain is prepared as described in Preparing the IdM domain for installing Samba on domain members.

- If IdM has a trust configured with AD, enable the AES encryption type for Kerberos. For example, use a group policy object (GPO) to enable the AES encryption type. For details, see Enabling AES encryption in Active Directory using a GPO .

Procedure

1. Install the **ipa-client-samba** package:

   [root@idm_client]# **yum install ipa-client-samba**

2. Use the **ipa-client-samba** utility to prepare the client and create an initial Samba configuration:

   [root@idm_client]# **ipa-client-samba**
   Searching for IPA server...
   IPA server: DNS discovery
   Chosen IPA master: *idm_server.idm.example.com*
   SMB principal to be created: cifs/*idm_client.idm.example.com@IDM.EXAMPLE.COM*
   NetBIOS name to be used: *IDM_CLIENT*
   Discovered domains to use:

    Domain name: *idm.example.com*
   NetBIOS name: *IDM*
        SID: *S-1-5-21-525930803-952335037-206501584*
     ID range: *212000000 - 212199999*

    Domain name: *ad.example.com*
   NetBIOS name: *AD*
        SID: None
     ID range: *1918400000 - 1918599999*

   Continue to configure the system with these values? [no]: **yes**
   Samba domain member is configured. Please check configuration at /etc/samba/smb.conf
   and start smb and winbind services

3. By default, **ipa-client-samba** automatically adds the **[homes]** section to the /**etc/samba/smb.conf** file that dynamically shares a user's home directory when the user connects. If users do not have home directories on this server, or if you do not want to share them, remove the following lines from /**etc/samba/smb.conf**:

   [homes]
       read only = no

4. Share directories and printers. For details, see:

   - Setting up a Samba file share that uses POSIX ACLs

   - Setting up a share that uses Windows ACLs

   - Setting up Samba as a print server

5. Open the ports required for a Samba client in the local firewall:

> [root@idm_client]# **firewall-cmd --permanent --add-service=samba-client**
> [root@idm_client]# **firewall-cmd --reload**

6. Enable and start the **smb** and **winbind** services:

> [root@idm_client]# **systemctl enable --now smb winbind**

Verification steps

Run the following verification step on a different IdM domain member that has the **samba-client** package installed:

- List the shares on the Samba server using Kerberos authentication:

> $ **smbclient -L _idm_client.idm.example.com_ -U _user_name_ --use-kerberos=required**
> lp_load_ex: changing to config backend registry
>
> Sharename   Type   Comment
> ---------   ----   -------
> _example_   Disk
> IPC$ IPC   IPC Service (Samba 4.15.2)
> ...

Additional resources

- **ipa-client-samba(1)** man page

3.6.4. Manually adding an ID mapping configuration if IdM trusts a new domain

Samba requires an ID mapping configuration for each domain from which users access resources. On an existing Samba server running on an IdM client, you must manually add an ID mapping configuration after the administrator added a new trust to an Active Directory (AD) domain.

Prerequisites

- You configured Samba on an IdM client. Afterward, a new trust was added to IdM.

- The DES and RC4 encryption types for Kerberos must be disabled in the trusted AD domain. For security reasons, RHEL 8 does not support these weak encryption types.

Procedure

1. Authenticate using the host's keytab:

> [root@idm_client]# **kinit -k**

2. Use the **ipa idrange-find** command to display both the base ID and the ID range size of the new domain. For example, the following command displays the values for the **ad.example.com** domain:

> [root@idm_client]# **ipa idrange-find --name="_AD.EXAMPLE.COM__id_range" --raw**
> ---------------

```
1 range matched
---------------
  cn: AD.EXAMPLE.COM_id_range
  ipabaseid: 1918400000
  ipaidrangesize: 200000
  ipabaserid: 0
  ipanttrusteddomainsid: S-1-5-21-968346183-862388825-1738313271
  iparangetype: ipa-ad-trust
----------------------------
Number of entries returned 1
----------------------------
```

You need the values from the **ipabaseid** and **ipaidrangesize** attributes in the next steps.

3. To calculate the highest usable ID, use the following formula:

   maximum_range = ipabaseid + ipaidrangesize - 1

   With the values from the previous step, the highest usable ID for the **ad.example.com** domain is **1918599999** (1918400000 + 200000 – 1).

4. Edit the **/etc/samba/smb.conf** file, and add the ID mapping configuration for the domain to the **[global]** section:

   idmap config *AD* : range = *1918400000 - 1918599999*
   idmap config *AD* : backend = sss

   Specify the value from **ipabaseid** attribute as the lowest and the computed value from the previous step as the highest value of the range.

5. Restart the **smb** and **winbind** services:

   [root@idm_client]# **systemctl restart smb winbind**

Verification steps

- List the shares on the Samba server using Kerberos authentication:

  $ **smbclient -L** *idm_client.idm.example.com* **-U** *user_name* **--use-kerberos=required**
  lp_load_ex: changing to config backend registry

  ```
    Sharename     Type     Comment
    ---------     ----     -------
    example       Disk
    IPC$          IPC      IPC Service (Samba 4.15.2)
    ...
  ```

3.6.5. Additional resources

- [Installing an Identity Management client](#)

3.7. SETTING UP A SAMBA FILE SHARE THAT USES POSIX ACLS

As a Linux service, Samba supports shares with POSIX ACLs. They enable you to manage permissions locally on the Samba server using utilities, such as **chmod**. If the share is stored on a file system that supports extended attributes, you can define ACLs with multiple users and groups.

> **NOTE**
>
> If you need to use fine-granular Windows ACLs instead, see Setting up a share that uses Windows ACLs.

Parts of this section were adopted from the Setting up a Share Using POSIX ACLs documentation published in the Samba Wiki. License: CC BY 4.0. Authors and contributors: See the history tab on the Wiki page.

3.7.1. Adding a share that uses POSIX ACLs

This section describes how to create a share named **example** that provides the content of the **/srv/samba/example/** directory, and uses POSIX ACLs.

Prerequisites

Samba has been set up in one of the following modes:

- Standalone server

- Domain member

Procedure

1. Create the folder if it does not exist. For example:

   ```
   # mkdir -p /srv/samba/example/
   ```

2. If you run SELinux in **enforcing** mode, set the **samba_share_t** context on the directory:

   ```
   # semanage fcontext -a -t samba_share_t "/srv/samba/example(/.*)?"
   # restorecon -Rv /srv/samba/example/
   ```

3. Set file system ACLs on the directory. For details, see:

   - Setting standard ACLs on a Samba Share that uses POSIX ACLs

   - Setting extended ACLs on a share that uses POSIX ACLs .

4. Add the example share to the **/etc/samba/smb.conf** file. For example, to add the share write-enabled:

   ```
   [example]
   path = /srv/samba/example/
   read only = no
   ```
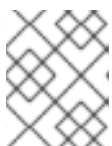
   > **NOTE**
   >
   > Regardless of the file system ACLs; if you do not set **read only = no**, Samba shares the directory in read-only mode.

5. Verify the **/etc/samba/smb.conf** file:

```
# testparm
```

6. Open the required ports and reload the firewall configuration using the **firewall-cmd** utility:

```
# firewall-cmd --permanent --add-service=samba
# firewall-cmd --reload
```

7. Restart the **smb** service:

```
# systemctl restart smb
```

3.7.2. Setting standard Linux ACLs on a Samba share that uses POSIX ACLs

The standard ACLs on Linux support setting permissions for one owner, one group, and for all other undefined users. You can use the **chown**, **chgrp**, and **chmod** utility to update the ACLs. If you require precise control, then you use the more complex POSIX ACLs, see

Setting extended ACLs on a Samba share that uses POSIX ACLs .

The following procedure sets the owner of the **/srv/samba/example/** directory to the **root** user, grants read and write permissions to the **Domain Users** group, and denies access to all other users.

Prerequisites

- The Samba share on which you want to set the ACLs exists.

Procedure

```
# chown root:"Domain Users" /srv/samba/example/
# chmod 2770 /srv/samba/example/
```

> **NOTE**
>
> Enabling the set-group-ID (SGID) bit on a directory automatically sets the default group for all new files and subdirectories to that of the directory group, instead of the usual behavior of setting it to the primary group of the user who created the new directory entry.

Additional resources

- **chown(1)** and **chmod(1)** man pages

3.7.3. Setting extended ACLs on a Samba share that uses POSIX ACLs

If the file system the shared directory is stored on supports extended ACLs, you can use them to set complex permissions. Extended ACLs can contain permissions for multiple users and groups.

Extended POSIX ACLs enable you to configure complex ACLs with multiple users and groups. However, you can only set the following permissions:

- No access

- Read access

- Write access

- Full control

If you require the fine-granular Windows permissions, such as **Create folder** / **append data**, configure the share to use Windows ACLs. See Setting up a share that uses Windows ACLs .

The following procedure shows how to enable extended ACLs on a share. Additionally, it contains an example about setting extended ACLs.

Prerequisites

- The Samba share on which you want to set the ACLs exists.

Procedure

1. Enable the following parameter in the share's section in the **/etc/samba/smb.conf** file to enable ACL inheritance of extended ACLs:

   inherit acls = yes

   For details, see the parameter description in the **smb.conf(5**) man page.

2. Restart the **smb** service:

   # **systemctl restart smb**

3. Set the ACLs on the directory. For example:

   Example 3.2. Setting Extended ACLs

   The following procedure sets read, write, and execute permissions for the **Domain Admins** group, read, and execute permissions for the **Domain Users** group, and deny access to everyone else on the **/srv/samba/example/** directory:

   1. Disable auto-granting permissions to the primary group of user accounts:

      # **setfacl -m group::--- /srv/samba/example/**
      # **setfacl -m default:group::--- /srv/samba/example/**

      The primary group of the directory is additionally mapped to the dynamic **CREATOR GROUP** principal. When you use extended POSIX ACLs on a Samba share, this principal is automatically added and you cannot remove it.

   2. Set the permissions on the directory:

      a. Grant read, write, and execute permissions to the **Domain Admins** group:

         # **setfacl -m group:"*DOMAIN*\Domain Admins":rwx /srv/samba/example/**

      b. Grant read and execute permissions to the **Domain Users** group:

```
# setfacl -m group:"DOMAIN\Domain Users":r-x /srv/samba/example/
```

c. Set permissions for the **other** ACL entry to deny access to users that do not match the other ACL entries:

```
# setfacl -R -m other::--- /srv/samba/example/
```

These settings apply only to this directory. In Windows, these ACLs are mapped to the **This folder only** mode.

3. To enable the permissions set in the previous step to be inherited by new file system objects created in this directory:

```
# setfacl -m default:group:"DOMAIN\Domain Admins":rwx /srv/samba/example/
# setfacl -m default:group:"DOMAIN\Domain Users":r-x /srv/samba/example/
# setfacl -m default:other::--- /srv/samba/example/
```

With these settings, the **This folder only** mode for the principals is now set to **This folder, subfolders, and files**.

Samba maps the permissions set in the procedure to the following Windows ACLs:

Principal	Access	Applies to
Domain\Domain Admins	Full control	This folder, subfolders, and files
Domain\Domain Users	Read & execute	This folder, subfolders, and files
Everyone [a]	None	This folder, subfolders, and files
owner (*Unix User*\owner) [b]	Full control	This folder only
primary_group (*Unix User*\primary_group) [c]	None	This folder only
CREATOR OWNER [d] [e]	Full control	Subfolders and files only
CREATOR GROUP [e] [f]	None	Subfolders and files only

[a] Samba maps the permissions for this principal from the **other** ACL entry.

[b] Samba maps the owner of the directory to this entry.

[c] Samba maps the primary group of the directory to this entry.

[d] On new file system objects, the creator inherits automatically the permissions of this principal.

[e] Configuring or removing these principals from the ACLs not supported on shares that use POSIX ACLs.

[f] On new file system objects, the creator's primary group inherits automatically the permissions of this principal.

3.8. SETTING PERMISSIONS ON A SHARE THAT USES POSIX ACLS

Optionally, to limit or grant access to a Samba share, you can set certain parameters in the share's section in the **/etc/samba/smb.conf** file.

> **NOTE**
>
> Share-based permissions manage if a user, group, or host is able to access a share. These settings do not affect file system ACLs.

Use share-based settings to restrict access to shares, for example, to deny access from specific hosts.

Prerequisites

- A share with POSIX ACLs has been set up.

3.8.1. Configuring user and group-based share access

User and group-based access control enables you to grant or deny access to a share for certain users and groups.

Prerequisites

- The Samba share on which you want to set user or group-based access exists.

Procedure

1. For example, to enable all members of the **Domain Users** group to access a share while access is denied for the **user** account, add the following parameters to the share's configuration:

   ```
   valid users = +DOMAIN\"Domain Users"
   invalid users = DOMAIN\user
   ```

   The **invalid users** parameter has a higher priority than the **valid users** parameter. For example, if the **user** account is a member of the **Domain Users** group, access is denied to this account when you use the previous example.

2. Reload the Samba configuration:

   ```
   # smbcontrol all reload-config
   ```

Additional resources

- **smb.conf(5)** man page

3.8.2. Configuring host-based share access

Host-based access control enables you to grant or deny access to a share based on client's host names, IP addresses, or IP range.

The following procedure explains how to enable the **127.0.0.1** IP address, the **192.0.2.0/24** IP range, and the **client1.example.com** host to access a share, and additionally deny access for the **client2.example.com** host:

Prerequisites

- The Samba share on which you want to set host-based access exists.

Procedure

1. Add the following parameters to the configuration of the share in the **/etc/samba/smb.conf** file:

   ```
   hosts allow = 127.0.0.1 192.0.2.0/24 client1.example.com
   hosts deny = client2.example.com
   ```

   The **hosts deny** parameter has a higher priority than **hosts allow**. For example, if **client1.example.com** resolves to an IP address that is listed in the **hosts allow** parameter, access for this host is denied.

2. Reload the Samba configuration:

   ```
   # smbcontrol all reload-config
   ```

Additional resources

- **smb.conf(5)** man page

3.9. SETTING UP A SHARE THAT USES WINDOWS ACLS

Samba supports setting Windows ACLs on shares and file system object. This enables you to:

- Use the fine-granular Windows ACLs

- Manage share permissions and file system ACLs using Windows

Alternatively, you can configure a share to use POSIX ACLs. For details, see Setting up a Samba file share that uses POSIX ACLs.

Parts of this section were adopted from the Setting up a Share Using Windows ACLs documentation published in the Samba Wiki. License: CC BY 4.0. Authors and contributors: See the history tab on the Wiki page.

3.9.1. Granting the SeDiskOperatorPrivilege privilege

Only users and groups having the **SeDiskOperatorPrivilege** privilege granted can configure permissions on shares that use Windows ACLs.

Procedure

1. For example, to grant the **SeDiskOperatorPrivilege** privilege to the *DOMAIN*\Domain Admins group:

   ```
   # net rpc rights grant "DOMAIN\Domain Admins" SeDiskOperatorPrivilege -U
   "DOMAIN\administrator"
   ```

Enter *DOMAIN*\administrator's password:
Successfully granted rights.

> **NOTE**
>
> In a domain environment, grant **SeDiskOperatorPrivilege** to a domain group.
> This enables you to centrally manage the privilege by updating a user's group
> membership.

2. To list all users and groups having **SeDiskOperatorPrivilege** granted:

   ```
   # net rpc rights list privileges SeDiskOperatorPrivilege -U "DOMAIN\administrator"
   Enter administrator's password:
   SeDiskOperatorPrivilege:
     BUILTIN\Administrators
     DOMAIN\Domain Admins
   ```

3.9.2. Enabling Windows ACL support

To configure shares that support Windows ACLs, you must enable this feature in Samba.

Prerequisites

- A user share is configured on the Samba server.

Procedure

1. To enable it globally for all shares, add the following settings to the **[global]** section of the
   /etc/samba/smb.conf file:

   ```
   vfs objects = acl_xattr
   map acl inherit = yes
   store dos attributes = yes
   ```

   Alternatively, you can enable Windows ACL support for individual shares, by adding the same
   parameters to a share's section instead.

2. Restart the **smb** service:

   ```
   # systemctl restart smb
   ```

3.9.3. Adding a share that uses Windows ACLs

This section describes how to create a share named **example**, that shares the content of the
/srv/samba/example/ directory, and uses Windows ACLs.

Procedure

1. Create the folder if it does not exists. For example:

   ```
   # mkdir -p /srv/samba/example/
   ```

2. If you run SELinux in **enforcing** mode, set the **samba_share_t** context on the directory:

```
# semanage fcontext -a -t samba_share_t "/srv/samba/example(/.*)?"
# restorecon -Rv /srv/samba/example/
```

3. Add the example share to the **/etc/samba/smb.conf** file. For example, to add the share write-enabled:

```
[example]
 path = /srv/samba/example/
 read only = no
```

> **NOTE**
>
> Regardless of the file system ACLs; if you do not set **read only = no**, Samba shares the directory in read-only mode.

4. If you have not enabled Windows ACL support in the **[global]** section for all shares, add the following parameters to the **[example]** section to enable this feature for this share:

```
vfs objects = acl_xattr
map acl inherit = yes
store dos attributes = yes
```

5. Verify the **/etc/samba/smb.conf** file:

```
# testparm
```

6. Open the required ports and reload the firewall configuration using the **firewall-cmd** utility:

```
# firewall-cmd --permanent --add-service=samba
# firewall-cmd --reload
```

7. Restart the **smb** service:

```
# systemctl restart smb
```

3.9.4. Managing share permissions and file system ACLs of a share that uses Windows ACLs

To manage share permissions and file system ACLs on a Samba share that uses Windows ACLs, use a Windows applications, such as **Computer Management**. For details, see the Windows documentation. Alternatively, use the **smbcacls** utility to manage ACLs.

> **NOTE**
>
> To modify the file system permissions from Windows, you must use an account that has the **SeDiskOperatorPrivilege** privilege granted.

Additional resources

- [Section 3.10, "Managing ACLs on an SMB share using smbcacls"](#)

- [Section 3.9.1, "Granting the SeDiskOperatorPrivilege privilege"](#)

3.10. MANAGING ACLS ON AN SMB SHARE USING SMBCACLS

The **smbcacls** utility can list, set, and delete ACLs of files and directories stored on an SMB share. You can use **smbcacls** to manage file system ACLs:

- On a local or remote Samba server that uses advanced Windows ACLs or POSIX ACLs

- On Red Hat Enterprise Linux to remotely manage ACLs on a share hosted on Windows

3.10.1. Access control entries

Each ACL entry of a file system object contains Access Control Entries (ACE) in the following format:

> *security_principal*:*access_right*/*inheritance_information*/*permissions*

> **Example 3.3. Access control entries**
>
> If the **AD\Domain Users** group has **Modify** permissions that apply to **This folder, subfolders, and files** on Windows, the ACL contains the following ACE:
>
> > AD\Domain Users:ALLOWED/OI|CI/CHANGE

An ACE contains the following parts:

Security principal

The security principal is the user, group, or SID the permissions in the ACL are applied to.

Access right

Defines if access to an object is granted or denied. The value can be **ALLOWED** or **DENIED**.

Inheritance information

The following values exist:

Table 3.1. Inheritance settings

Value	Description	Maps to
OI	Object Inherit	This folder and files
CI	Container Inherit	This folder and subfolders
IO	Inherit Only	The ACE does not apply to the current file or directory
ID	Inherited	The ACE was inherited from the parent directory

Additionally, the values can be combined as follows:

Table 3.2. Inheritance settings combinations

Value combinations	Maps to the Windows **Applies to** setting
OI|CI	This folder, subfolders, and files
OI|CI|IO	Subfolders and files only
CI|IO	Subfolders only
OI|IO	Files only

Permissions

This value can be either a hex value that represents one or more Windows permissions or an **smbcacls** alias:

- A hex value that represents one or more Windows permissions.
  The following table displays the advanced Windows permissions and their corresponding value in hex format:

Table 3.3. Windows permissions and their corresponding smbcacls value in hex format

Windows permissions	Hex values
Full control	**0x001F01FF**
Traverse folder / execute file	**0x00100020**
List folder / read data	**0x00100001**
Read attributes	**0x00100080**
Read extended attributes	**0x00100008**
Create files / write data	**0x00100002**
Create folders / append data	**0x00100004**
Write attributes	**0x00100100**
Write extended attributes	**0x00100010**
Delete subfolders and files	**0x00100040**
Delete	**0x00110000**
Read permissions	**0x00120000**
Change permissions	**0x00140000**

Windows permissions	Hex values
Take ownership	**0x00180000**

Multiple permissions can be combined as a single hex value using the bit-wise **OR** operation. For details, see ACE mask calculation.

- An **smbcacls** alias. The following table displays the available aliases:

Table 3.4. Existing smbcacls aliases and their corresponding Windows permission

smbcacls alias	Maps to Windows permission
R	Read
READ	Read & execute
W	Special: ○ Create files / write data ○ Create folders / append data ○ Write attributes ○ Write extended attributes ○ Read permissions
D	Delete
P	Change permissions
O	Take ownership
X	Traverse / execute
CHANGE	Modify
FULL	Full control

NOTE

You can combine single-letter aliases when you set permissions. For example, you can set **RD** to apply the Windows permission **Read** and **Delete**. However, you can neither combine multiple non-single-letter aliases nor combine aliases and hex values.

3.10.2. Displaying ACLs using smbcacls

To display ACLs on an SMB share, use the **smbcacls** utility. If you run **smbcacls** without any operation parameter, such as **--add**, the utility displays the ACLs of a file system object.

Procedure

For example, to list the ACLs of the root directory of the //**server**/**example** share:

```
# smbcacls //server/example / -U "DOMAIN\administrator"
Enter DOMAIN\administrator's password:
REVISION:1
CONTROL:SR|PD|DI|DP
OWNER:AD\Administrators
GROUP:AD\Domain Users
ACL:AD\Administrator:ALLOWED/OI|CI/FULL
ACL:AD\Domain Users:ALLOWED/OI|CI/CHANGE
ACL:AD\Domain Guests:ALLOWED/OI|CI/0x00100021
```

The output of the command displays:

- **REVISION**: The internal Windows NT ACL revision of the security descriptor

- **CONTROL**: Security descriptor control

- **OWNER**: Name or SID of the security descriptor's owner

- **GROUP**: Name or SID of the security descriptor's group

- **ACL** entries. For details, see Access control entries .

3.10.3. ACE mask calculation

In most situations, when you add or update an ACE, you use the **smbcacls** aliases listed in Existing smbcacls aliases and their corresponding Windows permission.

However, if you want to set advanced Windows permissions as listed in Windows permissions and their corresponding smbcacls value in hex format, you must use the bit-wise **OR** operation to calculate the correct value. You can use the following shell command to calculate the value:

```
# echo $(printf '0x%X' $(( hex_value_1 | hex_value_2 | ... )))
```

Example 3.4. Calculating an ACE Mask

You want to set the following permissions:

- Traverse folder / execute file (0x00100020)

- List folder / read data (0x00100001)

- Read attributes (0x00100080)

To calculate the hex value for the previous permissions, enter:

```
# echo $(printf '0x%X' $(( 0x00100020 | 0x00100001 | 0x00100080 )))
0x1000A1
```

Use the returned value when you set or update an ACE.

3.10.4. Adding, updating, and removing an ACL using smbcacls

Depending on the parameter you pass to the **smbcacls** utility, you can add, update, and remove ACLs from a file or directory.

Adding an ACL
To add an ACL to the root of the //**server**/**example** share that grants **CHANGE** permissions for **This folder, subfolders, and files** to the **AD\Domain Users** group:

> # **smbcacls //server/example / -U "*DOMAIN*\administrator --add ACL:"AD\Domain Users":ALLOWED/OI|CI/CHANGE**

Updating an ACL
Updating an ACL is similar to adding a new ACL. You update an ACL by overriding the ACL using the **--modify** parameter with an existing security principal. If **smbcacls** finds the security principal in the ACL list, the utility updates the permissions. Otherwise the command fails with an error:

> ACL for SID *principal_name* not found

For example, to update the permissions of the **AD\Domain Users** group and set them to **READ** for **This folder, subfolders, and files**:

> # **smbcacls //server/example / -U "*DOMAIN*\administrator --modify ACL:"AD\Domain Users":ALLOWED/OI|CI/READ**

Deleting an ACL
To delete an ACL, pass the **--delete** parameter with the exact ACL to the **smbcacls** utility. For example:

> # **smbcacls //server/example / -U "*DOMAIN*\administrator --delete ACL:"AD\Domain Users":ALLOWED/OI|CI/READ**

3.11. ENABLING USERS TO SHARE DIRECTORIES ON A SAMBA SERVER

On a Samba server, you can configure that users can share directories without root permissions.

3.11.1. Enabling the user shares feature

Before users can share directories, the administrator must enable user shares in Samba.

For example, to enable only members of the local **example** group to create user shares.

Procedure

1. Create the local **example** group, if it does not exist:

   > # **groupadd example**

2. Prepare the directory for Samba to store the user share definitions and set its permissions properly. For example:

   a. Create the directory:

      ```
      # mkdir -p /var/lib/samba/usershares/
      ```

   b. Set write permissions for the **example** group:

      ```
      # chgrp example /var/lib/samba/usershares/
      # chmod 1770 /var/lib/samba/usershares/
      ```

   c. Set the sticky bit to prevent users to rename or delete files stored by other users in this directory.

3. Edit the **/etc/samba/smb.conf** file and add the following to the **[global]** section:

   a. Set the path to the directory you configured to store the user share definitions. For example:

      ```
      usershare path = /var/lib/samba/usershares/
      ```

   b. Set how many user shares Samba allows to be created on this server. For example:

      ```
      usershare max shares = 100
      ```

      If you use the default of **0** for the **usershare max shares** parameter, user shares are disabled.

   c. Optionally, set a list of absolute directory paths. For example, to configure that Samba only allows to share subdirectories of the **/data** and **/srv** directory to be shared, set:

      ```
      usershare prefix allow list = /data /srv
      ```

   For a list of further user share-related parameters you can set, see the **USERSHARES** section in the **smb.conf(5)** man page.

4. Verify the **/etc/samba/smb.conf** file:

   ```
   # testparm
   ```

5. Reload the Samba configuration:

   ```
   # smbcontrol all reload-config
   ```

   Users are now able to create user shares.

3.11.2. Adding a user share

After you enabled the user share feature in Samba, users can share directories on the Samba server without **root** permissions by running the **net usershare add** command.

Synopsis of the **net usershare add** command:

net usershare add share_name path [[comment] | [ACLs]] [guest_ok=y|n]
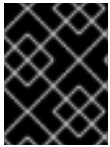
> **IMPORTANT**
>
> If you set ACLs when you create a user share, you must specify the comment parameter prior to the ACLs. To set an empty comment, use an empty string in double quotes.

Note that users can only enable guest access on a user share, if the administrator set **usershare allow guests = yes** in the **[global]** section in the **/etc/samba/smb.conf** file.

> Example 3.5. Adding a user share
>
> A user wants to share the **/srv/samba/** directory on a Samba server. The share should be named **example**, have no comment set, and should be accessible by guest users. Additionally, the share permissions should be set to full access for the **AD\Domain Users** group and read permissions for other users. To add this share, run as the user:
>
> $ **net usershare add example /srv/samba/ "" "AD\Domain Users":F,Everyone:R**
> **guest_ok=yes**

3.11.3. Updating settings of a user share

To update settings of a user share, override the share by using the **net usershare add** command with the same share name and the new settings. See Adding a user share .

3.11.4. Displaying information about existing user shares

Users can enter the **net usershare info** command on a Samba server to display user shares and their settings.

Prerequisites

- A user share is configured on the Samba server.

Procedure

1. To display all user shares created by any user:

   $ **net usershare info -l**
   [*share_1*]
   path=*/srv/samba/*
   comment=
   usershare_acl=*Everyone:R,host_name\user:F,*
   guest_ok=*y*
   ...

   To list only shares created by the user who runs the command, omit the **-l** parameter.

2. To display only the information about specific shares, pass the share name or wild cards to the command. For example, to display the information about shares whose name starts with **share_**:

   $ **net usershare info -l share_***

3.11.5. Listing user shares

If you want to list only the available user shares without their settings on a Samba server, use the **net usershare list** command.

Prerequisites

- A user share is configured on the Samba server.

Procedure

1. To list the shares created by any user:

   ```
   $ net usershare list -l
   share_1
   share_2
   ...
   ```

   To list only shares created by the user who runs the command, omit the **-l** parameter.

2. To list only specific shares, pass the share name or wild cards to the command. For example, to list only shares whose name starts with **share_**:

   ```
   $ net usershare list -l share_*
   ```

3.11.6. Deleting a user share

To delete a user share, use the command **net usershare delete** command as the user who created the share or as the **root** user.

Prerequisites

- A user share is configured on the Samba server.

Procedure

```
$ net usershare delete share_name
```

3.12. CONFIGURING A SHARE TO ALLOW ACCESS WITHOUT AUTHENTICATION

In certain situations, you want to share a directory to which users can connect without authentication. To configure this, enable guest access on a share.

> **WARNING**
>
> Shares that do not require authentication can be a security risk.

3.12.1. Enabling guest access to a share

If guest access is enabled on a share, Samba maps guest connections to the operating system account set in the **guest account** parameter. Guest users can access files on this share if at least one of the following conditions is satisfied:

- The account is listed in file system ACLs

- The POSIX permissions for **other** users allow it

Example 3.6. Guest share permissions

If you configured Samba to map the guest account to **nobody**, which is the default, the ACLs in the following example:

- Allow guest users to read **file1.txt**

- Allow guest users to read and modify **file2.txt**

- Prevent guest users to read or modify **file3.txt**

```
-rw-r--r--. 1 root      root     1024 1. Sep 10:00 file1.txt
-rw-r-----. 1 nobody    root     1024 1. Sep 10:00 file2.txt
-rw-r-----. 1 root      root     1024 1. Sep 10:00 file3.txt
```

Procedure

1. Edit the **/etc/samba/smb.conf** file:

   a. If this is the first guest share you set up on this server:

      i. Set **map to guest = Bad User** in the **[global]** section:

         ```
         [global]
                 ...
                 map to guest = Bad User
         ```

         With this setting, Samba rejects login attempts that use an incorrect password unless the user name does not exist. If the specified user name does not exist and guest access is enabled on a share, Samba treats the connection as a guest log in.

      ii. By default, Samba maps the guest account to the **nobody** account on Red Hat Enterprise Linux. Alternatively, you can set a different account. For example:

         ```
         [global]
                 ...
                 guest account = user_name
         ```

         The account set in this parameter must exist locally on the Samba server. For security reasons, Red Hat recommends using an account that does not have a valid shell assigned.

   b. Add the **guest ok = yes** setting to the **[example]** share section:

```
[example]
    ...
    guest ok = yes
```

2. Verify the **/etc/samba/smb.conf** file:

   ```
   # testparm
   ```

3. Reload the Samba configuration:

   ```
   # smbcontrol all reload-config
   ```

3.13. CONFIGURING SAMBA FOR MACOS CLIENTS

The **fruit** virtual file system (VFS) Samba module provides enhanced compatibility with Apple server message block (SMB) clients.

3.13.1. Optimizing the Samba configuration for providing file shares for macOS clients

This section describes how to configure the **fruit** module for all Samba shares hosted on the server to optimize Samba file shares for macOS clients.
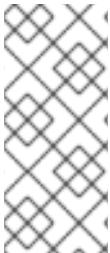
> **NOTE**
>
> Red Hat recommends enabling the **fruit** module globally. Clients using macOS negotiate the server server message block version 2 (SMB2) Apple (AAPL) protocol extensions when the client establishes the first connection to the server. If the client first connects to a share without AAPL extensions enabled, the client does not use the extensions for any share of the server.

Prerequisites

- Samba is configured as a file server.

Procedure

1. Edit the **/etc/samba/smb.conf** file, and enable the **fruit** and **streams_xattr** VFS modules in the **[global]** section:

   ```
   vfs objects = fruit streams_xattr
   ```

   > **IMPORTANT**
   >
   > You must enable the **fruit** module before enabling **streams_xattr**. The **fruit** module uses alternate data streams (ADS). For this reason, you must also enable the **streams_xattr** module.

2. Optionally, to provide macOS Time Machine support on a share, add the following setting to the share configuration in the **/etc/samba/smb.conf** file:

> fruit:time machine = yes

3. Verify the **/etc/samba/smb.conf** file:

   > # **testparm**

4. Reload the Samba configuration:

   > # **smbcontrol all reload-config**

Additional resources

- **vfs_fruit(8)** man page.

- Configuring file shares:

  - Setting up a Samba file share that uses POSIX ACLs

  - Setting up a share that uses Windows ACLs .

3.14. USING THE SMBCLIENT UTILITY TO ACCESS AN SMB SHARE

The smbclient utility enables you to access file shares on an SMB server, similarly to a command-line FTP client. You can use it, for example, to upload and download files to and from a share.

Prerequisites

- The **samba-client** package is installed.

3.14.1. How the smbclient interactive mode works

For example, to authenticate to the **example** share hosted on **server** using the *DOMAIN\user* account:

> # **smbclient -U "*DOMAIN\user*" //*server*/*example***
> Enter *domain\user*'s password:
> Try "help" to get a list of possible commands.
> smb: \>

After **smbclient** connected successfully to the share, the utility enters the interactive mode and shows the following prompt:

> smb: \>

To display all available commands in the interactive shell, enter:

> smb: \> **help**

To display the help for a specific command, enter:

> smb: \> **help** *command_name*

Additional resources

- **smbclient(1)** man page

3.14.2. Using smbclient in interactive mode

If you use **smbclient** without the **-c** parameter, the utility enters the interactive mode. The following procedure shows how to connect to an SMB share and download a file from a subdirectory.

Procedure

1. Connect to the share:

   ```
   # smbclient -U "DOMAIN\user_name" //server_name/share_name
   ```

2. Change into the **/example/** directory:

   ```
   smb: \> d /example/
   ```

3. List the files in the directory:

   ```
   smb: \example\> ls
     .              D       0  Thu Nov 1 10:00:00 2018
     ..             D       0  Thu Nov 1 10:00:00 2018
     example.txt    N  1048576  Thu Nov 1 10:00:00 2018

         9950208 blocks of size 1024. 8247144 blocks available
   ```

4. Download the **example.txt** file:

   ```
   smb: \example\> get example.txt
   getting file \directory\subdirectory\example.txt of size 1048576 as example.txt (511975,0
   KiloBytes/sec) (average 170666,7 KiloBytes/sec)
   ```

5. Disconnect from the share:

   ```
   smb: \example\> exit
   ```

3.14.3. Using smbclient in scripting mode

If you pass the **-c** parameter to **smbclient**, you can automatically execute the commands on the remote SMB share. This enables you to use **smbclient** in scripts.

The following procedure shows how to connect to an SMB share and download a file from a subdirectory.

Procedure

- Use the following command to connect to the share, change into the **example** directory, download the **example.txt** file:

```
# smbclient -U DOMAIN\user_name //server_name/share_name -c "cd /example/ ; get
example.txt ; exit"
```

3.15. SETTING UP SAMBA AS A PRINT SERVER

If you set up Samba as a print server, clients in your network can use Samba to print. Additionally, Windows clients can, if configured, download the driver from the Samba server.

Parts of this section were adopted from the Setting up Samba as a Print Server documentation published in the Samba Wiki. License: CC BY 4.0. Authors and contributors: See the history tab on the Wiki page.

Prerequisites

Samba has been set up in one of the following modes:

- Standalone server

- Domain member

3.15.1. Enabling print server support in Samba

By default, print server support is not enabled in Samba. To use Samba as a print server, you must configure Samba accordingly.

> **NOTE**
>
> Print jobs and printer operations require remote procedure calls (RPCs). By default, Samba starts the **rpcd_spoolss** service on demand to manage RPCs. During the first RPC call, or when you update the printer list in CUPS, Samba retrieves the printer information from CUPS. This can require approximately 1 second per printer. Therefore, if you have more than 50 printers, tune the **rpcd_spoolss** settings.

Prerequisites

- The printers are configured in a CUPS server.
  For details about configuring printers in CUPS, see the documentation provided in the CUPS web console (https://*printserver*:631/help) on the print server.

Procedure

1. Edit the **/etc/samba/smb.conf** file:

   a. Add the **[printers]** section to enable the printing backend in Samba:

      ```
      [printers]
          comment = All Printers
          path = /var/tmp/
          printable = yes
          create mask = 0600
      ```

      > **IMPORTANT**
      >
      > The **[printers]** share name is hard-coded and cannot be changed.

b. If the CUPS server runs on a different host or port, specify the setting in the **[printers]** section:

> **cups server = *printserver.example.com*:*631***

c. If you have many printers, set the number of idle seconds to a higher value than the numbers of printers connected to CUPS. For example, if you have 100 printers, set in the **[global]** section:

> **rpcd_spoolss:idle_seconds = 200**

If this setting does not scale in your environment, also increase the number of **rpcd_spoolss** workers in the **[global]** section:

> **rpcd_spoolss:num_workers = 10**

By default, **rpcd_spoolss** starts 5 workers.

2. Verify the **/etc/samba/smb.conf** file:

> **# testparm**

3. Open the required ports and reload the firewall configuration using the **firewall-cmd** utility:

> **# firewall-cmd --permanent --add-service=samba**
> **# firewall-cmd --reload**

4. Restart the **smb** service:

> **# systemctl restart smb**

After restarting the service, Samba automatically shares all printers that are configured in the CUPS back end. If you want to manually share only specific printers, see Manually sharing specific printers.

Verification

- Submit a print job. For example, to print a PDF file, enter:

> **# smbclient -U*user* //*sambaserver.example.com*/*printer_name* -c "print *example.pdf*"**

3.15.2. Manually sharing specific printers

If you configured Samba as a print server, by default, Samba shares all printers that are configured in the CUPS back end. The following procedure explains how to share only specific printers.

Prerequisites

- Samba is set up as a print server

Procedure

1. Edit the **/etc/samba/smb.conf** file:

a. In the **[global]** section, disable automatic printer sharing by setting:

```
load printers = no
```

b. Add a section for each printer you want to share. For example, to share the printer named **example** in the CUPS back end as **Example-Printer** in Samba, add the following section:

```
[Example-Printer]
        path = /var/tmp/
        printable = yes
        printer name = example
```

You do not need individual spool directories for each printer. You can set the same spool directory in the **path** parameter for the printer as you set in the **[printers]** section.

2. Verify the **/etc/samba/smb.conf** file:

```
# testparm
```

3. Reload the Samba configuration:

```
# smbcontrol all reload-config
```

3.16. SETTING UP AUTOMATIC PRINTER DRIVER DOWNLOADS FOR WINDOWS CLIENTS ON SAMBA PRINT SERVERS

If you are running a Samba print server for Windows clients, you can upload drivers and preconfigure printers. If a user connects to a printer, Windows automatically downloads and installs the driver locally on the client. The user does not require local administrator permissions for the installation. Additionally, Windows applies preconfigured driver settings, such as the number of trays.

Parts of this section were adopted from the Setting up Automatic Printer Driver Downloads for Windows Clients documentation published in the Samba Wiki. License: CC BY 4.0. Authors and contributors: See the history tab on the Wiki page.

Prerequisites

- Samba is set up as a print server

3.16.1. Basic information about printer drivers

This section provides general information about printer drivers.

Supported driver model version
Samba only supports the printer driver model version 3 which is supported in Windows 2000 and later, and Windows Server 2000 and later. Samba does not support the driver model version 4, introduced in Windows 8 and Windows Server 2012. However, these and later Windows versions also support version 3 drivers.

Package-aware drivers
Samba does not support package-aware drivers.

Preparing a printer driver for being uploaded

Before you can upload a driver to a Samba print server:

- Unpack the driver if it is provided in a compressed format.

- Some drivers require to start a setup application that installs the driver locally on a Windows host. In certain situations, the installer extracts the individual files into the operating system's temporary folder during the setup runs. To use the driver files for uploading:

  a. Start the installer.

  b. Copy the files from the temporary folder to a new location.

  c. Cancel the installation.

Ask your printer manufacturer for drivers that support uploading to a print server.

Providing 32-bit and 64-bit drivers for a printer to a client

To provide the driver for a printer for both 32-bit and 64-bit Windows clients, you must upload a driver with exactly the same name for both architectures. For example, if you are uploading the 32-bit driver named **Example PostScript** and the 64-bit driver named **Example PostScript (v1.0)**, the names do not match. Consequently, you can only assign one of the drivers to a printer and the driver will not be available for both architectures.

3.16.2. Enabling users to upload and preconfigure drivers

To be able to upload and preconfigure printer drivers, a user or a group needs to have the **SePrintOperatorPrivilege** privilege granted. A user must be added into the **printadmin** group. Red Hat Enterprise Linux automatically creates this group when you install the **samba** package. The **printadmin** group gets assigned the lowest available dynamic system GID that is lower than 1000.

Procedure

1. For example, to grant the **SePrintOperatorPrivilege** privilege to the **printadmin** group:

   > **# net rpc rights grant "printadmin" SePrintOperatorPrivilege -U "*DOMAIN*\administrator"**
   > Enter *DOMAIN*\administrator's password:
   > Successfully granted rights.

   > **NOTE**
   >
   > In a domain environment, grant **SePrintOperatorPrivilege** to a domain group. This enables you to centrally manage the privilege by updating a user's group membership.

2. To list all users and groups having **SePrintOperatorPrivilege** granted:

   > **# net rpc rights list privileges SePrintOperatorPrivilege -U "*DOMAIN*\administrator"**
   > Enter administrator's password:
   > SePrintOperatorPrivilege:
   >   BUILTIN\Administrators
   >   *DOMAIN*\printadmin

3.16.3. Setting up the print$ share

Windows operating systems download printer drivers from a share named **print$** from a print server. This share name is hard-coded in Windows and cannot be changed.

The following procedure explains how to share the **/var/lib/samba/drivers/** directory as **print$**, and enable members of the local **printadmin** group to upload printer drivers.

Procedure

1. Add the **[print$]** section to the **/etc/samba/smb.conf** file:

   ```
   [print$]
           path = /var/lib/samba/drivers/
           read only = no
           write list = @printadmin
           force group = @printadmin
           create mask = 0664
           directory mask = 2775
   ```

   Using these settings:

   - Only members of the **printadmin** group can upload printer drivers to the share.

   - The group of new created files and directories will be set to **printadmin**.

   - The permissions of new files will be set to **664**.

   - The permissions of new directories will be set to **2775**.

2. To upload only 64-bit drivers for all printers, include this setting in the **[global]** section in the **/etc/samba/smb.conf** file:

   ```
   spoolss: architecture = Windows x64
   ```

   Without this setting, Windows only displays drivers for which you have uploaded at least the 32-bit version.

3. Verify the **/etc/samba/smb.conf** file:

   ```
   # testparm
   ```

4. Reload the Samba configuration

   ```
   # smbcontrol all reload-config
   ```

5. Create the **printadmin** group if it does not exists:

   ```
   # groupadd printadmin
   ```

6. Grant the **SePrintOperatorPrivilege** privilege to the **printadmin** group.

   ```
   # net rpc rights grant "printadmin" SePrintOperatorPrivilege -U
   "DOMAIN\administrator"
   Enter DOMAIN\administrator's password:
   Successfully granted rights.
   ```

7. If you run SELinux in **enforcing** mode, set the **samba_share_t** context on the directory:

   ```
   # semanage fcontext -a -t samba_share_t "/var/lib/samba/drivers(/.)?" # *restorecon -
   Rv /var/lib/samba/drivers/
   ```

8. Set the permissions on the **/var/lib/samba/drivers/** directory:

   - If you use POSIX ACLs, set:

     ```
     # chgrp -R "printadmin" /var/lib/samba/drivers/
     # chmod -R 2775 /var/lib/samba/drivers/
     ```

   - If you use Windows ACLs, set:

Principal	Access	Apply to
CREATOR OWNER	Full control	Subfolders and files only
Authenticated Users	Read & execute, List folder contents, Read	This folder, subfolders, and files
printadmin	Full control	This folder, subfolders, and files

     For details about setting ACLs on Windows, see the Windows documentation.

Additional resources

- Enabling users to upload and preconfigure drivers.

3.16.4. Creating a GPO to enable clients to trust the Samba print server

For security reasons, recent Windows operating systems prevent clients from downloading non-package-aware printer drivers from an untrusted server. If your print server is a member in an AD, you can create a Group Policy Object (GPO) in your domain to trust the Samba server.

Prerequisites

- The Samba print server is a member of an AD domain.

- The Windows computer you are using to create the GPO must have the Windows Remote Server Administration Tools (RSAT) installed. For details, see the Windows documentation.

Procedure

1. Log into a Windows computer using an account that is allowed to edit group policies, such as the AD domain **Administrator** user.

2. Open the **Group Policy Management Console**.

3. Right-click to your AD domain and select **Create a GPO in this domain, and Link it here**.
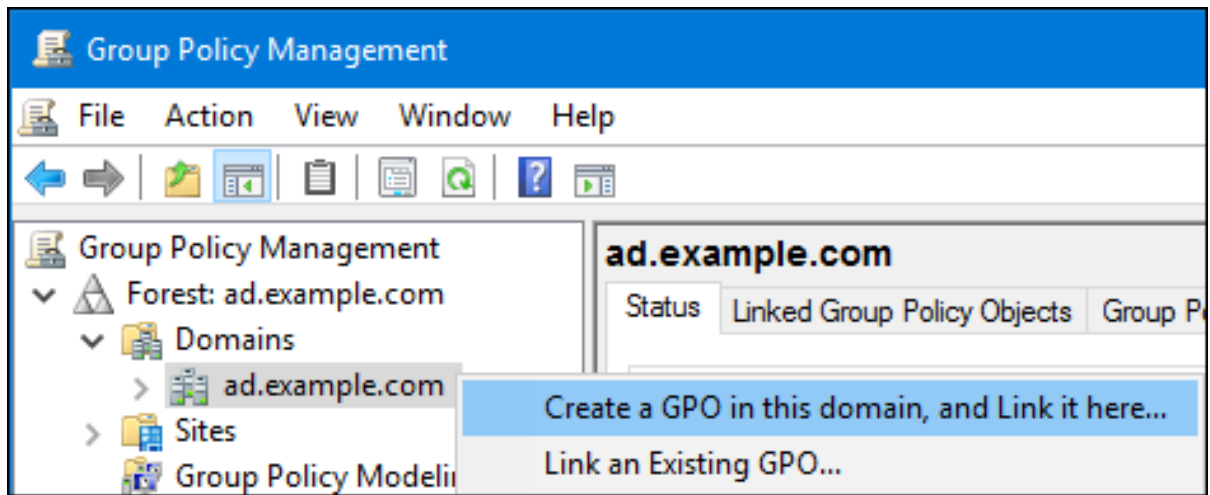
4. Enter a name for the GPO, such as **Legacy Printer Driver Policy** and click **OK**. The new GPO will be displayed under the domain entry.

5. Right-click to the newly-created GPO and select **Edit** to open the **Group Policy Management Editor**.

6. Navigate to **Computer Configuration → Policies → Administrative Templates → Printers**.
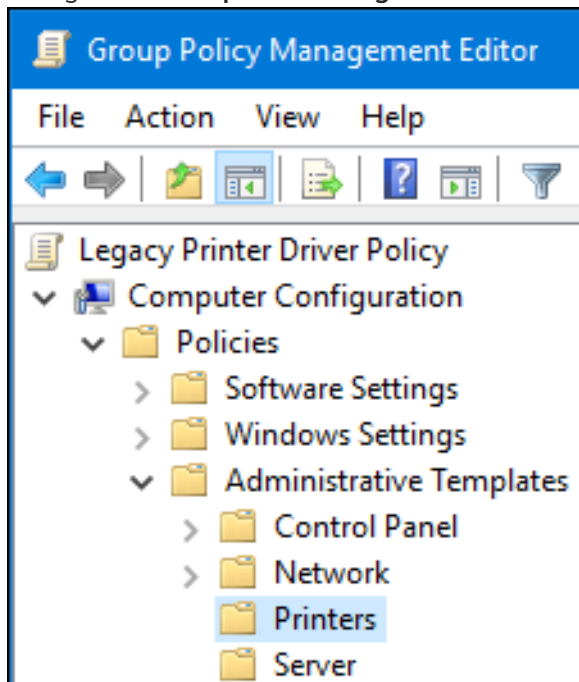
7. On the right side of the window, double-click **Point and Print Restriction** to edit the policy:

   a. Enable the policy and set the following options:

      i. Select **Users can only point and print to these servers** and enter the fully-qualified domain name (FQDN) of the Samba print server to the field next to this option.

      ii. In both check boxes under **Security Prompts**, select **Do not show warning or elevation prompt**.

b. Click OK.

8. Double-click **Package Point and Print - Approved servers** to edit the policy:

   a. Enable the policy and click the **Show** button.

   b. Enter the FQDN of the Samba print server.
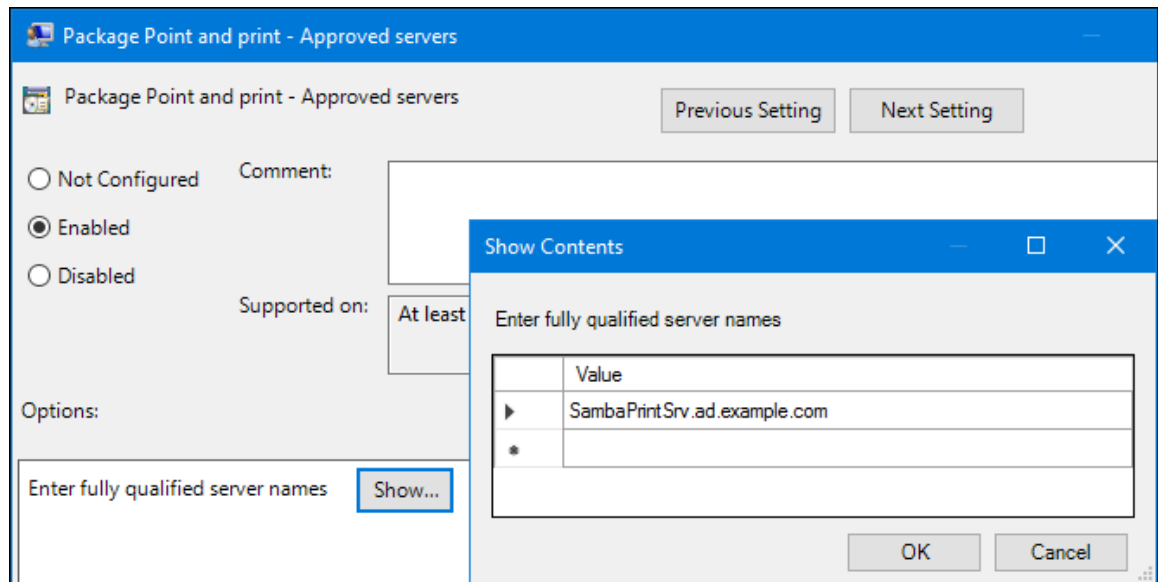
c. Close both the **Show Contents** and the policy's properties window by clicking **OK**.

9. Close the **Group Policy Management Editor**.

10. Close the **Group Policy Management Console**.

After the Windows domain members applied the group policy, printer drivers are automatically downloaded from the Samba server when a user connects to a printer.

Additional resources

- For using group policies, see the Windows documentation.

3.16.5. Uploading drivers and preconfiguring printers

Use the **Print Management** application on a Windows client to upload drivers and preconfigure printers hosted on the Samba print server. For further details, see the Windows documentation.

3.17. RUNNING SAMBA ON A SERVER WITH FIPS MODE ENABLED

This section provides an overview of the limitations of running Samba with FIPS mode enabled. It also provides the procedure for enabling FIPS mode on a Red Hat Enterprise Linux host running Samba.

3.17.1. Limitations of using Samba in FIPS mode

The following Samba modes and features work in FIPS mode under the indicated conditions:

- Samba as a domain member only in Active Directory (AD) or Red Hat Identity Management (IdM) environments with Kerberos authentication that uses AES ciphers.

- Samba as a file server on an Active Directory domain member. However, this requires that clients use Kerberos to authenticate to the server.

Due to the increased security of FIPS, the following Samba features and modes do not work if FIPS mode is enabled:

- NT LAN Manager (NTLM) authentication because RC4 ciphers are blocked

- The server message block version 1 (SMB1) protocol

- The stand-alone file server mode because it uses NTLM authentication

- NT4-style domain controllers

- NT4-style domain members. Note that Red Hat continues supporting the primary domain controller (PDC) functionality IdM uses in the background.

- Password changes against the Samba server. You can only perform password changes using Kerberos against an Active Directory domain controller.

The following feature is not tested in FIPS mode and, therefore, is not supported by Red Hat:

- Running Samba as a print server

3.17.2. Using Samba in FIPS mode

This section describes how to enable the FIPS mode on a RHEL host that runs Samba.

Prerequisites

- Samba is configured on the Red Hat Enterprise Linux host.

- Samba runs in a mode that is supported in FIPS mode.

Procedure

1. Enable the FIPS mode on RHEL:

   # fips-mode-setup --enable

2. Reboot the server:

   # reboot

3. Use the **testparm** utility to verify the configuration:

   # testparm -s

   If the command displays any errors or incompatibilities, fix them to ensure that Samba works correctly.

Additional resources

- Section 3.17.1, "Limitations of using Samba in FIPS mode"

3.18. TUNING THE PERFORMANCE OF A SAMBA SERVER

This chapter describes what settings can improve the performance of Samba in certain situations, and which settings can have a negative performance impact.

Parts of this section were adopted from the Performance Tuning documentation published in the Samba Wiki. License: CC BY 4.0. Authors and contributors: See the history tab on the Wiki page.

Prerequisites

- Samba is set up as a file or print server

3.18.1. Setting the SMB protocol version

Each new SMB version adds features and improves the performance of the protocol. The recent Windows and Windows Server operating systems always supports the latest protocol version. If Samba also uses the latest protocol version, Windows clients connecting to Samba benefit from the performance improvements. In Samba, the default value of the server max protocol is set to the latest supported stable SMB protocol version.
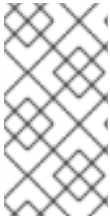
> **NOTE**
>
> To always have the latest stable SMB protocol version enabled, do not set the **server max protocol** parameter. If you set the parameter manually, you will need to modify the setting with each new version of the SMB protocol, to have the latest protocol version enabled.

The following procedure explains how to use the default value in the **server max protocol** parameter.

Procedure

1. Remove the **server max protocol** parameter from the **[global]** section in the **/etc/samba/smb.conf** file.

2. Reload the Samba configuration

   ```
   # smbcontrol all reload-config
   ```

3.18.2. Tuning shares with directories that contain a large number of files

Linux supports case-sensitive file names. For this reason, Samba needs to scan directories for uppercase and lowercase file names when searching or accessing a file. You can configure a share to create new files only in lowercase or uppercase, which improves the performance.

Prerequisites

- Samba is configured as a file server

Procedure

1. Rename all files on the share to lowercase.

   > **NOTE**
   >
   > Using the settings in this procedure, files with names other than in lowercase will no longer be displayed.

2. Set the following parameters in the share's section:

   ```
   case sensitive = true
   default case = lower
   ```

> preserve case = no
> short preserve case = no

For details about the parameters, see their descriptions in the **smb.conf(5)** man page.

3. Verify the **/etc/samba/smb.conf** file:

> # **testparm**

4. Reload the Samba configuration:

> # **smbcontrol all reload-config**

After you applied these settings, the names of all newly created files on this share use lowercase. Because of these settings, Samba no longer needs to scan the directory for uppercase and lowercase, which improves the performance.

3.18.3. Settings that can have a negative performance impact

By default, the kernel in Red Hat Enterprise Linux is tuned for high network performance. For example, the kernel uses an auto-tuning mechanism for buffer sizes. Setting the **socket options** parameter in the **/etc/samba/smb.conf** file overrides these kernel settings. As a result, setting this parameter decreases the Samba network performance in most cases.

To use the optimized settings from the Kernel, remove the **socket options** parameter from the **[global]** section in the **/etc/samba/smb.conf**.

3.19. CONFIGURING SAMBA TO BE COMPATIBLE WITH CLIENTS THAT REQUIRE AN SMB VERSION LOWER THAN THE DEFAULT

Samba uses a reasonable and secure default value for the minimum server message block (SMB) version it supports. However, if you have clients that require an older SMB version, you can configure Samba to support it.

3.19.1. Setting the minimum SMB protocol version supported by a Samba server

In Samba, the **server min protocol** parameter in the **/etc/samba/smb.conf** file defines the minimum server message block (SMB) protocol version the Samba server supports. This section describes how to change the minimum SMB protocol version.

> NOTE
>
> By default, Samba on RHEL 8.2 and later supports only SMB2 and newer protocol versions. Red Hat recommends to not use the deprecated SMB1 protocol. However, if your environment requires SMB1, you can manually set the **server min protocol** parameter to **NT1** to re-enable SMB1.

Prerequisites

- Samba is installed and configured.

Procedure

1. Edit the **/etc/samba/smb.conf** file, add the **server min protocol** parameter, and set the parameter to the minimum SMB protocol version the server should support. For example, to set the minimum SMB protocol version to **SMB3**, add:

   ```
   server min protocol = SMB3
   ```

2. Restart the **smb** service:

   ```
   # systemctl restart smb
   ```

Additional resources

- **smb.conf(5)** man page

3.20. FREQUENTLY USED SAMBA COMMAND-LINE UTILITIES

This chapter describes frequently used commands when working with a Samba server.

3.20.1. Using the net ads join and net rpc join commands

Using the **join** subcommand of the **net** utility, you can join Samba to an AD or NT4 domain. To join the domain, you must create the **/etc/samba/smb.conf** file manually, and optionally update additional configurations, such as PAM.

> **IMPORTANT**
>
> Red Hat recommends using the **realm** utility to join a domain. The **realm** utility automatically updates all involved configuration files.

Procedure

1. Manually create the **/etc/samba/smb.conf** file with the following settings:

   - For an AD domain member:

     ```
     [global]
     workgroup = domain_name
     security = ads
     passdb backend = tdbsam
     realm = AD_REALM
     ```

   - For an NT4 domain member:

     ```
     [global]
     workgroup = domain_name
     security = user
     passdb backend = tdbsam
     ```

2. Add an ID mapping configuration for the ***** default domain and for the domain you want to join to the **[global**] section in the **/etc/samba/smb.conf** file.

3. Verify the **/etc/samba/smb.conf** file:

> # **testparm**

4. Join the domain as the domain administrator:

   - To join an AD domain:

     > # **net ads join -U "DOMAIN\administrator"**

   - To join an NT4 domain:

     > # **net rpc join -U "DOMAIN\administrator"**

5. Append the **winbind** source to the **passwd** and **group** database entry in the **/etc/nsswitch.conf** file:

   > passwd: files **winbind**
   > group: files **winbind**

6. Enable and start the **winbind** service:

   > # **systemctl enable --now winbind**

7. Optionally, configure PAM using the **authselect** utility.
   For details, see the **authselect(8)** man page.

8. Optionally for AD environments, configure the Kerberos client.
   For details, see the documentation of your Kerberos client.

Additional resources

- Joining Samba to a domain .

- Understanding and configuring Samba ID mapping .

3.20.2. Using the net rpc rights command

In Windows, you can assign privileges to accounts and groups to perform special operations, such as setting ACLs on a share or upload printer drivers. On a Samba server, you can use the **net rpc rights** command to manage privileges.

Listing privileges you can set
To list all available privileges and their owners, use the **net rpc rights list** command. For example:

```
# net rpc rights list -U "DOMAIN\administrator"
Enter DOMAIN\administrator's password:
    SeMachineAccountPrivilege  Add machines to domain
     SeTakeOwnershipPrivilege  Take ownership of files or other objects
            SeBackupPrivilege  Back up files and directories
           SeRestorePrivilege  Restore files and directories
    SeRemoteShutdownPrivilege  Force shutdown from a remote system
      SePrintOperatorPrivilege  Manage printers
```

> SeAddUsersPrivilege Add users and groups to the domain
> SeDiskOperatorPrivilege Manage disk shares
> SeSecurityPrivilege System security

Granting privileges

To grant a privilege to an account or group, use the **net rpc rights grant** command.

For example, grant the **SePrintOperatorPrivilege** privilege to the *DOMAIN\printadmin* group:

> # **net rpc rights grant "*DOMAIN*\printadmin" SePrintOperatorPrivilege -U "*DOMAIN*\administrator"**
> Enter *DOMAIN*\administrator's password:
> Successfully granted rights.

Revoking privileges

To revoke a privilege from an account or group, use the **net rpc rights revoke** command.

For example, to revoke the **SePrintOperatorPrivilege** privilege from the *DOMAIN\printadmin* group:

> # **net rpc rights remoke "*DOMAIN*\printadmin" SePrintOperatorPrivilege -U "*DOMAIN*\administrator"**
> Enter *DOMAIN*\administrator's password:
> Successfully revoked rights.

3.20.3. Using the net rpc share command

The **net rpc share** command provides the capability to list, add, and remove shares on a local or remote Samba or Windows server.

Listing shares

To list the shares on an SMB server, use the **net rpc share list** command. Optionally, pass the **-S *server_name*** parameter to the command to list the shares of a remote server. For example:

> # **net rpc share list -U "*DOMAIN*\administrator" -S *server_name***
> Enter *DOMAIN*\administrator's password:
> IPC$
> *share_1*
> *share_2*
> ...

> **NOTE**
>
> Shares hosted on a Samba server that have **browseable = no** set in their section in the **/etc/samba/smb.conf** file are not displayed in the output.

Adding a share

The **net rpc share add** command enables you to add a share to an SMB server.

For example, to add a share named **example** on a remote Windows server that shares the **C:\example** directory:

> # **net rpc share add example="C:\example" -U "*DOMAIN*\administrator" -S *server_name***

> **NOTE**
>
> You must omit the trailing backslash in the path when specifying a Windows directory name.

To use the command to add a share to a Samba server:

- The user specified in the **-U** parameter must have the **SeDiskOperatorPrivilege** privilege granted on the destination server.

- You must write a script that adds a share section to the **/etc/samba/smb.conf** file and reloads Samba. The script must be set in the **add share command** parameter in the **[global]** section in **/etc/samba/smb.conf**. For further details, see the **add share command** description in the **smb.conf(5)** man page.

Removing a share

The **net rpc share delete** command enables you to remove a share from an SMB server.

For example, to remove the share named example from a remote Windows server:

```
# net rpc share delete example -U "DOMAIN\administrator" -S server_name
```

To use the command to remove a share from a Samba server:

- The user specified in the **-U** parameter must have the **SeDiskOperatorPrivilege** privilege granted.

- You must write a script that removes the share's section from the **/etc/samba/smb.conf** file and reloads Samba. The script must be set in the **delete share command** parameter in the **[global]** section in **/etc/samba/smb.conf**. For further details, see the **delete share command** description in the **smb.conf(5)** man page.

3.20.4. Using the net user command

The **net user** command enables you to perform the following actions on an AD DC or NT4 PDC:

- List all user accounts

- Add users

- Remove Users

> **NOTE**
>
> Specifying a connection method, such as **ads** for AD domains or **rpc** for NT4 domains, is only required when you list domain user accounts. Other user-related subcommands can auto-detect the connection method.

Pass the **-U** *user_name* parameter to the command to specify a user that is allowed to perform the requested action.

Listing domain user accounts

To list all users in an AD domain:

```
# net ads user -U "DOMAIN\administrator"
```

To list all users in an NT4 domain:

> # **net rpc user -U "*DOMAIN*\administrator"**

Adding a user account to the domain
On a Samba domain member, you can use the **net user add** command to add a user account to the domain.

For example, add the **user** account to the domain:

1. Add the account:

   > # **net user add *user password* -U "*DOMAIN*\administrator"**
   > User *user* added

2. Optionally, use the remote procedure call (RPC) shell to enable the account on the AD DC or NT4 PDC. For example:

   > # **net rpc shell -U *DOMAIN*\administrator -S *DC_or_PDC_name***
   > Talking to domain DOMAIN (*S-1-5-21-1424831554-512457234-5642315751*)
   >
   > net rpc> **user edit disabled *user*: no**
   > Set *user*'s disabled flag from [yes] to [no]
   >
   > net rpc> **exit**

Deleting a user account from the domain
On a Samba domain member, you can use the **net user delete** command to remove a user account from the domain.

For example, to remove the **user** account from the domain:

> # **net user delete *user* -U "*DOMAIN*\administrator"**
> User *user* deleted

3.20.5. Using the rpcclient utility

The **rpcclient** utility enables you to manually execute client-side Microsoft Remote Procedure Call (MS-RPC) functions on a local or remote SMB server. However, most of the features are integrated into separate utilities provided by Samba. Use **rpcclient** only for testing MS-PRC functions.

Prerequisites

- The **samba-client** package is installed.

Examples
For example, you can use the **rpcclient** utility to:

- Manage the printer Spool Subsystem (SPOOLSS).

  Example 3.7. Assigning a Driver to a Printer

  > # **rpcclient *server_name* -U "*DOMAIN*\administrator" -c 'setdriver "*printer_name*" "*driver_name*"'**

> Enter *DOMAIN*\administrators password:
> Successfully set *printer_name* to driver *driver_name*.

- Retrieve information about an SMB server.

  Example 3.8. Listing all File Shares and Shared Printers

  > # **rpcclient** *server_name* **-U "*DOMAIN*\administrator" -c 'netshareenum'**
  > Enter *DOMAIN*\administrators password:
  > netname: *Example_Share*
  > remark:
  > path: *C:\srv\samba\example_share*
  > password:
  > netname: *Example_Printer*
  > remark:
  > path: *C:\var\spool\samba*
  > password:

- Perform actions using the Security Account Manager Remote (SAMR) protocol.

  Example 3.9. Listing Users on an SMB Server

  > # **rpcclient** *server_name* **-U "*DOMAIN*\administrator" -c 'enumdomusers'**
  > Enter *DOMAIN*\administrators password:
  > user:[*user1*] rid:[*0x3e8*]
  > user:[*user2*] rid:[*0x3e9*]

  If you run the command against a standalone server or a domain member, it lists the users in the local database. Running the command against an AD DC or NT4 PDC lists the domain users.

Additional resources

- **rpcclient(1)** man page

3.20.6. Using the samba-regedit application

Certain settings, such as printer configurations, are stored in the registry on the Samba server. You can use the ncurses-based **samba-regedit** application to edit the registry of a Samba server.

```
Path: ...AL_MACHINE/SOFTWARE/Microsoft/Windows NT/CurrentVersion/Print/Printers/
 ┌─Key────────────────────┬─Value─────────────────────────────────────────────────┐
 │Name                    │Name                │ Type       │ Data                │
 │+Example-Printer        │Attributes          │ REG_DWORD  │ 0x00001848 (6216)   │
 │                        │ChangeID            │ REG_DWORD  │ 0x00160374 (1442676)│
 │                        │Datatype            │ REG_SZ     │ RAW                 │
 │                        │Default Priority    │ REG_DWORD  │ 0x00000001 (1)      │
 │                        │Description         │ REG_SZ     │                     │
 │                        │Location            │ REG_SZ     │                     │
 │                        │Name                │ REG_SZ     │ Example-Printer     │
 │                        │Parameters          │ REG_SZ     │                     │
 │                        │Port                │ REG_SZ     │ Samba Printer Port  │
 │                        │Print Processor     │ REG_SZ     │ winprint            │
 │                        │Printer Driver      │ REG_SZ     │ Example Printer Driver│
 │                        │Priority            │ REG_DWORD  │ 0x00000001 (1)      │
 │                        │Security            │ REG_BINARY │ (248 bytes)         │
 │                        │Separator File      │ REG_SZ     │                     │
 │                        │Share Name          │ REG_SZ     │ Example-Printer     │
 │                        │StartTime           │ REG_DWORD  │ 0x00000000 (0)      │
 │                        │Status              │ REG_DWORD  │ 0x00000000 (0)      │
 │                        │UntilTime           │ REG_DWORD  │ 0x00000000 (0)      │
 ├────────────────────────┴───────────────────────────────────────────────────────┤
 │[n] New Value [d] Del Value [ENTER] Edit [b] Edit binary            VALUES        │
 │[TAB] Switch sections [q] Quit [UP] List up [DOWN] List down [/] Search [x] Next  │
 └──────────────────────────────────────────────────────────────────────────────────┘
```

Prerequisites

- The **samba-client** package is installed.

Procedure

To start the application, enter:

> # **samba-regedit**

Use the following keys:

- Cursor up and cursor down: Navigate through the registry tree and the values.

- **Enter**: Opens a key or edits a value.

- **Tab**: Switches between the **Key** and **Value** pane.

- **Ctrl**+**C**: Closes the application.

3.20.7. Using the smbcontrol utility

The **smbcontrol** utility enables you to send command messages to the **smbd**, **nmbd**, **winbindd**, or all of these services. These control messages instruct the service, for example, to reload its configuration.

The procedure in this section shows how to to reload the configuration of the **smbd**, **nmbd**, **winbindd** services by sending the **reload-config** message type to the **all** destination.

Prerequisites

- The **samba-common-tools** package is installed.

Procedure

-

> # **smbcontrol all reload-config**

Additional resources

- **smbcontrol(1)** man page

3.20.8. Using the smbpasswd utility

The **smbpasswd** utility manages user accounts and passwords in the local Samba database.

Prerequisites

- The **samba-common-tools** package is installed.

Procedure

1. If you run the command as a user, **smbpasswd** changes the Samba password of the user who run the command. For example:

   > [user@server ~]$ **smbpasswd**
   > New SMB password: *password*
   > Retype new SMB password: *password*

2. If you run **smbpasswd** as the **root** user, you can use the utility, for example, to:

   - Create a new user:

     > [root@server ~]# **smbpasswd -a *user_name***
     > New SMB password: ***password***
     > Retype new SMB password: ***password***
     > Added user *user_name*.

     > **NOTE**
     >
     > Before you can add a user to the Samba database, you must create the account in the local operating system. See the Adding a new user from the command line section in the Configuring basic system settings guide.

   - Enable a Samba user:

     > [root@server ~]# **smbpasswd -e *user_name***
     > Enabled user *user_name*.

   - Disable a Samba user:

     > [root@server ~]# **smbpasswd -x *user_name***
     > Disabled user *user_name*

   - Delete a user:

     > [root@server ~]# **smbpasswd -x *user_name***
     > Deleted user *user_name*.

Additional resources

- **smbpasswd(8)** man page

3.20.9. Using the smbstatus utility

The **smbstatus** utility reports on:

- Connections per PID of each **smbd** daemon to the Samba server. This report includes the user name, primary group, SMB protocol version, encryption, and signing information.

- Connections per Samba share. This report includes the PID of the **smbd** daemon, the IP of the connecting machine, the time stamp when the connection was established, encryption, and signing information.

- A list of locked files. The report entries include further details, such as opportunistic lock (oplock) types

Prerequisites

- The **samba** package is installed.

- The **smbd** service is running.

Procedure

```
# smbstatus

Samba version 4.15.2
PID  Username           Group           Machine                    Protocol Version  Encryption
Signing
....----------------------------------------------------------------------------------------------------------
-
963  DOMAIN\administrator  DOMAIN\domain users  client-pc  (ipv4:192.0.2.1:57786)  SMB3_02
-          AES-128-CMAC

Service  pid  Machine    Connected at              Encryption  Signing:
....---------------------------------------------------------------------------
example  969  192.0.2.1  Thu Nov  1 10:00:00 2018 CEST  -          AES-128-CMAC

Locked files:
Pid  Uid   DenyMode    Access     R/W    Oplock      SharePath        Name       Time
....---------------------------------------------------------------------------------------------------
969  10000  DENY_WRITE 0x120089  RDONLY  LEASE(RWH)  /srv/samba/example  file.txt  Thu
Nov  1 10:00:00 2018
```

Additional resources

- **smbstatus(1)** man page

3.20.10. Using the smbtar utility

The **smbtar** utility backs up the content of an SMB share or a subdirectory of it and stores the content in a **tar** archive. Alternatively, you can write the content to a tape device.

Prerequisites

- The **samba-client** package is installed.

Procedure

- Use the following command to back up the content of the **demo** directory on the //**server**/**example**/ share and store the content in the /**root**/**example.tar** archive:

      # **smbtar -s** *server* **-x** *example* **-u** *user_name* **-p** *password* **-t** */root/example.tar*

Additional resources

- **smbtar(1)** man page

3.20.11. Using the wbinfo utility

The **wbinfo** utility queries and returns information created and used by the **winbindd** service.

Prerequisites

- The **samba-winbind-clients** package is installed.

Procedure

You can use **wbinfo**, for example, to:

- List domain users:

      # **wbinfo -u**
      AD\administrator
      AD\guest
      ...

- List domain groups:

      # **wbinfo -g**
      AD\domain computers
      AD\domain admins
      AD\domain users
      ...

- Display the SID of a user:

      # **wbinfo --name-to-sid="AD\administrator"**
      S-1-5-21-1762709870-351891212-3141221786-500 SID_USER (1)

- Display information about domains and trusts:

      # **wbinfo --trusted-domains --verbose**
      Domain Name   DNS Domain   Trust Type Transitive In   Out
      BUILTIN   None   Yes   Yes Yes

```
server                         None       Yes        Yes  Yes
DOMAIN1      domain1.example.com  None       Yes        Yes  Yes
DOMAIN2      domain2.example.com  External   No         Yes  Yes
```

Additional resources

- **wbinfo(1)** man page

3.21. ADDITIONAL RESOURCES

- **smb.conf(5)** man page

- **/usr/share/docs/samba-version/** directory contains general documentation, example scripts, and LDAP schema files, provided by the Samba project

- Setting up Samba and the Clustered Trivial Database (CDTB) to share directories stored on an GlusterFS volume

- Mounting an SMB Share on Red Hat Enterprise Linux

CHAPTER 4. SETTING UP AND CONFIGURING A BIND DNS SERVER

BIND is a feature-rich DNS server that is fully compliant with the Internet Engineering Task Force (IETF) DNS standards and draft standards. For example, administrators frequently use BIND as:

- Caching DNS server in the local network

- Authoritative DNS server for zones

- Secondary server to provide high availability for zones

4.1. CONSIDERATIONS ABOUT PROTECTING BIND WITH SELINUX OR RUNNING IT IN A CHANGE-ROOT ENVIRONMENT

To secure a BIND installation, you can:

- Run the **named** service without a change-root environment. In this case, SELinux in **enforcing** mode prevents exploitation of known BIND security vulnerabilities. By default, Red Hat Enterprise Linux uses SELinux in **enforcing** mode.
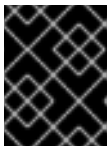
  > **IMPORTANT**
  >
  > Running BIND on RHEL with SELinux in **enforcing** mode is more secure than running BIND in a change-root environment.

- Run the **named-chroot** service in a change-root environment.
  Using the change-root feature, administrators can define that the root directory of a process and its sub-processes is different to the / directory. When you start the **named-chroot** service, BIND switches its root directory to **/var/named/chroot/**. As a consequence, the service uses **mount --bind** commands to make the files and directories listed in **/etc/named-chroot.files** available in **/var/named/chroot/**, and the process has no access to files outside of **/var/named/chroot/**.

If you decide to use BIND:

- In normal mode, use the **named** service.

- In a change-root environment, use the **named-chroot** service. This requires that you install, additionally, the **named-chroot** package.

Additional resources

- The **Red Hat SELinux BIND security profile** section in the **named(8)** man page

4.2. THE BIND ADMINISTRATOR REFERENCE MANUAL

The comprehensive **BIND Administrator Reference Manual**, that is included in the **bind** package, provides:

- Configuration examples

- Documentation on advanced features

- A configuration reference

- Security considerations

To display the **BIND Administrator Reference Manual** on a host that has the **bind** package installed, open the **/usr/share/doc/bind/Bv9ARM.html** file in a browser.

4.3. CONFIGURING BIND AS A CACHING DNS SERVER

By default, the BIND DNS server resolves and caches successful and failed lookups. The service then answers requests to the same records from its cache. This significantly improves the speed of DNS lookups.

Prerequisites

- The IP address of the server is static.

Procedure

1. Install the **bind** and **bind-utils** packages:

   ```
   # yum install bind bind-utils
   ```

   These packages provide BIND 9.11. If you require BIND 9.16, install the **bind9.16** and **bind9.16-utils** packages.

2. If you want to run BIND in a change-root environment install the **bind-chroot** package:

   ```
   # yum install bind-chroot
   ```

   Note that running BIND on a host with SELinux in **enforcing** mode, which is default, is more secure.

3. Edit the **/etc/named.conf** file, and make the following changes in the **options** statement:

   a. Update the **listen-on** and **listen-on-v6** statements to specify on which IPv4 and IPv6 interfaces BIND should listen:

      ```
      listen-on port 53 { 127.0.0.1; 192.0.2.1; };
      listen-on-v6 port 53 { ::1; 2001:db8:1::1; };
      ```

   b. Update the **allow-query** statement to configure from which IP addresses and ranges clients can query this DNS server:

      ```
      allow-query { localhost; 192.0.2.0/24; 2001:db8:1::/64; };
      ```

   c. Add an **allow-recursion** statement to define from which IP addresses and ranges BIND accepts recursive queries:

      ```
      allow-recursion { localhost; 192.0.2.0/24; 2001:db8:1::/64; };
      ```

> **WARNING**
>
> Do not allow recursion on public IP addresses of the server. Otherwise, the server can become part of large-scale DNS amplification attacks.

d. By default, BIND resolves queries by recursively querying from the root servers to an authoritative DNS server. Alternatively, you can configure BIND to forward queries to other DNS servers, such as the ones of your provider. In this case, add a **forwarders** statement with the list of IP addresses of the DNS servers that BIND should forward queries to:

> **forwarders {** *198.51.100.1; 203.0.113.5;* **};**

As a fall-back behavior, BIND resolves queries recursively if the forwarder servers do not respond. To disable this behavior, add a **forward only;** statement.

4. Verify the syntax of the **/etc/named.conf** file:

> **# named-checkconf**

If the command displays no output, the syntax is correct.

5. Update the **firewalld** rules to allow incoming DNS traffic:

> **# firewall-cmd --permanent --add-service=dns**
> **# firewall-cmd --reload**

6. Start and enable BIND:

> **# systemctl enable --now named**

If you want to run BIND in a change-root environment, use the **systemctl enable --now named-chroot** command to enable and start the service.

Verification

1. Use the newly set up DNS server to resolve a domain:

> **# dig @***localhost www.example.org*
>
> ...
> *www.example.org.   86400   IN   A   198.51.100.34*
>
> ;; Query time: *917 msec*
>
> ...

This example assumes that BIND runs on the same host and responds to queries on the **localhost** interface.

After querying a record for the first time, BIND adds the entry to its cache.

2. Repeat the previous query:

```
# dig @localhost www.example.org
...
www.example.org.    85332    IN    A    198.51.100.34

;; Query time: 1 msec
...
```

Because of the cached entry, further requests for the same record are significantly faster until the entry expires.

Next steps

- Configure the clients in your network to use this DNS server. If a DHCP server provides the DNS server setting to the clients, update the DHCP server's configuration accordingly.

Additional resources

- Considerations about protecting BIND with SELinux or running it in a change-root environment

- **named.conf(5)** man page

- **/usr/share/doc/bind/sample/etc/named.conf**

- The BIND Administrator Reference Manual

4.4. CONFIGURING LOGGING ON A BIND DNS SERVER

The configuration in the default **/etc/named.conf** file, as provided by the **bind** package, uses the **default_debug** channel and logs messages to the **/var/named/data/named.run** file. The **default_debug** channel only logs entries when the server's debug level is non-zero.

Using different channels and categories, you can configure BIND to write different events with a defined severity to separate files.

Prerequisites

- BIND is already configured, for example, as a caching name server.

- The **named** or **named-chroot** service is running.

Procedure

1. Edit the **/etc/named.conf** file, and add **category** and **channel** phrases to the **logging** statement, for example:

```
logging {
    ...

  category notify { zone_transfer_log; };
  category xfer-in { zone_transfer_log; };
  category xfer-out { zone_transfer_log; };
  channel zone_transfer_log {
      file "/var/named/log/transfer.log" versions 10 size 50m;
      print-time yes;
```

```
        print-category yes;
        print-severity yes;
        severity info;
    };


    ...
};
```

With this example configuration, BIND logs messages related to zone transfers to
/var/named/log/transfer.log. BIND creates up to **10** versions of the log file and rotates them if
they reach a maximum size of **50** MB.

The **category** phrase defines to which channels BIND sends messages of a category.

The **channel** phrase defines the destination of log messages including the number of versions,
the maximum file size, and the severity level BIND should log to a channel. Additional settings,
such as enabling logging the time stamp, category, and severity of an event are optional, but
useful for debugging purposes.

2. Create the log directory if it does not exist, and grant write permissions to the **named** user on
   this directory:

   ```
   # mkdir /var/named/log/
   # chown named:named /var/named/log/
   # chmod 700 /var/named/log/
   ```

3. Verify the syntax of the **/etc/named.conf** file:

   ```
   # named-checkconf
   ```

   If the command displays no output, the syntax is correct.

4. Restart BIND:

   ```
   # systemctl restart named
   ```

   If you run BIND in a change-root environment, use the **systemctl restart named-chroot**
   command to restart the service.

Verification

- Display the content of the log file:

  ```
  # cat /var/named/log/transfer.log
  ...
  06-Jul-2022 15:08:51.261 xfer-out: info: client @0x7fecbc0b0700 192.0.2.2#36121/key
  example-transfer-key (example.com): transfer of 'example.com/IN': AXFR started: TSIG
  example-transfer-key (serial 2022070603)
  06-Jul-2022 15:08:51.261 xfer-out: info: client @0x7fecbc0b0700 192.0.2.2#36121/key
  example-transfer-key (example.com): transfer of 'example.com/IN': AXFR ended
  ```

Additional resources

- **named.conf(5)** man page

- The BIND Administrator Reference Manual

4.5. WRITING BIND ACLS

Controlling access to certain features of BIND can prevent unauthorized access and attacks, such as denial of service (DoS). BIND access control list (**acl**) statements are lists of IP addresses and ranges. Each ACL has a nickname that you can use in several statements, such as **allow-query**, to refer to the specified IP addresses and ranges.

> **WARNING**
>
> BIND uses only the first matching entry in an ACL. For example, if you define an ACL **{ 192.0.2/24; !192.0.2.1; }** and the host with IP address **192.0.2.1** connects, access is granted even if the second entry excludes this address.

BIND has the following built-in ACLs:

- **none**: Matches no hosts.

- **any**: Matches all hosts.

- **localhost**: Matches the loopback addresses **127.0.0.1** and **::1**, as well as the IP addresses of all interfaces on the server that runs BIND.

- **localnets**: Matches the loopback addresses **127.0.0.1** and **::1**, as well as all subnets the server that runs BIND is directly connected to.

Prerequisites

- BIND is already configured, for example, as a caching name server.

- The **named** or **named-chroot** service is running.

Procedure

1. Edit the **/etc/named.conf** file and make the following changes:

   a. Add **acl** statements to the file. For example, to create an ACL named **internal-networks** for **127.0.0.1**, **192.0.2.0/24**, and **2001:db8:1::/64**, enter:

      ```
      acl internal-networks { 127.0.0.1; 192.0.2.0/24; 2001:db8:1::/64; };
      acl dmz-networks { 198.51.100.0/24; 2001:db8:2::/64; };
      ```

   b. Use the ACL's nickname in statements that support them, for example:

      ```
      allow-query { internal-networks; dmz-networks; };
      allow-recursion { internal-networks; };
      ```

2. Verify the syntax of the **/etc/named.conf** file:

> # **named-checkconf**

If the command displays no output, the syntax is correct.

3. Reload BIND:

> # **systemctl reload named**

If you run BIND in a change-root environment, use the **systemctl reload named-chroot** command to reload the service.

Verification

- Execute an action that triggers a feature which uses the configured ACL. For example, the ACL in this procedure allows only recursive queries from the defined IP addresses. In this case, enter the following command on a host that is not within the ACL's definition to attempt resolving an external domain:

> # **dig +short** *@192.0.2.1 www.example.com*

If the command returns no output, BIND denied access, and the ACL works. For a verbose output on the client, use the command without **+short** option:

> # **dig** *@192.0.2.1 www.example.com*
> ...
> ;; WARNING: recursion requested but not available
> ...

Additional resources

- The **Access control lists** section in the The BIND Administrator Reference Manual .

4.6. CONFIGURING ZONES ON A BIND DNS SERVER

A DNS zone is a database with resource records for a specific sub-tree in the domain space. For example, if you are responsible for the **example.com** domain, you can set up a zone for it in BIND. As a result, clients can, resolve **www.example.com** to the IP address configured in this zone.

4.6.1. The SOA record in zone files

The start of authority (SOA) record is a required record in a DNS zone. This record is important, for example, if multiple DNS servers are authoritative for a zone but also to DNS resolvers.

A SOA record in BIND has the following syntax:

> *name class type mname rname serial refresh retry expire minimum*

For better readability, administrators typically split the record in zone files into multiple lines with comments that start with a semicolon (**;**). Note that, if you split a SOA record, parentheses keep the record together:

> **@ IN SOA** *ns1.example.com. hostmaster.example.com.* **(**
>   **2022070601 ; serial number**

```
1d        ; refresh period
3h        ; retry period
3d        ; expire time
3h )      ; minimum TTL
```

> **IMPORTANT**
>
> Note the trailing dot at the end of the fully-qualified domain names (FQDNs). FQDNs consist of multiple domain labels, separated by dots. Because the DNS root has an empty label, FQDNs end with a dot. Therefore, BIND appends the zone name to names without a trailing dot. A hostname without a trailing dot, for example, **ns1.example.com** would be expanded to **ns1.example.com.example.com.**, which is not the correct address of the primary name server.

These are the fields in a SOA record:

- **name**: The name of the zone, the so-called **origin**. If you set this field to **@**, BIND expands it to the zone name defined in **/etc/named.conf**.

- **class**: In SOA records, you must set this field always to Internet (**IN**).

- **type**: In SOA records, you must set this field always to **SOA**.

- **mname** (master name): The hostname of the primary name server of this zone.

- **rname** (responsible name): The email address of who is responsible for this zone. Note that the format is different. You must replace the at sign (**@**) with a dot (**.**).

- **serial**: The version number of this zone file. Secondary name servers only update their copies of the zone if the serial number on the primary server is higher.
  The format can be any numeric value. A commonly-used format is **<year><month><day><two-digit-number>**. With this format, you can, theoretically, change the zone file up to a hundred times per day.

- **refresh**: The amount of time secondary servers should wait before checking the primary server if the zone was updated.

- **retry**: The amount of time after that a secondary server retries to query the primary server after a failed attempt.

- **expire**: The amount of time after that a secondary server stops querying the primary server, if all previous attempts failed.

- **minimum**: RFC 2308 changed the meaning of this field to the negative caching time. Compliant resolvers use it to determine how long to cache **NXDOMAIN** name errors.

> **NOTE**
>
> A numeric value in the **refresh**, **retry**, **expire**, and **minimum** fields define a time in seconds. However, for better readability, use time suffixes, such as **m** for minute, **h** for hours, and **d** for days. For example, **3h** stands for 3 hours.

Additional resources

- RFC 1035: Domain names - implementation and specification

- **RFC 1034**: Domain names - concepts and facilities

- **RFC 2308**: Negative caching of DNS queries (DNS cache)

4.6.2. Setting up a forward zone on a BIND primary server

Forward zones map names to IP addresses and other information. For example, if you are responsible for the domain **example.com**, you can set up a forward zone in BIND to resolve names, such as **www.example.com**.

Prerequisites

- BIND is already configured, for example, as a caching name server.

- The **named** or **named-chroot** service is running.

Procedure

1. Add a zone definition to the **/etc/named.conf** file:

   ```
   zone "example.com" {
       type master;
       file "example.com.zone";
       allow-query { any; };
       allow-transfer { none; };
   };
   ```

   These settings define:

   - This server as the primary server (**type master**) for the **example.com** zone.

   - The **/var/named/example.com.zone** file is the zone file. If you set a relative path, as in this example, this path is relative to the directory you set in **directory** in the **options** statement.

   - Any host can query this zone. Alternatively, specify IP ranges or BIND access control list (ACL) nicknames to limit the access.

   - No host can transfer the zone. Allow zone transfers only when you set up secondary servers and only for the IP addresses of the secondary servers.

2. Verify the syntax of the **/etc/named.conf** file:

   ```
   # named-checkconf
   ```

   If the command displays no output, the syntax is correct.

3. Create the **/var/named/example.com.zone** file, for example, with the following content:

   ```
   $TTL 8h
   @ IN SOA ns1.example.com. hostmaster.example.com. (
               2022070601 ; serial number
               1d       ; refresh period
               3h       ; retry period
               3d       ; expire time
               3h )     ; minimum TTL
   ```

```
            IN NS   ns1.example.com.
            IN MX   10 mail.example.com.

www             IN A    192.0.2.30
www             IN AAAA 2001:db8:1::30
ns1         IN A    192.0.2.1
ns1         IN AAAA 2001:db8:1::1
mail        IN A    192.0.2.20
mail        IN AAAA 2001:db8:1::20
```

This zone file:

- Sets the default time-to-live (TTL) value for resource records to 8 hours. Without a time suffix, such as **h** for hour, BIND interprets the value as seconds.

- Contains the required SOA resource record with details about the zone.

- Sets **ns1.example.com** as an authoritative DNS server for this zone. To be functional, a zone requires at least one name server (**NS**) record. However, to be compliant with RFC 1912, you require at least two name servers.

- Sets **mail.example.com** as the mail exchanger (**MX**) of the **example.com** domain. The numeric value in front of the host name is the priority of the record. Entries with a lower value have a higher priority.

- Sets the IPv4 and IPv6 addresses of **www.example.com**, **mail.example.com**, and **ns1.example.com**.

4. Set secure permissions on the zone file that allow only the **named** group to read it:

   ```
   # chown root:named /var/named/example.com.zone
   # chmod 640 /var/named/example.com.zone
   ```

5. Verify the syntax of the **/var/named/example.com.zone** file:

   ```
   # named-checkzone example.com /var/named/example.com.zone
   zone example.com/IN: loaded serial 2022070601
   OK
   ```

6. Reload BIND:

   ```
   # systemctl reload named
   ```

   If you run BIND in a change-root environment, use the **systemctl reload named-chroot** command to reload the service.

Verification

- Query different records from the **example.com** zone, and verify that the output matches the records you have configured in the zone file:

   ```
   # dig +short @localhost AAAA www.example.com
   2001:db8:1::30
   ```

```
# dig +short @localhost NS example.com
ns1.example.com.

# dig +short @localhost A ns1.example.com
192.0.2.1
```

This example assumes that BIND runs on the same host and responds to queries on the **localhost** interface.

Additional resources

- The SOA record in zone files

- Writing BIND ACLs

- The BIND Administrator Reference Manual

- RFC 1912 - Common DNS operational and configuration errors

4.6.3. Setting up a reverse zone on a BIND primary server

Reverse zones map IP addresses to names. For example, if you are responsible for IP range **192.0.2.0/24**, you can set up a reverse zone in BIND to resolve IP addresses from this range to hostnames.

> **NOTE**
>
> If you create a reverse zone for whole classful networks, name the zone accordingly. For example, for the class C network **192.0.2.0/24**, the name of the zone is **2.0.192.in-addr.arpa**. If you want to create a reverse zone for a different network size, for example **190.0.2.0/28**, the name of the zone is **28-2.0.192.in-addr.arpa**.

Prerequisites

- BIND is already configured, for example, as a caching name server.

- The **named** or **named-chroot** service is running.

Procedure

1. Add a zone definition to the **/etc/named.conf** file:

   ```
   zone "2.0.192.in-addr.arpa" {
       type master;
       file "2.0.192.in-addr.arpa.zone";
       allow-query { any; };
       allow-transfer { none; };
   };
   ```

   These settings define:

   - This server as the primary server (**type master**) for the **2.0.192.in-addr.arpa** reverse zone.

   - The **/var/named/2.0.192.in-addr.arpa.zone** file is the zone file. If you set a relative path, as in this example, this path is relative to the directory you set in **directory** in the **options** statement.

- Any host can query this zone. Alternatively, specify IP ranges or BIND access control list (ACL) nicknames to limit the access.

- No host can transfer the zone. Allow zone transfers only when you set up secondary servers and only for the IP addresses of the secondary servers.

2. Verify the syntax of the **/etc/named.conf** file:

   > # **named-checkconf**

   If the command displays no output, the syntax is correct.

3. Create the **/var/named/2.0.192.in-addr.arpa.zone** file, for example, with the following content:

   > **$TTL** *8h*
   > **@ IN SOA** *ns1.example.com. hostmaster.example.com.* **(**
   >   *2022070601* **; serial number**
   >   *1d*   **; refresh period**
   >   *3h*   **; retry period**
   >   *3d*   **; expire time**
   >   *3h*) **; minimum TTL**
   >
   >   **IN NS   ns1.example.com.**
   >
   > **1   IN PTR ns1.example.com.**
   > **30   IN PTR www.example.com.**

   This zone file:

   - Sets the default time-to-live (TTL) value for resource records to 8 hours. Without a time suffix, such as **h** for hour, BIND interprets the value as seconds.

   - Contains the required SOA resource record with details about the zone.

   - Sets **ns1.example.com** as an authoritative DNS server for this reverse zone. To be functional, a zone requires at least one name server (**NS**) record. However, to be compliant with RFC 1912, you require at least two name servers.

   - Sets the pointer (**PTR**) record for the **192.0.2.1** and **192.0.2.30** addresses.

4. Set secure permissions on the zone file that only allow the **named** group to read it:

   > # **chown root:named** /var/named/*2.0.192.in-addr.arpa.zone*
   > # **chmod 640** /var/named/*2.0.192.in-addr.arpa.zone*

5. Verify the syntax of the **/var/named/2.0.192.in-addr.arpa.zone** file:

   > # **named-checkzone** *2.0.192.in-addr.arpa /var/named/2.0.192.in-addr.arpa.zone*
   > zone *2.0.192.in-addr.arpa/IN*: loaded serial *2022070601*
   > OK

6. Reload BIND:

   > # **systemctl reload named**

If you run BIND in a change-root environment, use the **systemctl reload named-chroot** command to reload the service.

Verification

- Query different records from the reverse zone, and verify that the output matches the records you have configured in the zone file:

  ```
  # dig +short @localhost -x 192.0.2.1
  ns1.example.com.

  # dig +short @localhost -x 192.0.2.30
  www.example.com.
  ```

  This example assumes that BIND runs on the same host and responds to queries on the **localhost** interface.

Additional resources

- The SOA record in zone files

- Writing BIND ACLs

- The BIND Administrator Reference Manual

- RFC 1912 – Common DNS operational and configuration errors

4.6.4. Updating a BIND zone file

In certain situations, for example if an IP address of a server changes, you must update a zone file. If multiple DNS servers are responsible for a zone, perform this procedure only on the primary server. Other DNS servers that store a copy of the zone will receive the update through a zone transfer.

Prerequisites

- The zone is configured.

- The **named** or **named-chroot** service is running.

Procedure

1. Optional: Identify the path to the zone file in the **/etc/named.conf** file:

   ```
   options {
       ...
       directory       "/var/named";
   }

   zone "example.com" {
       ...
       file "example.com.zone";
   };
   ```

   You find the path to the zone file in the **file** statement in the zone's definition. A relative path is relative to the directory set in **directory** in the **options** statement.

2. Edit the zone file:

   a. Make the required changes.

   b. Increment the serial number in the start of authority (SOA) record.

   > **IMPORTANT**
   >
   > If the serial number is equal to or lower than the previous value, secondary servers will not update their copy of the zone.

3. Verify the syntax of the zone file:

   > # **named-checkzone** *example.com* */var/named/example.com.zone*
   > zone *example.com/IN*: loaded serial *2022062802*
   > OK

4. Reload BIND:

   > # **systemctl reload named**

   If you run BIND in a change-root environment, use the **systemctl reload named-chroot** command to reload the service.

Verification

- Query the record you have added, modified, or removed, for example:

  > # **dig +short** *@localhost A ns2.example.com*
  > 192.0.2.2

  This example assumes that BIND runs on the same host and responds to queries on the **localhost** interface.

Additional resources

- The SOA record in zone files

- Setting up a forward zone on a BIND primary server

- Setting up a reverse zone on a BIND primary server

4.6.5. DNSSEC zone signing using the automated key generation and zone maintenance features

You can sign zones with domain name system security extensions (DNSSEC) to ensure authentication and data integrity. Such zones contain additional resource records. Clients can use them to verify the authenticity of the zone information.

If you enable the DNSSEC policy feature for a zone, BIND performs the following actions automatically:

- Creates the keys

- Signs the zone

- Maintains the zone, including re-signing and periodically replacing the keys.
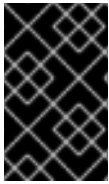
> **IMPORTANT**
>
> To enable external DNS servers to verify the authenticity of a zone, you must add the public key of the zone to the parent zone. Contact your domain provider or registry for further details on how to accomplish this.

This procedure uses the built-in **default** DNSSEC policy in BIND. This policy uses single **ECDSAP256SHA** key signatures. Alternatively, create your own policy to use custom keys, algorithms, and timings.

Prerequisites

- BIND 9.16 or later is installed. To meet this requirement, install the **bind9.16** package instead of **bind**.

- The zone for which you want to enable DNSSEC is configured.

- The **named** or **named-chroot** service is running.

- The server synchronizes the time with a time server. An accurate system time is important for DNSSEC validation.

Procedure

1. Edit the **/etc/named.conf** file, and add **dnssec-policy default;** to the zone for which you want to enable DNSSEC:

   ```
   zone "example.com" {
       ...
       dnssec-policy default;
   };
   ```

2. Reload BIND:

   ```
   # systemctl reload named
   ```

   If you run BIND in a change-root environment, use the **systemctl reload named-chroot** command to reload the service.

3. BIND stores the public key in the **/var/named/K*<zone_name>*.+*<algorithm>*+*<key_ID>*.key** file. Use this file to display the public key of the zone in the format that the parent zone requires:

   - DS record format:

     ```
     # dnssec-dsfromkey /var/named/Kexample.com.+013+61141.key
     example.com. IN DS 61141 13 2
     3E184188CF6D2521EDFDC3F07CFEE8D0195AACBD85E68BAE0620F638B4B1B027
     ```

   - DNSKEY format:

     ```
     # grep DNSKEY /var/named/Kexample.com.+013+61141.key
     example.com. 3600 IN DNSKEY 257 3 13
     ```

> *sjzT3jNEp120aSO4mPEHHSkReHUf7AABNnT8hNRTzD5cKMQSjDJin2I3*
> *5CaKVcWO1pm+HltxUEt+X9dfp8OZkg==*

4. Request to add the public key of the zone to the parent zone. Contact your domain provider or registry for further details on how to accomplish this.

Verification

1. Query your own DNS server for a record from the zone for which you enabled DNSSEC signing:

   > # **dig +dnssec +short @***localhost* **A** *www.example.com*
   > *192.0.2.30*
   > *A 13 3 28800 20220718081258 20220705120353 61141 example.com.*
   > *e7Cfh6GuOBMAWsgsHSVTPh+JJSOl/Y6zctzIuqIU1JqEgOOAfL/Qz474*
   > *M0sgi54m1Kmnr2ANBKJN9uvOs5eXYw==*

   This example assumes that BIND runs on the same host and responds to queries on the **localhost** interface.

2. After the public key has been added to the parent zone and propagated to other servers, verify that the server sets the authenticated data (**ad**) flag on queries to the signed zone:

   > # **dig @***localhost example.com* **+dnssec**
   > ...
   > ;; flags: qr rd ra **ad**; QUERY: 1, ANSWER: 2, AUTHORITY: 0, ADDITIONAL: 1
   > ...

Additional resources

- [Setting up a forward zone on a BIND primary server](#)

- [Setting up a reverse zone on a BIND primary server](#)

4.7. CONFIGURING ZONE TRANSFERS AMONG BIND DNS SERVERS

Zone transfers ensure that all DNS servers that have a copy of the zone use up-to-date data.

Prerequisites

- On the future primary server, the zone for which you want to set up zone transfers is already configured.

- On the future secondary server, BIND is already configured, for example, as a caching name server.

- On both servers, the **named** or **named-chroot** service is running.

Procedure

1. On the existing primary server:

   a. Create a shared key, and append it to the **/etc/named.conf** file:

      > # **tsig-keygen** *example-transfer-key* **| tee -a /etc/named.conf**

```
key "example-transfer-key" {
    algorithm hmac-sha256;
    secret "q7ANbnyliDMuvWgnKOxMLi313JGcTZB5ydMW5CyUGXQ=";
};
```

This command displays the output of the **tsig-keygen** command and automatically appends it to **/etc/named.conf**.

You will require the output of the command later on the secondary server as well.

b. Edit the zone definition in the **/etc/named.conf** file:

    i. In the **allow-transfer** statement, define that servers must provide the key specified in the **example-transfer-key** statement to transfer a zone:

```
zone "example.com" {
    ...
    allow-transfer { key example-transfer-key; };
};
```

Alternatively, use BIND access control list (ACL) nicknames in the **allow-transfer** statement.

    ii. By default, after a zone has been updated, BIND notifies all name servers which have a name server (**NS**) record in this zone. If you do not plan to add an **NS** record for the secondary server to the zone, you can, configure that BIND notifies this server anyway. For that, add the **also-notify** statement with the IP addresses of this secondary server to the zone:

```
zone "example.com" {
    ...
    also-notify { 192.0.2.2; 2001:db8:1::2; };
};
```

c. Verify the syntax of the **/etc/named.conf** file:

```
# named-checkconf
```

If the command displays no output, the syntax is correct.

d. Reload BIND:

```
# systemctl reload named
```

If you run BIND in a change-root environment, use the **systemctl reload named-chroot** command to reload the service.

2. On the future secondary server:

a. Edit the **/etc/named.conf** file as follows:

    i. Add the same key definition as on the primary server:

```
key "example-transfer-key" {
    algorithm hmac-sha256;
    secret "q7ANbnyliDMuvWgnKOxMLi313JGcTZB5ydMW5CyUGXQ=";
```

```
};
```

ii. Add the zone definition to the /**etc**/**named.conf** file:

```
zone "example.com" {
  type slave;
  file "slaves/example.com.zone";
  allow-query { any; };
  allow-transfer { none; };
  masters {
    192.0.2.1 key example-transfer-key;
    2001:db8:1::1 key example-transfer-key;
  };
};
```

These settings state:

- This server is a secondary server (**type slave**) for the **example.com** zone.

- The /**var**/**named**/**slaves**/**example.com.zone** file is the zone file. If you set a relative path, as in this example, this path is relative to the directory you set in **directory** in the **options** statement. To separate zone files for which this server is secondary from primary ones, you can store them, for example, in the /**var**/**named**/**slaves**/ directory.

- Any host can query this zone. Alternatively, specify IP ranges or ACL nicknames to limit the access.

- No host can transfer the zone from this server.

- The IP addresses of the primary server of this zone are **192.0.2.1** and **2001:db8:1::2**. Alternatively, you can specify ACL nicknames. This secondary server will use the key named **example-transfer-key** to authenticate to the primary server.

b. Verify the syntax of the /**etc**/**named.conf** file:

```
# named-checkconf
```

c. Reload BIND:

```
# systemctl reload named
```

If you run BIND in a change-root environment, use the **systemctl reload named-chroot** command to reload the service.

3. Optional: Modify the zone file on the primary server and add an **NS** record for the new secondary server.

Verification

On the secondary server:

1. Display the **systemd** journal entries of the **named** service:

```
# journalctl -u named
...
```

```
Jul 06 15:08:51 ns2.example.com named[2024]: zone example.com/IN: Transfer started.
Jul 06 15:08:51 ns2.example.com named[2024]: transfer of 'example.com/IN' from
192.0.2.1#53: connected using 192.0.2.2#45803
Jul 06 15:08:51 ns2.example.com named[2024]: zone example.com/IN: transferred serial
2022070101
Jul 06 15:08:51 ns2.example.com named[2024]: transfer of 'example.com/IN' from
192.0.2.1#53: Transfer status: success
Jul 06 15:08:51 ns2.example.com named[2024]: transfer of 'example.com/IN' from
192.0.2.1#53: Transfer completed: 1 messages, 29 records, 2002 bytes, 0.003 secs (667333
bytes/sec)
```

If you run BIND in a change-root environment, use the **journalctl -u named-chroot** command to display the journal entries.

2. Verify that BIND created the zone file:

```
# ls -l /var/named/slaves/
total 4
-rw-r--r--. 1 named named 2736 Jul  6 15:08 example.com.zone
```

Note that, by default, secondary servers store zone files in a binary raw format.

3. Query a record of the transferred zone from the secondary server:

```
# dig +short @192.0.2.2 AAAA www.example.com
2001:db8:1::30
```

This example assumes that the secondary server you set up in this procedure listens on IP address **192.0.2.2**.

Additional resources

- Setting up a forward zone on a BIND primary server

- Setting up a reverse zone on a BIND primary server

- Writing BIND ACLs

- Updating a BIND zone file

4.8. CONFIGURING RESPONSE POLICY ZONES IN BIND TO OVERRIDE DNS RECORDS

Using DNS blocking and filtering, administrators can rewrite a DNS response to block access to certain domains or hosts. In BIND, response policy zones (RPZs) provide this feature. You can configure different actions for blocked entries, such as returning an **NXDOMAIN** error or not responding to the query.

If you have multiple DNS servers in your environment, use this procedure to configure the RPZ on the primary server, and later configure zone transfers to make the RPZ available on your secondary servers.

Prerequisites

- BIND is already configured, for example, as a caching name server.

- The **named** or **named-chroot** service is running.

Procedure

1. Edit the **/etc/named.conf** file, and make the following changes:

   a. Add a **response-policy** definition to the **options** statement:

   ```
   options {
       ...

       response-policy {
           zone "rpz.local";
       };

       ...
   }
   ```

   You can set a custom name for the RPZ in the **zone** statement in **response-policy**. However, you must use the same name in the zone definition in the next step.

   b. Add a **zone** definition for the RPZ you set in the previous step:

   ```
   zone "rpz.local" {
       type master;
       file "rpz.local";
       allow-query { localhost; 192.0.2.0/24; 2001:db8:1::/64; };
       allow-transfer { none; };
   };
   ```

   These settings state:

   - This server is the primary server (**type master**) for the RPZ named **rpz.local**.

   - The **/var/named/rpz.local** file is the zone file. If you set a relative path, as in this example, this path is relative to the directory you set in **directory** in the **options** statement.

   - Any hosts defined in **allow-query** can query this RPZ. Alternatively, specify IP ranges or BIND access control list (ACL) nicknames to limit the access.

   - No host can transfer the zone. Allow zone transfers only when you set up secondary servers and only for the IP addresses of the secondary servers.

2. Verify the syntax of the **/etc/named.conf** file:

   ```
   # named-checkconf
   ```

   If the command displays no output, the syntax is correct.

3. Create the **/var/named/rpz.local** file, for example, with the following content:

   ```
   $TTL 10m
   @ IN SOA ns1.example.com. hostmaster.example.com. (
               2022070601 ; serial number
               1h         ; refresh period
   ```

```
                    1m       ; retry period
                    3d       ; expire time
                    1m )      ; minimum TTL


             IN NS   ns1.example.com.


example.org     IN CNAME .
*.example.org   IN CNAME .
example.net     IN CNAME rpz-drop.
*.example.net   IN CNAME rpz-drop.
```

This zone file:

- Sets the default time-to-live (TTL) value for resource records to 10 minutes. Without a time suffix, such as **h** for hour, BIND interprets the value as seconds.

- Contains the required start of authority (SOA) resource record with details about the zone.

- Sets **ns1.example.com** as an authoritative DNS server for this zone. To be functional, a zone requires at least one name server (**NS**) record. However, to be compliant with RFC 1912, you require at least two name servers.

- Return an **NXDOMAIN** error for queries to **example.org** and hosts in this domain.

- Drop queries to **example.net** and hosts in this domain.

For a full list of actions and examples, see IETF draft: DNS Response Policy Zones (RPZ) .

4. Verify the syntax of the **/var/named/rpz.local** file:

   ```
   # named-checkzone rpz.local /var/named/rpz.local
   zone rpz.local/IN: loaded serial 2022070601
   OK
   ```

5. Reload BIND:

   ```
   # systemctl reload named
   ```

   If you run BIND in a change-root environment, use the **systemctl reload named-chroot** command to reload the service.

Verification

1. Attempt to resolve a host in **example.org**, that is configured in the RPZ to return an **NXDOMAIN** error:

   ```
   # dig @localhost www.example.org
   ...
   ;; ->>HEADER<<- opcode: QUERY, status: NXDOMAIN, id: 30286
   ...
   ```

   This example assumes that BIND runs on the same host and responds to queries on the **localhost** interface.

2. Attempt to resolve a host in the **example.net** domain, that is configured in the RPZ to drop queries:

> # **dig @localhost www.example.net**
> ...
> ;; connection timed out; no servers could be reached
> ...

Additional resources

- IETF draft: DNS Response Policy Zones (RPZ)

CHAPTER 5. EXPORTING NFS SHARES

As a system administrator, you can use the NFS server to share a directory on your system over network.

5.1. INTRODUCTION TO NFS

This section explains the basic concepts of the NFS service.

A Network File System (NFS) allows remote hosts to mount file systems over a network and interact with those file systems as though they are mounted locally. This enables you to consolidate resources onto centralized servers on the network.

The NFS server refers to the **/etc/exports** configuration file to determine whether the client is allowed to access any exported file systems. Once verified, all file and directory operations are available to the user.

5.2. SUPPORTED NFS VERSIONS

This section lists versions of NFS supported in Red Hat Enterprise Linux and their features.

Currently, Red Hat Enterprise Linux 8 supports the following major versions of NFS:

- NFS version 3 (NFSv3) supports safe asynchronous writes and is more robust at error handling than the previous NFSv2; it also supports 64-bit file sizes and offsets, allowing clients to access more than 2 GB of file data.

- NFS version 4 (NFSv4) works through firewalls and on the Internet, no longer requires an **rpcbind** service, supports Access Control Lists (ACLs), and utilizes stateful operations.

NFS version 2 (NFSv2) is no longer supported by Red Hat.

Default NFS version
The default NFS version in Red Hat Enterprise Linux 8 is 4.2. NFS clients attempt to mount using NFSv4.2 by default, and fall back to NFSv4.1 when the server does not support NFSv4.2. The mount later falls back to NFSv4.0 and then to NFSv3.

Features of minor NFS versions
Following are the features of NFSv4.2 in Red Hat Enterprise Linux 8:

Server-side copy

Enables the NFS client to efficiently copy data without wasting network resources using the **copy_file_range()** system call.

Sparse files

Enables files to have one or more *holes*, which are unallocated or uninitialized data blocks consisting only of zeroes. The **lseek()** operation in NFSv4.2 supports **seek_hole()** and **seek_data()**, which enables applications to map out the location of holes in the sparse file.

Space reservation

Permits storage servers to reserve free space, which prohibits servers to run out of space. NFSv4.2 supports the **allocate()** operation to reserve space, the **deallocate()** operation to unreserve space, and the **fallocate()** operation to preallocate or deallocate space in a file.

Labeled NFS

Enforces data access rights and enables SELinux labels between a client and a server for individual files on an NFS file system.

Layout enhancements

Provides the **layoutstats()** operation, which enables some Parallel NFS (pNFS) servers to collect better performance statistics.

Following are the features of NFSv4.1:

- Enhances performance and security of network, and also includes client-side support for pNFS.

- No longer requires a separate TCP connection for callbacks, which allows an NFS server to grant delegations even when it cannot contact the client: for example, when NAT or a firewall interferes.

- Provides exactly once semantics (except for reboot operations), preventing a previous issue whereby certain operations sometimes returned an inaccurate result if a reply was lost and the operation was sent twice.

5.3. THE TCP AND UDP PROTOCOLS IN NFSV3 AND NFSV4

NFSv4 requires the Transmission Control Protocol (TCP) running over an IP network.

NFSv3 could also use the User Datagram Protocol (UDP) in earlier Red Hat Enterprise Linux versions. In Red Hat Enterprise Linux 8, NFS over UDP is no longer supported. By default, UDP is disabled in the NFS server.

5.4. SERVICES REQUIRED BY NFS

This section lists system services that are required for running an NFS server or mounting NFS shares. Red Hat Enterprise Linux starts these services automatically.

Red Hat Enterprise Linux uses a combination of kernel-level support and service processes to provide NFS file sharing. All NFS versions rely on Remote Procedure Calls (RPC) between clients and servers. To share or mount NFS file systems, the following services work together depending on which version of NFS is implemented:

nfsd

The NFS server kernel module that services requests for shared NFS file systems.

rpcbind

Accepts port reservations from local RPC services. These ports are then made available (or advertised) so the corresponding remote RPC services can access them. The **rpcbind** service responds to requests for RPC services and sets up connections to the requested RPC service. This is not used with NFSv4.

rpc.mountd

This process is used by an NFS server to process **MOUNT** requests from NFSv3 clients. It checks that the requested NFS share is currently exported by the NFS server, and that the client is allowed to access it. If the mount request is allowed, the **nfs-mountd** service replies with a Success status and provides the File-Handle for this NFS share back to the NFS client.

rpc.nfsd

This process enables explicit NFS versions and protocols the server advertises to be defined. It works with the Linux kernel to meet the dynamic demands of NFS clients, such as providing server threads each time an NFS client connects. This process corresponds to the **nfs-server** service.

lockd

This is a kernel thread that runs on both clients and servers. It implements the Network Lock Manager (NLM) protocol, which enables NFSv3 clients to lock files on the server. It is started automatically whenever the NFS server is run and whenever an NFS file system is mounted.

rpc.statd

This process implements the Network Status Monitor (NSM) RPC protocol, which notifies NFS clients when an NFS server is restarted without being gracefully brought down. The **rpc-statd** service is started automatically by the **nfs-server** service, and does not require user configuration. This is not used with NFSv4.

rpc.rquotad

This process provides user quota information for remote users. The **rpc-rquotad** service, which is provided by the **quota-rpc** package, has to be started by user when the **nfs-server** is started.

rpc.idmapd

This process provides NFSv4 client and server upcalls, which map between on-the-wire NFSv4 names (strings in the form of *user@domain*) and local UIDs and GIDs. For **idmapd** to function with NFSv4, the **/etc/idmapd.conf** file must be configured. At a minimum, the **Domain** parameter should be specified, which defines the NFSv4 mapping domain. If the NFSv4 mapping domain is the same as the DNS domain name, this parameter can be skipped. The client and server must agree on the NFSv4 mapping domain for ID mapping to function properly.

Only the NFSv4 server uses **rpc.idmapd**, which is started by the **nfs-idmapd** service. The NFSv4 client uses the keyring-based **nfsidmap** utility, which is called by the kernel on-demand to perform ID mapping. If there is a problem with **nfsidmap**, the client falls back to using **rpc.idmapd**.

The RPC services with NFSv4

The mounting and locking protocols have been incorporated into the NFSv4 protocol. The server also listens on the well-known TCP port 2049. As such, NFSv4 does not need to interact with **rpcbind**, **lockd**, and **rpc-statd** services. The **nfs-mountd** service is still required on the NFS server to set up the exports, but is not involved in any over-the-wire operations.

Additional resources

- Configuring an NFSv4 only server without **rpcbind**.

5.5. NFS HOST NAME FORMATS

This section describes different formats that you can use to specify a host when mounting or exporting an NFS share.

You can specify the host in the following formats:

Single machine

Either of the following:

- A fully-qualified domain name (that can be resolved by the server)

- Host name (that can be resolved by the server)

- An IP address.

IP networks

Either of the following formats is valid:

- *a.b.c.d/z*, where *a.b.c.d* is the network and *z* is the number of bits in the netmask; for

example **192.168.0.0/24**.

- **a.b.c.d/netmask**, where **a.b.c.d** is the network and **netmask** is the netmask; for example, **192.168.100.8/255.255.255.0**.

Netgroups

The **@group-name** format , where **group-name** is the NIS netgroup name.

5.6. NFS SERVER CONFIGURATION

This section describes the syntax and options of two ways to configure exports on an NFS server:

- Manually editing the **/etc/exports** configuration file

- Using the **exportfs** utility on the command line

5.6.1. The /etc/exports configuration file

The **/etc/exports** file controls which file systems are exported to remote hosts and specifies options. It follows the following syntax rules:

- Blank lines are ignored.

- To add a comment, start a line with the hash mark (**#**).

- You can wrap long lines with a backslash (\).

- Each exported file system should be on its own individual line.

- Any lists of authorized hosts placed after an exported file system must be separated by space characters.

- Options for each of the hosts must be placed in parentheses directly after the host identifier, without any spaces separating the host and the first parenthesis.

Export entry
Each entry for an exported file system has the following structure:

    export host(options)

It is also possible to specify multiple hosts, along with specific options for each host. To do so, list them on the same line as a space-delimited list, with each host name followed by its respective options (in parentheses), as in:

    export host1(options1) host2(options2) host3(options3)

In this structure:

export

The directory being exported

host

The host or network to which the export is being shared

options

The options to be used for host

> **Example 5.1. A simple /etc/exports file**
>
> In its simplest form, the **/etc/exports** file only specifies the exported directory and the hosts permitted to access it:
>
> ```
> /exported/directory bob.example.com
> ```
>
> Here, **bob.example.com** can mount **/exported/directory/** from the NFS server. Because no options are specified in this example, NFS uses default options.

> **IMPORTANT**
>
> The format of the **/etc/exports** file is very precise, particularly in regards to use of the space character. Remember to always separate exported file systems from hosts and hosts from one another with a space character. However, there should be no other space characters in the file except on comment lines.
>
> For example, the following two lines do not mean the same thing:
>
> ```
> /home bob.example.com(rw)
> /home bob.example.com (rw)
> ```
>
> The first line allows only users from **bob.example.com** read and write access to the /**home** directory. The second line allows users from   **bob.example.com** to mount the directory as read-only (the default), while the rest of the world can mount it read/write.

Default options
The default options for an export entry are:

ro

The exported file system is read-only. Remote hosts cannot change the data shared on the file system. To allow hosts to make changes to the file system (that is, read and write), specify the rw option.

sync

The NFS server will not reply to requests before changes made by previous requests are written to disk. To enable asynchronous writes instead, specify the option **async**.

wdelay

The NFS server will delay writing to the disk if it suspects another write request is imminent. This can improve performance as it reduces the number of times the disk must be accessed by separate write commands, thereby reducing write overhead. To disable this, specify the **no_wdelay** option, which is available only if the default sync option is also specified.

root_squash

This prevents root users connected remotely (as opposed to locally) from having root privileges; instead, the NFS server assigns them the user ID **nobody**. This effectively "squashes" the power of the remote root user to the lowest local user, preventing possible unauthorized writes on the remote server. To disable root squashing, specify the **no_root_squash** option.
To squash every remote user (including root), use the **all_squash** option. To specify the user and group IDs that the NFS server should assign to remote users from a particular host, use the **anonuid** and **anongid** options, respectively, as in:

> *export host*(anonuid=*uid*,anongid=*gid*)

Here, *uid* and *gid* are user ID number and group ID number, respectively. The **anonuid** and **anongid** options enable you to create a special user and group account for remote NFS users to share.

By default, access control lists (ACLs) are supported by NFS under Red Hat Enterprise Linux. To disable this feature, specify the **no_acl** option when exporting the file system.

Default and overridden options

Each default for every exported file system must be explicitly overridden. For example, if the **rw** option is not specified, then the exported file system is shared as read-only. The following is a sample line from **/etc/exports** which overrides two default options:

> /another/exported/directory 192.168.0.3(rw,async)

In this example, **192.168.0.3** can mount **/another/exported/directory/** read and write, and all writes to disk are asynchronous.

5.6.2. The exportfs utility

The **exportfs** utility enables the root user to selectively export or unexport directories without restarting the NFS service. When given the proper options, the **exportfs** utility writes the exported file systems to **/var/lib/nfs/xtab**. Because the **nfs-mountd** service refers to the **xtab** file when deciding access privileges to a file system, changes to the list of exported file systems take effect immediately.

Common exportfs options

The following is a list of commonly-used options available for **exportfs**:

-r

Causes all directories listed in **/etc/exports** to be exported by constructing a new export list in **/var/lib/nfs/etab**. This option effectively refreshes the export list with any changes made to **/etc/exports**.

-a

Causes all directories to be exported or unexported, depending on what other options are passed to **exportfs**. If no other options are specified, **exportfs** exports all file systems specified in **/etc/exports**.

-o *file-systems*

Specifies directories to be exported that are not listed in **/etc/exports**. Replace *file-systems* with additional file systems to be exported. These file systems must be formatted in the same way they are specified in **/etc/exports**. This option is often used to test an exported file system before adding it permanently to the list of exported file systems.

-i

Ignores **/etc/exports**; only options given from the command line are used to define exported file systems.

-u

Unexports all shared directories. The command **exportfs -ua** suspends NFS file sharing while keeping all NFS services up. To re-enable NFS sharing, use **exportfs -r**.

-v

Verbose operation, where the file systems being exported or unexported are displayed in greater detail when the **exportfs** command is executed.

If no options are passed to the **exportfs** utility, it displays a list of currently exported file systems.

Additional resources

- [NFS host name formats](#) .

5.7. NFS AND RPCBIND

This section explains the purpose of the **rpcbind** service, which is required by NFSv3.

The **rpcbind** service maps Remote Procedure Call (RPC) services to the ports on which they listen. RPC processes notify **rpcbind** when they start, registering the ports they are listening on and the RPC program numbers they expect to serve. The client system then contacts **rpcbind** on the server with a particular RPC program number. The **rpcbind** service redirects the client to the proper port number so it can communicate with the requested service.

Because RPC-based services rely on **rpcbind** to make all connections with incoming client requests, **rpcbind** must be available before any of these services start.

Access control rules for **rpcbind** affect all RPC-based services. Alternatively, it is possible to specify access control rules for each of the NFS RPC daemons.

Additional resources

- **rpc.mountd(8)** man page

- **rpc.statd(8)** man page

5.8. INSTALLING NFS

This procedure installs all packages necessary to mount or export NFS shares.

Procedure

- Install the **nfs-utils** package:

  ```
  # yum install nfs-utils
  ```

5.9. STARTING THE NFS SERVER

This procedure describes how to start the NFS server, which is required to export NFS shares.

Prerequisites

- For servers that support NFSv3 connections, the **rpcbind** service must be running. To verify that **rpcbind** is active, use the following command:

  ```
  $ systemctl status rpcbind
  ```

  If the service is stopped, start and enable it:

  ```
  $ systemctl enable --now rpcbind
  ```

Procedure

- To start the NFS server and enable it to start automatically at boot, use the following command:

  ```
  # systemctl enable --now nfs-server
  ```

Additional resources

- Configuring an NFSv4-only server.

5.10. TROUBLESHOOTING NFS AND RPCBIND

Because the **rpcbind** service provides coordination between RPC services and the port numbers used to communicate with them, it is useful to view the status of current RPC services using **rpcbind** when troubleshooting. The **rpcinfo** utility shows each RPC-based service with port numbers, an RPC program number, a version number, and an IP protocol type (TCP or UDP).

Procedure

1. To make sure the proper NFS RPC-based services are enabled for **rpcbind**, use the following command:

   ```
   # rpcinfo -p
   ```

   Example 5.2. rpcinfo -p command output

   The following is sample output from this command:

   ```
   program vers proto   port  service
    100000   4   tcp    111  portmapper
    100000   3   tcp    111  portmapper
    100000   2   tcp    111  portmapper
    100000   4   udp    111  portmapper
    100000   3   udp    111  portmapper
    100000   2   udp    111  portmapper
    100005   1   udp  20048  mountd
    100005   1   tcp  20048  mountd
    100005   2   udp  20048  mountd
    100005   2   tcp  20048  mountd
    100005   3   udp  20048  mountd
    100005   3   tcp  20048  mountd
    100024   1   udp  37769  status
    100024   1   tcp  49349  status
    100003   3   tcp   2049  nfs
    100003   4   tcp   2049  nfs
    100227   3   tcp   2049  nfs_acl
    100021   1   udp  56691  nlockmgr
    100021   3   udp  56691  nlockmgr
    100021   4   udp  56691  nlockmgr
    100021   1   tcp  46193  nlockmgr
    100021   3   tcp  46193  nlockmgr
    100021   4   tcp  46193  nlockmgr
   ```

If one of the NFS services does not start up correctly, **rpcbind** will be unable to map RPC requests from clients for that service to the correct port.

2. In many cases, if NFS is not present in **rpcinfo** output, restarting NFS causes the service to correctly register with **rpcbind** and begin working:

```
# systemctl restart nfs-server
```

Additional resources

- Configuring an NFSv4-only server.

5.11. CONFIGURING THE NFS SERVER TO RUN BEHIND A FIREWALL

NFS requires the **rpcbind** service, which dynamically assigns ports for RPC services and can cause issues for configuring firewall rules. The following sections describe how to configure NFS versions to work behind a firewall if you want to support:

- NFSv3
  This includes any servers that support NFSv3:

  - NFSv3-only servers

  - Servers that support both NFSv3 and NFSv4

- NFSv4-only

5.11.1. Configuring the NFSv3-enabled server to run behind a firewall

The following procedure describes how to configure servers that support NFSv3 to run behind a firewall. This includes NFSv3-only servers and servers that support both NFSv3 and NFSv4.

Procedure

1. To allow clients to access NFS shares behind a firewall, configure the firewall by running the following commands on the NFS server:

   ```
   firewall-cmd --permanent --add-service mountd
   firewall-cmd --permanent --add-service rpc-bind
   firewall-cmd --permanent --add-service nfs
   ```

2. Specify the ports to be used by the RPC service **nlockmgr** in the **/etc/nfs.conf** file as follows:

   ```
   [lockd]

   port=tcp-port-number
   udp-port=udp-port-number
   ```

   Alternatively, you can specify **nlm_tcpport** and **nlm_udpport** in the **/etc/modprobe.d/lockd.conf** file.

3. Open the specified ports in the firewall by running the following commands on the NFS server:

```
firewall-cmd --permanent --add-port=<lockd-tcp-port>/tcp
firewall-cmd --permanent --add-port=<lockd-udp-port>/udp
```

4. Add static ports for **rpc.statd** by editing the **[statd]** section of the **/etc/nfs.conf** file as follows:

```
[statd]

port=port-number
```

5. Open the added ports in the firewall by running the following commands on the NFS server:

```
firewall-cmd --permanent --add-port=<statd-tcp-port>/tcp
firewall-cmd --permanent --add-port=<statd-udp-port>/udp
```

6. Reload the firewall configuration:

```
firewall-cmd --reload
```

7. Restart the **rpc-statd** service first, and then restart the **nfs-server** service:

```
# systemctl restart rpc-statd.service
# systemctl restart nfs-server.service
```

Alternatively, if you specified the **lockd** ports in the **/etc/modprobe.d/lockd.conf** file:

a. Update the current values of **/proc/sys/fs/nfs/nlm_tcpport** and **/proc/sys/fs/nfs/nlm_udpport**:

```
# sysctl -w fs.nfs.nlm_tcpport=<tcp-port>
# sysctl -w fs.nfs.nlm_udpport=<udp-port>
```

b. Restart the **rpc-statd** and **nfs-server** services:

```
# systemctl restart rpc-statd.service
# systemctl restart nfs-server.service
```

5.11.2. Configuring the NFSv4-only server to run behind a firewall

The following procedure describes how to configure the NFSv4-only server to run behind a firewall.

Procedure

1. To allow clients to access NFS shares behind a firewall, configure the firewall by running the following command on the NFS server:

```
firewall-cmd --permanent --add-service nfs
```

2. Reload the firewall configuration:

```
firewall-cmd --reload
```

3. Restart the nfs-server:

> # **systemctl restart nfs-server**

5.11.3. Configuring an NFSv3 client to run behind a firewall

The procedure to configure an NFSv3 client to run behind a firewall is similar to the procedure to configure an NFSv3 server to run behind a firewall.

If the machine you are configuring is both an NFS client and an NFS server, follow the procedure described in Configuring the NFSv3-enabled server to run behind a firewall .

The following procedure describes how to configure a machine that is an NFS client only to run behind a firewall.

Procedure

1. To allow the NFS server to perform callbacks to the NFS client when the client is behind a firewall, add the **rpc-bind** service to the firewall by running the following command on the NFS client:

   > **firewall-cmd --permanent --add-service rpc-bind**

2. Specify the ports to be used by the RPC service **nlockmgr** in the /**etc**/**nfs.conf** file as follows:

   > [lockd]
   >
   > port=*port-number*
   > udp-port=*upd-port-number*

   Alternatively, you can specify **nlm_tcpport** and **nlm_udpport** in the /**etc**/**modprobe.d**/**lockd.conf** file.

3. Open the specified ports in the firewall by running the following commands on the NFS client:

   > **firewall-cmd --permanent --add-port=*<lockd-tcp-port>*/tcp**
   > **firewall-cmd --permanent --add-port=*<lockd-udp-port>*/udp**

4. Add static ports for **rpc.statd** by editing the **[statd]** section of the /**etc**/**nfs.conf** file as follows:

   > [statd]
   >
   > port=*port-number*

5. Open the added ports in the firewall by running the following commands on the NFS client:

   > **firewall-cmd --permanent --add-port=*<statd-tcp-port>*/tcp**
   > **firewall-cmd --permanent --add-port=*<statd-udp-port>*/udp**

6. Reload the firewall configuration:

   > **firewall-cmd --reload**

7. Restart the **rpc-statd** service:

```
# systemctl restart rpc-statd.service
```

Alternatively, if you specified the **lockd** ports in the **/etc/modprobe.d/lockd.conf** file:

a. Update the current values of **/proc/sys/fs/nfs/nlm_tcpport** and
   /proc/sys/fs/nfs/nlm_udpport:

   ```
   # sysctl -w fs.nfs.nlm_tcpport=<tcp-port>
   # sysctl -w fs.nfs.nlm_udpport=<udp-port>
   ```

b. Restart the **rpc-statd** service:

   ```
   # systemctl restart rpc-statd.service
   ```

5.11.4. Configuring an NFSv4 client to run behind a firewall

Perform this procedure only if the client is using NFSv4.0. In that case, it is necessary to open a port for NFSv4.0 callbacks.

This procedure is not needed for NFSv4.1 or higher because in the later protocol versions the server performs callbacks on the same connection that was initiated by the client.

Procedure

1. To allow NFSv4.0 callbacks to pass through firewalls, set
   /proc/sys/fs/nfs/nfs_callback_tcpport and allow the server to connect to that port on the
   client as follows:

   ```
   # echo "fs.nfs.nfs_callback_tcpport = <callback-port>" >/etc/sysctl.d/90-nfs-callback-port.conf
   # sysctl -p /etc/sysctl.d/90-nfs-callback-port.conf
   ```

2. Open the specified port in the firewall by running the following command on the NFS client:

   ```
   firewall-cmd --permanent --add-port=<callback-port>/tcp
   ```

3. Reload the firewall configuration:

   ```
   firewall-cmd --reload
   ```

5.12. EXPORTING RPC QUOTA THROUGH A FIREWALL

If you export a file system that uses disk quotas, you can use the quota Remote Procedure Call (RPC) service to provide disk quota data to NFS clients.

Procedure

1. Enable and start the **rpc-rquotad** service:

   ```
   # systemctl enable --now rpc-rquotad
   ```

> **NOTE**
>
> The **rpc-rquotad** service is, if enabled, started automatically after starting the nfs-server service.

2. To make the quota RPC service accessible behind a firewall, the TCP (or UDP, if UDP is enabled) port 875 need to be open. The default port number is defined in the **/etc/services** file. You can override the default port number by appending **-p port-number** to the **RPCRQUOTADOPTS** variable in the **/etc/sysconfig/rpc-rquotad** file.

3. By default, remote hosts can only read quotas. If you want to allow clients to set quotas, append the **-S** option to the **RPCRQUOTADOPTS** variable in the **/etc/sysconfig/rpc-rquotad** file.

4. Restart **rpc-rquotad** for the changes in the **/etc/sysconfig/rpc-rquotad** file to take effect:

   ```
   # systemctl restart rpc-rquotad
   ```

5.13. ENABLING NFS OVER RDMA (NFSORDMA)

Remote direct memory access (RDMA) service works automatically on RDMA-capable hardware in Red Hat Enterprise Linux 8.

Procedure

1. Install the **rdma-core** package:

   ```
   # yum install rdma-core
   ```

2. Verify the lines with **xprtrdma** and **svcrdma** are commented out in the **/etc/rdma/modules/rdma.conf** file:

   ```
   # NFS over RDMA client support
   xprtrdma
   # NFS over RDMA server support
   svcrdma
   ```

3. On the NFS server, create directory **/mnt/nfsordma** and export it to **/etc/exports**:

   ```
   # mkdir /mnt/nfsordma
   # echo "/mnt/nfsordma *(fsid=0,rw,async,insecure,no_root_squash)" >> /etc/exports
   ```

4. On the NFS client, mount the nfs-share with server IP address, for example, **172.31.0.186**:

   ```
   # mount -o rdma,port=20049 172.31.0.186:/mnt/nfs-share /mnt/nfs
   ```

5. Restart the **nfs-server** service:

   ```
   # systemctl restart nfs-server
   ```

Additional resources

- [The RFC 5667 standard](#)

5.14. ADDITIONAL RESOURCES

- The Linux NFS wiki

CHAPTER 6. SECURING NFS

To minimize NFS security risks and protect data on the server, consider the following sections when exporting NFS file systems on a server or mounting them on a client.

6.1. NFS SECURITY WITH AUTH_SYS AND EXPORT CONTROLS

NFS provides the following traditional options in order to control access to exported files:

- The server restricts which hosts are allowed to mount which file systems either by IP address or by host name.

- The server enforces file system permissions for users on NFS clients in the same way it does for local users. Traditionally, NFS does this using the **AUTH_SYS** call message (also called **AUTH_UNIX**), which relies on the client to state the UID and GIDs of the user. Be aware that this means that a malicious or misconfigured client might easily get this wrong and allow a user access to files that it should not.

To limit the potential risks, administrators often limits the access to read-only or squash user permissions to a common user and group ID. Unfortunately, these solutions prevent the NFS share from being used in the way it was originally intended.

Additionally, if an attacker gains control of the DNS server used by the system exporting the NFS file system, they can point the system associated with a particular hostname or fully qualified domain name to an unauthorized machine. At this point, the unauthorized machine *is* the system permitted to mount the NFS share, because no username or password information is exchanged to provide additional security for the NFS mount.

Wildcards should be used sparingly when exporting directories through NFS, as it is possible for the scope of the wildcard to encompass more systems than intended.

Additional resources

- To secure NFS and **rpcbind**, use, for example, **nftables** and **firewalld**.

- **nft(8)** man page

- **firewalld-cmd(1)** man page

6.2. NFS SECURITY WITH AUTH_GSS

All version of NFS support RPCSEC_GSS and the Kerberos mechanism.

Unlike AUTH_SYS, with the RPCSEC_GSS Kerberos mechanism, the server does not depend on the client to correctly represent which user is accessing the file. Instead, cryptography is used to authenticate users to the server, which prevents a malicious client from impersonating a user without having that user's Kerberos credentials. Using the RPCSEC_GSS Kerberos mechanism is the most straightforward way to secure mounts because after configuring Kerberos, no additional setup is needed.

6.3. CONFIGURING AN NFS SERVER AND CLIENT TO USE KERBEROS

Kerberos is a network authentication system that allows clients and servers to authenticate to each other by using symmetric encryption and a trusted third party, the KDC. Red Hat recommends using Identity Management (IdM) for setting up Kerberos.

Prerequisites

- The Kerberos Key Distribution Centre (**KDC**) is installed and configured.

Procedure

1. - Create the **nfs**/**hostname.*domain@REALM*** principal on the NFS server side.

   - Create the **host**/**hostname.*domain@REALM*** principal on both the server and the client side.

   - Add the corresponding keys to keytabs for the client and server.

2. On the server side, use the **sec=** option to enable the wanted security flavors. To enable all security flavors as well as non-cryptographic mounts:

   > /export *(sec=sys:krb5:krb5i:krb5p)

   Valid security flavors to use with the **sec=** option are:

   - **sys**: no cryptographic protection, the default

   - **krb5**: authentication only

   - **krb5i**: integrity protection

     - uses Kerberos V5 for user authentication and performs integrity checking of NFS operations using secure checksums to prevent data tampering.

   - **krb5p**: privacy protection

     - uses Kerberos V5 for user authentication, integrity checking, and encrypts NFS traffic to prevent traffic sniffing. This is the most secure setting, but it also involves the most performance overhead.

3. On the client side, add **sec=krb5** (or **sec=krb5i**, or **sec=krb5p**, depending on the setup) to the mount options:

   > # mount -o sec=krb5 server:/export /mnt

Additional resources

- [Creating files as root on krb5-secured NFS](#) . Not recommended.

- **exports(5)** man page

- **nfs(5)** man page

6.4. NFSV4 SECURITY OPTIONS

NFSv4 includes ACL support based on the Microsoft Windows NT model, not the POSIX model, because of the Microsoft Windows NT model's features and wide deployment.

Another important security feature of NFSv4 is the removal of the use of the **MOUNT** protocol for mounting file systems. The **MOUNT** protocol presented a security risk because of the way the protocol processed file handles.

6.5. FILE PERMISSIONS ON MOUNTED NFS EXPORTS

Once the NFS file system is mounted as either read or read and write by a remote host, the only protection each shared file has is its permissions. If two users that share the same user ID value mount the same NFS file system on different client systems, they can modify each others' files. Additionally, anyone logged in as root on the client system can use the **su -** command to access any files with the NFS share.

By default, access control lists (ACLs) are supported by NFS under Red Hat Enterprise Linux. Red Hat recommends to keep this feature enabled.

By default, NFS uses *root squashing* when exporting a file system. This sets the user ID of anyone accessing the NFS share as the root user on their local machine to **nobody**. Root squashing is controlled by the default option **root_squash**; for more information about this option, see NFS server configuration.

When exporting an NFS share as read-only, consider using the **all_squash** option. This option makes every user accessing the exported file system take the user ID of the **nobody** user.

CHAPTER 7. ENABLING PNFS SCSI LAYOUTS IN NFS

You can configure the NFS server and client to use the pNFS SCSI layout for accessing data. pNFS SCSI is beneficial in use cases that involve longer-duration single-client access to a file.

Prerequisites

- Both the client and the server must be able to send SCSI commands to the same block device. That is, the block device must be on a shared SCSI bus.

- The block device must contain an XFS file system.

- The SCSI device must support SCSI Persistent Reservations as described in the SCSI-3 Primary Commands specification.

7.1. THE PNFS TECHNOLOGY

The pNFS architecture improves the scalability of NFS. When a server implements pNFS, the client is able to access data through multiple servers concurrently. This can lead to performance improvements.

pNFS supports the following storage protocols or layouts on RHEL:

- Files

- Flexfiles

- SCSI

7.2. PNFS SCSI LAYOUTS

The SCSI layout builds on the work of pNFS block layouts. The layout is defined across SCSI devices. It contains a sequential series of fixed-size blocks as logical units (LUs) that must be capable of supporting SCSI persistent reservations. The LU devices are identified by their SCSI device identification.

pNFS SCSI performs well in use cases that involve longer-duration single-client access to a file. An example might be a mail server or a virtual machine housing a cluster.

Operations between the client and the server
When an NFS client reads from a file or writes to it, the client performs a **LAYOUTGET** operation. The server responds with the location of the file on the SCSI device. The client might need to perform an additional operation of **GETDEVICEINFO** to determine which SCSI device to use. If these operations work correctly, the client can issue I/O requests directly to the SCSI device instead of sending **READ** and **WRITE** operations to the server.

Errors or contention between clients might cause the server to recall layouts or not issue them to the clients. In those cases, the clients fall back to issuing **READ** and **WRITE** operations to the server instead of sending I/O requests directly to the SCSI device.

To monitor the operations, see Monitoring pNFS SCSI layouts functionality.

Device reservations
pNFS SCSI handles fencing through the assignment of reservations. Before the server issues layouts to clients, it reserves the SCSI device to ensure that only registered clients may access the device. If a client can issue commands to that SCSI device but is not registered with the device, many operations

from the client on that device fail. For example, the **blkid** command on the client fails to show the UUID of the XFS file system if the server has not given a layout for that device to the client.

The server does not remove its own persistent reservation. This protects the data within the file system on the device across restarts of clients and servers. In order to repurpose the SCSI device, you might need to manually remove the persistent reservation on the NFS server.

7.3. CHECKING FOR A SCSI DEVICE COMPATIBLE WITH PNFS

This procedure checks if a SCSI device supports the pNFS SCSI layout.

Prerequisites

- Install the **sg3_utils** package:

  ```
  # yum install sg3_utils
  ```

Procedure

- On both the server and client, check for the proper SCSI device support:

  ```
  # sg_persist --in --report-capabilities --verbose path-to-scsi-device
  ```

  Ensure that the *Persist Through Power Loss Active* (**PTPL_A**) bit is set.

  > **Example 7.1. A SCSI device that supports pNFS SCSI**
  >
  > The following is an example of **sg_persist** output for a SCSI device that supports pNFS SCSI. The **PTPL_A** bit reports **1**.
  >
  > ```
  >     inquiry cdb: 12 00 00 00 24 00
  >     Persistent Reservation In cmd: 5e 02 00 00 00 00 00 20 00 00
  >   LIO-ORG   block11          4.0
  >    Peripheral device type: disk
  >  Report capabilities response:
  >   Compatible Reservation Handling(CRH): 1
  >   Specify Initiator Ports Capable(SIP_C): 1
  >   All Target Ports Capable(ATP_C): 1
  >   Persist Through Power Loss Capable(PTPL_C): 1
  >   Type Mask Valid(TMV): 1
  >   Allow Commands: 1
  >   Persist Through Power Loss Active(PTPL_A): 1
  >     Support indicated in Type mask:
  >       Write Exclusive, all registrants: 1
  >       Exclusive Access, registrants only: 1
  >       Write Exclusive, registrants only: 1
  >       Exclusive Access: 1
  >       Write Exclusive: 1
  >       Exclusive Access, all registrants: 1
  > ```

Additional resources

- **sg_persist(8)** man page

7.4. SETTING UP PNFS SCSI ON THE SERVER

This procedure configures an NFS server to export a pNFS SCSI layout.

Procedure

1. On the server, mount the XFS file system created on the SCSI device.

2. Configure the NFS server to export NFS version 4.1 or higher. Set the following option in the **[nfsd]** section of the **/etc/nfs.conf** file:

   ```
   [nfsd]

   vers4.1=y
   ```

3. Configure the NFS server to export the XFS file system over NFS with the **pnfs** option:

   Example 7.2. An entry in /etc/exports to export pNFS SCSI

   The following entry in the **/etc/exports** configuration file exports the file system mounted at **/exported/directory/** to the **allowed.example.com** client as a pNFS SCSI layout:

   ```
   /exported/directory allowed.example.com(pnfs)
   ```

   > **NOTE**
   >
   > The exported file system must be created on the whole block device, not only on a partition.

Additional resources

- [Exporting NFS shares](#) .

7.5. SETTING UP PNFS SCSI ON THE CLIENT

This procedure configures an NFS client to mount a pNFS SCSI layout.

Prerequisites

- The NFS server is configured to export an XFS file system over pNFS SCSI. See [Setting up pNFS SCSI on the server](#).

Procedure

- On the client, mount the exported XFS file system using NFS version 4.1 or higher:

  ```
  # mount -t nfs -o nfsvers=4.1 host:/remote/export /local/directory
  ```

  Do not mount the XFS file system directly without NFS.

Additional resources

- [Mounting NFS shares](#).

7.6. RELEASING THE PNFS SCSI RESERVATION ON THE SERVER

This procedure releases the persistent reservation that an NFS server holds on a SCSI device. This enables you to repurpose the SCSI device when you no longer need to export pNFS SCSI.

You must remove the reservation from the server. It cannot be removed from a different IT Nexus.

Prerequisites

- Install the **sg3_utils** package:

  ```
  # yum install sg3_utils
  ```

Procedure

1. Query an existing reservation on the server:

   ```
   # sg_persist --read-reservation path-to-scsi-device
   ```

   Example 7.3. Querying a reservation on /dev/sda

   ```
   # *sg_persist --read-reservation /dev/sda*

     LIO-ORG   block_1        4.0
     Peripheral device type: disk
     PR generation=0x8, Reservation follows:
       Key=0x100000000000000
       scope: LU_SCOPE,  type: Exclusive Access, registrants only
   ```

2. Remove the existing registration on the server:

   ```
   # sg_persist --out \
           --release \
           --param-rk=reservation-key \
           --prout-type=6 \
           path-to-scsi-device
   ```

   Example 7.4. Removing a reservation on /dev/sda

   ```
   # sg_persist --out \
           --release \
           --param-rk=0x100000000000000 \
           --prout-type=6 \
           /dev/sda

     LIO-ORG   block_1        4.0
     Peripheral device type: disk
   ```

■

Additional resources

- **sg_persist(8)** man page

CHAPTER 8. CONFIGURING THE SQUID CACHING PROXY SERVER

Squid is a proxy server that caches content to reduce bandwidth and load web pages more quickly. This chapter describes how to set up Squid as a proxy for the HTTP, HTTPS, and FTP protocol, as well as authentication and restricting access.

8.1. SETTING UP SQUID AS A CACHING PROXY WITHOUT AUTHENTICATION

This section describes a basic configuration of Squid as a caching proxy without authentication. The procedure limits access to the proxy based on IP ranges.

Prerequisites

- The procedure assumes that the **/etc/squid/squid.conf** file is as provided by the **squid** package. If you edited this file before, remove the file and reinstall the package.

Procedure

1. Install the **squid** package:

   ```
   # yum install squid
   ```

2. Edit the **/etc/squid/squid.conf** file:

   a. Adapt the **localnet** access control lists (ACL) to match the IP ranges that should be allowed to use the proxy:

      ```
      acl localnet src 192.0.2.0/24
      acl localnet 2001:db8:1::/64
      ```

      By default, the **/etc/squid/squid.conf** file contains the **http_access allow localnet** rule that allows using the proxy from all IP ranges specified in **localnet** ACLs. Note that you must specify all **localnet** ACLs before the **http_access allow localnet** rule.
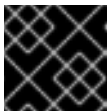
      > **IMPORTANT**
      >
      > Remove all existing **acl localnet** entries that do not match your environment.

   b. The following ACL exists in the default configuration and defines **443** as a port that uses the HTTPS protocol:

      ```
      acl SSL_ports port 443
      ```

      If users should be able to use the HTTPS protocol also on other ports, add an ACL for each of these port:

      ```
      acl SSL_ports port port_number
      ```

   c. Update the list of **acl Safe_ports** rules to configure to which ports Squid can establish a connection. For example, to configure that clients using the proxy can only access

resources on port 21 (FTP), 80 (HTTP), and 443 (HTTPS), keep only the following **acl Safe_ports** statements in the configuration:

```
acl Safe_ports port 21
acl Safe_ports port 80
acl Safe_ports port 443
```

By default, the configuration contains the **http_access deny !Safe_ports** rule that defines access denial to ports that are not defined in **Safe_ports** ACLs.

d. Configure the cache type, the path to the cache directory, the cache size, and further cache type-specific settings in the **cache_dir** parameter:

```
cache_dir ufs /var/spool/squid 10000 16 256
```

With these settings:

- Squid uses the **ufs** cache type.

- Squid stores its cache in the **/var/spool/squid/** directory.

- The cache grows up to **10000** MB.

- Squid creates **16** level-1 sub-directories in the **/var/spool/squid/** directory.

- Squid creates **256** sub-directories in each level-1 directory.
  If you do not set a **cache_dir** directive, Squid stores the cache in memory.

3. If you set a different cache directory than **/var/spool/squid/** in the **cache_dir** parameter:

   a. Create the cache directory:

      ```
      # mkdir -p path_to_cache_directory
      ```

   b. Configure the permissions for the cache directory:

      ```
      # chown squid:squid path_to_cache_directory
      ```

   c. If you run SELinux in **enforcing** mode, set the **squid_cache_t** context for the cache directory:

      ```
      # semanage fcontext -a -t squid_cache_t "path_to_cache_directory(/.*)?"
      # restorecon -Rv path_to_cache_directory
      ```

      If the **semanage** utility is not available on your system, install the **policycoreutils-python-utils** package.

4. Open the **3128** port in the firewall:

   ```
   # firewall-cmd --permanent --add-port=3128/tcp
   # firewall-cmd --reload
   ```

5. Enable and start the **squid** service:

   ```
   # systemctl enable --now squid
   ```

Verification steps

To verify that the proxy works correctly, download a web page using the **curl** utility:

```
# curl -O -L "https://www.redhat.com/index.html" -x "proxy.example.com:3128"
```

If **curl** does not display any error and the **index.html** file was downloaded to the current directory, the proxy works.

8.2. SETTING UP SQUID AS A CACHING PROXY WITH LDAP AUTHENTICATION

This section describes a basic configuration of Squid as a caching proxy that uses LDAP to authenticate users. The procedure configures that only authenticated users can use the proxy.

Prerequisites

- The procedure assumes that the **/etc/squid/squid.conf** file is as provided by the **squid** package. If you edited this file before, remove the file and reinstall the package.

- An service user, such as **uid=proxy_user,cn=users,cn=accounts,dc=example,dc=com** exists in the LDAP directory. Squid uses this account only to search for the authenticating user. If the authenticating user exists, Squid binds as this user to the directory to verify the authentication.

Procedure

1. Install the **squid** package:

   ```
   # yum install squid
   ```

2. Edit the **/etc/squid/squid.conf** file:

   a. To configure the **basic_ldap_auth** helper utility, add the following configuration entry to the top of **/etc/squid/squid.conf**:

      ```
      auth_param basic program /usr/lib64/squid/basic_ldap_auth -b
      "cn=users,cn=accounts,dc=example,dc=com" -D
      "uid=proxy_user,cn=users,cn=accounts,dc=example,dc=com" -W
      /etc/squid/ldap_password -f "(&(objectClass=person)(uid=%s))" -ZZ -H
      ldap://ldap_server.example.com:389
      ```

      The following describes the parameters passed to the **basic_ldap_auth** helper utility in the example above:

      - **-b** *base_DN* sets the LDAP search base.

      - **-D** *proxy_service_user_DN* sets the distinguished name (DN) of the account Squid uses to search for the authenticating user in the directory.

      - **-W** *path_to_password_file* sets the path to the file that contains the password of the proxy service user. Using a password file prevents that the password is visible in the operating system's process list.

- **-f** *LDAP_filter* specifies the LDAP search filter. Squid replaces the **%s** variable with the user name provided by the authenticating user.
  The **(&(objectClass=person)(uid=%s))** filter in the example defines that the user name must match the value set in the **uid** attribute and that the directory entry contains the **person** object class.

- **-ZZ** enforces a TLS-encrypted connection over the LDAP protocol using the **STARTTLS** command. Omit the **-ZZ** in the following situations:

  - The LDAP server does not support encrypted connections.

  - The port specified in the URL uses the LDAPS protocol.

- The -H LDAP_URL parameter specifies the protocol, the host name or IP address, and the port of the LDAP server in URL format.

b. Add the following ACL and rule to configure that Squid allows only authenticated users to use the proxy:

> **acl *ldap-auth* proxy_auth REQUIRED**
> **http_access allow *ldap-auth***

> **IMPORTANT**
>
> Specify these settings before the **http_access deny** all rule.

c. Remove the following rule to disable bypassing the proxy authentication from IP ranges specified in **localnet** ACLs:

> **http_access allow localnet**

d. The following ACL exists in the default configuration and defines **443** as a port that uses the HTTPS protocol:

> acl SSL_ports port 443

If users should be able to use the HTTPS protocol also on other ports, add an ACL for each of these port:

> **acl SSL_ports port port_number**

e. Update the list of **acl Safe_ports** rules to configure to which ports Squid can establish a connection. For example, to configure that clients using the proxy can only access resources on port 21 (FTP), 80 (HTTP), and 443 (HTTPS), keep only the following **acl Safe_ports** statements in the configuration:

> **acl Safe_ports port 21**
> **acl Safe_ports port 80**
> **acl Safe_ports port 443**

By default, the configuration contains the **http_access deny !Safe_ports** rule that defines access denial to ports that are not defined in **Safe_ports ACLs**.

f. Configure the cache type, the path to the cache directory, the cache size, and further cache type-specific settings in the **cache_dir** parameter:

> **cache_dir ufs /var/spool/squid 10000 16 256**

With these settings:

- Squid uses the **ufs** cache type.

- Squid stores its cache in the **/var/spool/squid/** directory.

- The cache grows up to **10000** MB.

- Squid creates **16** level-1 sub-directories in the **/var/spool/squid/** directory.

- Squid creates **256** sub-directories in each level-1 directory.
  If you do not set a **cache_dir** directive, Squid stores the cache in memory.

3. If you set a different cache directory than **/var/spool/squid/** in the **cache_dir** parameter:

   a. Create the cache directory:

   > **# mkdir -p *path_to_cache_directory***

   b. Configure the permissions for the cache directory:

   > **# chown squid:squid *path_to_cache_directory***

   c. If you run SELinux in **enforcing** mode, set the **squid_cache_t** context for the cache directory:

   > **# semanage fcontext -a -t squid_cache_t "*path_to_cache_directory*(/.*)?"**
   > **# restorecon -Rv *path_to_cache_directory***

   If the **semanage** utility is not available on your system, install the **policycoreutils-python-utils** package.

4. Store the password of the LDAP service user in the **/etc/squid/ldap_password** file, and set appropriate permissions for the file:

   > **# echo "*password*" > /etc/squid/ldap_password**
   > **# chown root:squid /etc/squid/ldap_password**
   > **# chmod 640 /etc/squid/ldap_password**

5. Open the **3128** port in the firewall:

   > **# firewall-cmd --permanent --add-port=3128/tcp**
   > **# firewall-cmd --reload**

6. Enable and start the **squid** service:

   > **# systemctl enable --now squid**

Verification steps

To verify that the proxy works correctly, download a web page using the **curl** utility:

```
# curl -O -L "https://www.redhat.com/index.html" -x
"user_name:password@proxy.example.com:3128"
```

If curl does not display any error and the **index.html** file was downloaded to the current directory, the proxy works.

Troubleshooting steps

To verify that the helper utility works correctly:

1. Manually start the helper utility with the same settings you used in the **auth_param** parameter:

   ```
   # /usr/lib64/squid/basic_ldap_auth -b "cn=users,cn=accounts,dc=example,dc=com" -D
   "uid=proxy_user,cn=users,cn=accounts,dc=example,dc=com" -W
   /etc/squid/ldap_password -f "(&(objectClass=person)(uid=%s))" -ZZ -H
   ldap://ldap_server.example.com:389
   ```

2. Enter a valid user name and password, and press **Enter**:

   ```
   user_name password
   ```

   If the helper utility returns **OK**, authentication succeeded.

8.3. SETTING UP SQUID AS A CACHING PROXY WITH KERBEROS AUTHENTICATION

This section describes a basic configuration of Squid as a caching proxy that authenticates users to an Active Directory (AD) using Kerberos. The procedure configures that only authenticated users can use the proxy.

Prerequisites

- The procedure assumes that the **/etc/squid/squid.conf** file is as provided by the **squid** package. If you edited this file before, remove the file and reinstall the package.

- The server on which you want to install Squid is a member of the AD domain. For details, see Setting up Samba as a Domain Member in the Red Hat Enterprise Linux 8 **Deploying different types of servers** documentation.

Procedure

1. Install the following packages:

   ```
   # yum install squid krb5-workstation
   ```

2. Authenticate as the AD domain administrator:

   ```
   # kinit administrator@AD.EXAMPLE.COM
   ```

3. Create a keytab for Squid and store it in the **/etc/squid/HTTP.keytab** file:

> **# export KRB5_KTNAME=FILE:/etc/squid/HTTP.keytab**
> **# net ads keytab CREATE -U administrator**

4. Add the **HTTP** service principal to the keytab:

> **# net ads keytab ADD HTTP -U administrator**

5. Set the owner of the keytab file to the **squid** user:

> **# chown squid /etc/squid/HTTP.keytab**

6. Optionally, verify that the keytab file contains the **HTTP** service principal for the fully-qualified domain name (FQDN) of the proxy server:

```
 klist -k /etc/squid/HTTP.keytab
Keytab name: FILE:/etc/squid/HTTP.keytab
KVNO Principal
---- --------------------------------------------------
...
   2 HTTP/proxy.ad.example.com@AD.EXAMPLE.COM
...
```

7. Edit the **/etc/squid/squid.conf** file:

   a. To configure the **negotiate_kerberos_auth** helper utility, add the following configuration entry to the top of **/etc/squid/squid.conf**:

   > **auth_param negotiate program /usr/lib64/squid/negotiate_kerberos_auth -k**
   > **/etc/squid/HTTP.keytab -s HTTP/*proxy.ad.example.com@AD.EXAMPLE.COM***

   The following describes the parameters passed to the **negotiate_kerberos_auth** helper utility in the example above:

   - **-k** *file* sets the path to the key tab file. Note that the squid user must have read permissions on this file.

   - **-s HTTP/*host_name@kerberos_realm*** sets the Kerberos principal that Squid uses. Optionally, you can enable logging by passing one or both of the following parameters to the helper utility:

   - **-i** logs informational messages, such as the authenticating user.

   - **-d** enables debug logging.
     Squid logs the debugging information from the helper utility to the **/var/log/squid/cache.log** file.

   b. Add the following ACL and rule to configure that Squid allows only authenticated users to use the proxy:

   > **acl *kerb-auth* proxy_auth REQUIRED**
   > **http_access allow *kerb-auth***

> **IMPORTANT**
>
> Specify these settings before the **http_access deny all** rule.

c. Remove the following rule to disable bypassing the proxy authentication from IP ranges specified in **localnet** ACLs:

```
http_access allow localnet
```

d. The following ACL exists in the default configuration and defines **443** as a port that uses the HTTPS protocol:

```
acl SSL_ports port 443
```

If users should be able to use the HTTPS protocol also on other ports, add an ACL for each of these port:

```
acl SSL_ports port port_number
```

e. Update the list of **acl Safe_ports** rules to configure to which ports Squid can establish a connection. For example, to configure that clients using the proxy can only access resources on port 21 (FTP), 80 (HTTP), and 443 (HTTPS), keep only the following **acl Safe_ports** statements in the configuration:

```
acl Safe_ports port 21
acl Safe_ports port 80
acl Safe_ports port 443
```

By default, the configuration contains the **http_access deny !Safe_ports** rule that defines access denial to ports that are not defined in **Safe_ports** ACLs.

f. Configure the cache type, the path to the cache directory, the cache size, and further cache type-specific settings in the **cache_dir** parameter:

```
cache_dir ufs /var/spool/squid 10000 16 256
```

With these settings:

- Squid uses the **ufs** cache type.

- Squid stores its cache in the **/var/spool/squid/** directory.

- The cache grows up to **10000** MB.

- Squid creates **16** level-1 sub-directories in the **/var/spool/squid/** directory.

- Squid creates **256** sub-directories in each level-1 directory.
  If you do not set a **cache_dir** directive, Squid stores the cache in memory.

8. If you set a different cache directory than **/var/spool/squid/** in the **cache_dir** parameter:

a. Create the cache directory:

```
# mkdir -p path_to_cache_directory
```

b. Configure the permissions for the cache directory:

```
# chown squid:squid path_to_cache_directory
```

c. If you run SELinux in **enforcing** mode, set the **squid_cache_t** context for the cache directory:

```
# semanage fcontext -a -t squid_cache_t "path_to_cache_directory(/.*)?"
# restorecon -Rv path_to_cache_directory
```

If the **semanage** utility is not available on your system, install the **policycoreutils-python-utils** package.

9. Open the **3128** port in the firewall:

```
# firewall-cmd --permanent --add-port=3128/tcp
# firewall-cmd --reload
```

10. Enable and start the **squid** service:

```
# systemctl enable --now squid
```

Verification steps

To verify that the proxy works correctly, download a web page using the **curl** utility:

```
# curl -O -L "https://www.redhat.com/index.html" --proxy-negotiate -u : -x "proxy.ad.example.com:3128"
```

If **curl** does not display any error and the **index.html** file exists in the current directory, the proxy works.

Troubleshooting steps

To manually test Kerberos authentication:

1. Obtain a Kerberos ticket for the AD account:

```
# kinit user@AD.EXAMPLE.COM
```

2. Optionally, display the ticket:

```
# klist
```

3. Use the **negotiate_kerberos_auth_test** utility to test the authentication:

```
# /usr/lib64/squid/negotiate_kerberos_auth_test proxy.ad.example.com
```

If the helper utility returns a token, the authentication succeeded:

```
Token: YIIFtAYGKwYBBBQUCoIIFqDC...
```

8.4. CONFIGURING A DOMAIN DENY LIST IN SQUID

Frequently, administrators want to block access to specific domains. This section describes how to configure a domain deny list in Squid.

Prerequisites

- Squid is configured, and users can use the proxy.

Procedure

1. Edit the **/etc/squid/squid.conf** file and add the following settings:

   > **acl domain_deny_list dstdomain "/etc/squid/domain_deny_list.txt"**
   > **http_access deny all domain_deny_list**

   > **IMPORTANT**
   >
   > Add these entries before the first **http_access allow** statement that allows access to users or clients.

2. Create the **/etc/squid/domain_deny_list.txt** file and add the domains you want to block. For example, to block access to **example.com** including subdomains and to block **example.net**, add:

   > **.example.com**
   > **example.net**

   > **IMPORTANT**
   >
   > If you referred to the **/etc/squid/domain_deny_list.txt** file in the squid configuration, this file must not be empty. If the file is empty, Squid fails to start.

3. Restart the **squid** service:

   > **# systemctl restart squid**

8.5. CONFIGURING THE SQUID SERVICE TO LISTEN ON A SPECIFIC PORT OR IP ADDRESS

By default, the Squid proxy service listens on the **3128** port on all network interfaces. This section describes how to change the port and configuring Squid to listen on a specific IP address.

Prerequisites

- The **squid** package is installed.

Procedure

1. Edit the **/etc/squid/squid.conf** file:

   - To set the port on which the Squid service listens, set the port number in the **http_port** parameter. For example, to set the port to **8080**, set:

> **http_port 8080**

- To configure on which IP address the Squid service listens, set the IP address and port number in the **http_port** parameter. For example, to configure that Squid listens only on the **192.0.2.1** IP address on port **3128**, set:

> **http_port 192.0.2.1:3128**

Add multiple **http_port** parameters to the configuration file to configure that Squid listens on multiple ports and IP addresses:

> **http_port 192.0.2.1:3128**
> **http_port 192.0.2.1:8080**

2. If you configured that Squid uses a different port as the default (**3128**):

   a. Open the port in the firewall:

   > **# firewall-cmd --permanent --add-port=*port_number*/tcp**
   > **# firewall-cmd --reload**

   b. If you run SELinux in enforcing mode, assign the port to the **squid_port_t** port type definition:

   > **# semanage port -a -t squid_port_t -p tcp *port_number***

   If the **semanage** utility is not available on your system, install the **policycoreutils-python-utils** package.

3. Restart the **squid** service:

> **# systemctl restart squid**

8.6. ADDITIONAL RESOURCES

- Configuration parameters **usr/share/doc/squid-<version>/squid.conf.documented**

CHAPTER 9. DATABASE SERVERS

9.1. INTRODUCTION TO DATABASE SERVERS

A database server is a service that provides features of a database management system (DBMS). DBMS provides utilities for database administration and interacts with end users, applications, and databases.

Red Hat Enterprise Linux 8 provides the following database management systems:

- **MariaDB 10.3**

- **MariaDB 10.5** – available since RHEL 8.4

- **MySQL 8.0**

- **PostgreSQL 10**

- **PostgreSQL 9.6**

- **PostgreSQL 12** – available since RHEL 8.1.1

- **PostgreSQL 13** – available since RHEL 8.4

9.2. USING MARIADB

The **MariaDB** server is an open source fast and robust database server that is based on the **MySQL** technology. This part describes how to install and configure **MariaDB** on a RHEL system, how to back up **MariaDB** data, how to migrate from an earlier **MariaDB** version, and how to replicate a database using the **MariaDB Galera Cluster**.

9.2.1. Getting started with MariaDB

MariaDB is a relational database that converts data into structured information and provides an SQL interface for accessing data. It includes multiple storage engines and plug-ins, as well as geographic information system (GIS) and JavaScript Object Notation (JSON) features.

For Red Hat Enterprise Linux 8 this part describes:

- How to install the **MariaDB** server in Installing MariaDB.

- How to adjust **MariaDB** configuration in Configuring MariaDB.

- How to back up **MariaDB** data in Backing up MariaDB data.

- How to migrate from the RHEL 7 default version, **MariaDB 5.5**, to the RHEL 8 default version, **MariaDB 10.3**, in Migrating to MariaDB 10.3.

- How to upgrade from **MariaDB 10.3** to **MariaDB 10.5** within RHEL 8 in Upgrading from MariaDB 10.3 to MariaDB 10.5.

- How to replicate a database using the **MariaDB Galera Cluster** in Replicating MariaDB with Galera.

9.2.2. Installing MariaDB

In RHEL 8, the **MariaDB** server is available in the following versions, each provided by a separate stream:

- **MariaDB 10.3**

- **MariaDB 10.5** – available since RHEL 8.4

To install **MariaDB**, use the following procedure.

Procedure

1. Install **MariaDB** server packages by selecting a stream (version) from the **mariadb** module and specifying the **server** profile. For example:

   > # **yum module install mariadb:10.3/server**

2. Start the **mariadb** service:

   > # **systemctl start mariadb.service**

3. Enable the **mariadb** service to start at boot:

   > # **systemctl enable mariadb.service**

4. *Recommended:* To improve security when installing **MariaDB**, run the following command:

   > $ **mysql_secure_installation**

   The command launches a fully interactive script, which prompts for each step in the process. The script enables you to improve security in the following ways:

   - Setting a password for root accounts

   - Removing anonymous users

   - Disallowing remote root logins (outside the local host)

> **NOTE**
>
> The **MariaDB** and **MySQL** database servers cannot be installed in parallel in RHEL 8 due to conflicting RPM packages. Parallel installation of components is possible in Red Hat Software Collections for RHEL 7. In RHEL 8, different versions of database servers can be used in containers.

9.2.3. Configuring MariaDB

To configure the **MariaDB** server for networking, use the following procedure.

Procedure

1. Edit the **[mysqld]** section of the **/etc/my.cnf.d/mariadb-server.cnf** file. You can set the following configuration directives:

   - **bind-address** – is the address on which the server listens. Possible options are:

- a host name

- an IPv4 address

- an IPv6 address

- **skip-networking** – controls whether the server listens for TCP/IP connections. Possible values are:

    - 0 – to listen for all clients

    - 1 – to listen for local clients only

- **port** – the port on which **MariaDB** listens for TCP/IP connections.

2. Restart the **mariadb** service:

    ```
    # systemctl restart mariadb.service
    ```

9.2.4. Setting up TLS encryption on a MariaDB server

By default, **MariaDB** uses unencrypted connections. For secure connections, enable TLS support on the **MariaDB** server and configure your clients to establish encrypted connections.

9.2.4.1. Placing the CA certificate, server certificate, and private key on the MariaDB server

Before you can enable TLS encryption in the **MariaDB** server, store the certificate authority (CA) certificate, the server certificate, and the private key on the **MariaDB** server.

Prerequisites

- The following files in Privacy Enhanced Mail (PEM) format have been copied to the server:

    - The private key of the server: **server.example.com.key.pem**

    - The server certificate: **server.example.com.crt.pem**

    - The Certificate Authority (CA) certificate: **ca.crt.pem**

    For details about creating a private key and certificate signing request (CSR), as well as about requesting a certificate from a CA, see your CA's documentation.

Procedure

1. Store the CA and server certificates in the **/etc/pki/tls/certs/** directory:

    ```
    # mv <path>/server.example.com.crt.pem /etc/pki/tls/certs/
    # mv <path>/ca.crt.pem /etc/pki/tls/certs/
    ```

2. Set permissions on the CA and server certificate that enable the **MariaDB** server to read the files:

    ```
    # chmod 644 /etc/pki/tls/certs/server.example.com.crt.pem /etc/pki/tls/certs/ca.crt.pem
    ```

Because certificates are part of the communication before a secure connection is established, any client can retrieve them without authentication. Therefore, you do not need to set strict permissions on the CA and server certificate files.

3. Store the server's private key in the **/etc/pki/tls/private/** directory:

> # mv *<path>*/server.example.com.key.pem /etc/pki/tls/private/

4. Set secure permissions on the server's private key:

> # chmod 640 /etc/pki/tls/private/server.example.com.key.pem
> # chgrp mysql /etc/pki/tls/private/server.example.com.key.pem

If unauthorized users have access to the private key, connections to the **MariaDB** server are no longer secure.

5. Restore the SELinux context:

> # restorecon -Rv /etc/pki/tls/

9.2.4.2. Configuring TLS on a MariaDB server

To improve security, enable TLS support on the **MariaDB** server. As a result, clients can transmit data with the server using TLS encryption.

Prerequisites

- You installed the **MariaDB** server.

- The **mariadb** service is running.

- The following files in Privacy Enhanced Mail (PEM) format exist on the server and are readable by the **mysql** user:

  - The private key of the server: **/etc/pki/tls/private/server.example.com.key.pem**

  - The server certificate: **/etc/pki/tls/certs/server.example.com.crt.pem**

  - The Certificate Authority (CA) certificate **/etc/pki/tls/certs/ca.crt.pem**

- The subject distinguished name (DN) or the subject alternative name (SAN) field in the server certificate matches the server's host name.

Procedure

1. Create the **/etc/my.cnf.d/mariadb-server-tls.cnf** file:

   a. Add the following content to configure the paths to the private key, server and CA certificate:

   > [mariadb]
   > ssl_key = */etc/pki/tls/private/server.example.com.key.pem*
   > ssl_cert = */etc/pki/tls/certs/server.example.com.crt.pem*
   > ssl_ca = */etc/pki/tls/certs/ca.crt.pem*

b. If you have a Certificate Revocation List (CRL), configure the **MariaDB** server to use it:

> **ssl_crl =** */etc/pki/tls/certs/example.crl.pem*

c. Optional: If you run **MariaDB 10.5.2** or later, you can reject connection attempts without encryption. To enable this feature, append:

> **require_secure_transport = on**

d. Optional: If you run **MariaDB 10.4.6** or later, you can set the TLS versions the server should support. For example, to support TLS 1.2 and TLS 1.3, append:

> **tls_version =** *TLSv1.2,TLSv1.3*

By default, the server supports TLS 1.1, TLS 1.2, and TLS 1.3.

2. Restart the **mariadb** service:

> **# systemctl restart mariadb**

Verification

To simplify troubleshooting, perform the following steps on the **MariaDB** server before you configure the local client to use TLS encryption:

1. Verify that **MariaDB** now has TLS encryption enabled:

```
# mysql -u root -p -e "SHOW GLOBAL VARIABLES LIKE 'have_ssl';"
+---------------+-----------------+
| Variable_name | Value           |
+---------------+-----------------+
| have_ssl      | YES             |
+---------------+-----------------+
```

If the **have_ssl** variable is set to **yes**, TLS encryption is enabled.

2. If you configured the **MariaDB** service to only support specific TLS versions, display the **tls_version** variable:

```
# mysql -u root -p -e "SHOW GLOBAL VARIABLES LIKE 'tls_version';"
+---------------+-----------------+
| Variable_name | Value           |
+---------------+-----------------+
| tls_version   | TLSv1.2,TLSv1.3 |
+---------------+-----------------+
```

Additional resources

- Placing the CA certificate, server certificate, and private key on the MariaDB server

- Upgrading from MariaDB 10.3 to MariaDB 10.5

9.2.4.3. Requiring TLS encrypted connections for specific user accounts

Users that have access to sensitive data should always use a TLS-encrypted connection to avoid sending data unencrypted over the network.

If you cannot configure on the server that a secure transport is required for all connections (**require_secure_transport = on**), configure individual user accounts to require TLS encryption.

Prerequisites

- The **MariaDB** server has TLS support enabled.

- The user you configure to require secure transport exists.

Procedure

1. Connect as an administrative user to the **MariaDB** server:

   > # **mysql -u *root* -p -h *server.example.com***

   If your administrative user has no permissions to access the server remotely, perform the command on the **MariaDB** server and connect to **localhost**.

2. Use the **REQUIRE SSL** clause to enforce that a user must connect using a TLS-encrypted connection:

   > MariaDB [(none)]> **ALTER USER __'example__'@__'%'__ REQUIRE SSL;**

Verification

1. Connect to the server as the **example** user using TLS encryption:

   > # **mysql -u *example* -p -h server.example.com --ssl**
   > ...
   > MariaDB [(none)]>

   If no error is shown and you have access to the interactive **MariaDB** console, the connection with TLS succeeds.

2. Attempt to connect as the **example** user with TLS disabled:

   > # **mysql -u *example* -p -h server.example.com --skip-ssl**
   > ERROR 1045 (28000): Access denied for user 'example'@'server.example.com' (using password: YES)

   The server rejected the login attempt, because TLS is required for this user, but disabled (**--skip-ssl**).

Additional resources

- [Configuring TLS on a MariaDB server](#)

9.2.5. Globally enabling TLS encryption in MariaDB clients

If your **MariaDB** server supports TLS encryption, configure your clients to establish only secure connections and to verify the server certificate. This procedure describes how to enable TLS support for all users on the server.

9.2.5.1. Configuring the MariaDB client to use TLS encryption by default

On RHEL, you can globally configure that the **MariaDB** client uses TLS encryption and verifies that the Common Name (CN) in the server certificate matches the hostname the user connects to. This prevents man-in-the-middle attacks.

Prerequisites

- The **MariaDB** server has TLS support enabled.

- If the certificate authority (CA) that issued the server's certificate is not trusted by RHEL, the CA certificate has been copied to the client.

Procedure

1. If RHEL does not trust the CA that issued the server's certificate:

   a. Copy the CA certificate to the **/etc/pki/ca-trust/source/anchors/** directory:

      ```
      # cp <path>/ca.crt.pem /etc/pki/ca-trust/source/anchors/
      ```

   b. Set permissions that enable all users to read the CA certificate file:

      ```
      # chmod 644 /etc/pki/ca-trust/source/anchors/ca.crt.pem
      ```

   c. Rebuild the CA trust database:

      ```
      # update-ca-trust
      ```

2. Create the **/etc/my.cnf.d/mariadb-client-tls.cnf** file with the following content:

   ```
   [client-mariadb]
   ssl
   ssl-verify-server-cert
   ```

   These settings define that the **MariaDB** client uses TLS encryption (**ssl**) and that the client compares the hostname with the CN in the server certificate (**ssl-verify-server-cert**).

Verification

- Connect to the server using the hostname, and display the server status:

   ```
   # mysql -u root -p -h server.example.com -e status
   ...
   SSL:       Cipher in use is TLS_AES_256_GCM_SHA384
   ```

   If the **SSL** entry contains **Cipher in use is…**, the connection is encrypted.

   Note that the user you use in this command has permissions to authenticate remotely.

If the hostname you connect to does not match the hostname in the TLS certificate of the server, the **ssl-verify-server-cert** parameter causes the connection to fail. For example, if you connect to **localhost**:

```
# mysql -u root -p -h localhost -e status
ERROR 2026 (HY000): SSL connection error: Validation of SSL server certificate failed
```

Additional resources

- The **--ssl*** parameter descriptions in the **mysql(1)** man page.

9.2.6. Backing up MariaDB data

There are two main ways to back up data from a **MariaDB** database in Red Hat Enterprise Linux 8:

- Logical backup

- Physical backup

Logical backup consists of the SQL statements necessary to restore the data. This type of backup exports information and records in plain text files.

The main advantage of logical backup over physical backup is portability and flexibility. The data can be restored on other hardware configurations, **MariaDB** versions or Database Management System (DBMS), which is not possible with physical backups.

Note that logical backup can be performed if the **mariadb.service** is running. Logical backup does not include log and configuration files.

Physical backup consists of copies of files and directories that store the content.

Physical backup has the following advantages compared to logical backup:

- Output is more compact.

- Backup is smaller in size.

- Backup and restore are faster.

- Backup includes log and configuration files.

Note that physical backup must be performed when the **mariadb.service** is not running or all tables in the database are locked to prevent changes during the backup.

You can use one of the following **MariaDB** backup approaches to back up data from a **MariaDB** database:

- Logical backup with **mysqldump**

- Physical online backup using the **Mariabackup** utility

- File system backup

- Replication as a backup solution

9.2.6.1. Performing logical backup with mysqldump

The **mysqldump** client is a backup utility, which can be used to dump a database or a collection of databases for the purpose of a backup or transfer to another database server. The output of **mysqldump** typically consists of SQL statements to re-create the server table structure, populate it with data, or both. **mysqldump** can also generate files in other formats, including XML and delimited text formats, such as CSV.

To perform the **mysqldump** backup, you can use one of the following options:

- Back up one or more selected databases

- Back up all databases

- Back up a subset of tables from one database

Procedure

- To dump a single database, run:

  ```
  # mysqldump [options] --databases db_name > backup-file.sql
  ```

- To dump multiple databases at once, run:

  ```
  # mysqldump [options] --databases db_name1 [db_name2 …] > backup-file.sql
  ```

- To dump all databases, run:

  ```
  # mysqldump [options] --all-databases > backup-file.sql
  ```

- To load one or more dumped full databases back into a server, run:

  ```
  # mysql < backup-file.sql
  ```

- To load a database to a remote **MariaDB** server, run:

  ```
  # mysql --host=remote_host < backup-file.sql
  ```

- To dump a subset of tables from one database, add a list of the chosen tables at the end of the **mysqldump** command:

  ```
  # mysqldump [options] db_name [tbl_name …] > backup-file.sql
  ```

- To load a subset of tables dumped from one database, run:

  ```
  # mysql db_name < backup-file.sql
  ```

  > **NOTE**
  >
  > The *db_name* database must exist at this point.

- To see a list of the options that **mysqldump** supports, run:

  ```
  $ mysqldump --help
  ```

Additional resources

- For more information on logical backup with **mysqldump**, see the [MariaDB Documentation](#).

9.2.6.2. Performing physical online backup using the Mariabackup utility

Mariabackup is a utility based on the Percona XtraBackup technology, which enables performing physical online backups of InnoDB, Aria, and MyISAM tables. This utility is provided by the **mariadb-backup** package from the AppStream repository.

Mariabackup supports full backup capability for **MariaDB** server, which includes encrypted and compressed data.

Prerequisites

- The **mariadb-backup** package is installed on the system:

  # **yum install mariadb-backup**

- You must provide **Mariabackup** with credentials for the user under which the backup will be run. You can provide the credentials either on the command line or by a configuration file.

- Users of **Mariabackup** must have the **RELOAD**, **LOCK TABLES**, and **REPLICATION CLIENT** privileges.

To create a backup of a database using **Mariabackup**, use the following procedure.

Procedure

- To create a backup while providing credentials on the command line, run:

  $ **mariabackup --backup --target-dir <backup_directory> --user <backup_user> --password <backup_passwd>**

  The **target-dir** option defines the directory where the backup files are stored. If you want to perform a full backup, the target directory must be empty or not exist.

  The **user** and **password** options allow you to configure the user name and the password.

- To create a backup with credentials set in a configuration file:

  1. Create a configuration file in the /**etc/my.cnf.d**/ directory, for example, /**etc/my.cnf.d/mariabackup.cnf**.

  2. Add the following lines into the **[xtrabackup]** or **[mysqld]** section of the new file:

     ```
     [xtrabackup]
     user=myuser
     password=mypassword
     ```

  3. Perform the backup:

     $ **mariabackup --backup --target-dir <backup_directory>**

> **IMPORTANT**
>
> **Mariabackup** does not read options in the **[mariadb]** section of configuration files. If a non-default data directory is specified on a **MariaDB** server, you must specify this directory in the **[xtrabackup]** or **[mysqld]** sections of configuration files so that **Mariabackup** is able to find the data directory.
>
> To specify a non-default data directory, include the following line in the **[xtrabackup]** or **[mysqld]** sections of **MariaDB** configuration files:
>
> ```
> datadir=/var/mycustomdatadir
> ```

Additional resources

- [Full Backup and Restore with Mariabackup](#)

9.2.6.3. Restoring data using the Mariabackup utility

When the backup is complete, you can restore the data from the backup by using the **mariabackup** command with one of the following options:

- **--copy-back** allows you to keep the original backup files.

- **--move-back** moves the backup files to the data directory and removes the original backup files.

To restore data using the **Mariabackup** utility, use the following procedure.

Prerequisites

- Verify that the **mariadb** service is not running:

  ```
  # systemctl stop mariadb.service
  ```

- Verify that the data directory is empty.

- Users of **Mariabackup** must have the **RELOAD**, **LOCK TABLES**, and **REPLICATION CLIENT** privileges.

Procedure

1. Run the **mariabackup** command:

   - To restore data and keep the original backup files, use the **--copy-back** option:

     ```
     $ mariabackup --copy-back --target-dir=/var/mariadb/backup/
     ```

   - To restore data and remove the original backup files, use the **--move-back** option:

     ```
     $ mariabackup --move-back --target-dir=/var/mariadb/backup/
     ```

2. Fix the file permissions.
   When restoring a database, **Mariabackup** preserves the file and directory privileges of the backup. However, **Mariabackup** writes the files to disk as the user and group restoring the

database. After restoring a backup, you may need to adjust the owner of the data directory to match the user and group for the **MariaDB** server, typically **mysql** for both.

For example, to recursively change ownership of the files to the **mysql** user and group:

```
# chown -R mysql:mysql /var/lib/mysql/
```

3. Start the **mariadb** service:

```
# systemctl start mariadb.service
```

Additional resources

- [Full Backup and Restore with Mariabackup](#)

9.2.6.4. Performing file system backup

To create a file system backup of **MariaDB** data files, copy the content of the **MariaDB** data directory to your backup location.

To back up also your current configuration or the log files, use the optional steps of the following procedure.

Procedure

1. Stop the **mariadb** service:

```
# systemctl stop mariadb.service
```

2. Copy the data files to the required location:

```
# cp -r /var/lib/mysql /backup-location
```

3. Optionally, copy the configuration files to the required location:

```
# cp -r /etc/my.cnf /etc/my.cnf.d /backup-location/configuration
```

4. Optionally, copy the log files to the required location:

```
# cp /var/log/mariadb/ /backup-location/logs*
```

5. Start the **mariadb** service:

```
# systemctl start mariadb.service
```

6. When loading the backed up data from the backup location to the **/var/lib/mysql** directory, ensure that **mysql:mysql** is an owner of all data in **/var/lib/mysql**:

```
# chown -R mysql:mysql /var/lib/mysql
```

9.2.6.5. Replication as a backup solution

Replication is an alternative backup solution for source servers. If a source server replicates to a replica server, backups can be run on the replica without any impact on the source. The source can still run while you shut down the replica and back the data up from the replica.

> **WARNING**
>
> Replication itself is not a sufficient backup solution. Replication protects source servers against hardware failures, but it does not ensure protection against data loss. It is recommended that you use any other backup solution on the replica together with this method.

Additional resources

- Replicating MariaDB with Galera

- Replication as a backup solution

9.2.7. Migrating to MariaDB 10.3

RHEL 7 contains **MariaDB 5.5** as the default implementation of a server from the MySQL databases family. Later versions of the **MariaDB** database server are available as Software Collections for RHEL 7. RHEL 8 provides **MariaDB 10.3**, **MariaDB 10.5**, and **MySQL 8.0**.

This part describes migration to **MariaDB 10.3** from a RHEL 7 or Red Hat Software Collections version of **MariaDB**. If you want to migrate from **MariaDB 10.3** to **MariaDB 10.5** within RHEL 8, see Upgrading from MariaDB 10.3 to MariaDB 10.5 instead.

9.2.7.1. Notable differences between the RHEL 7 and RHEL 8 versions of MariaDB

The most important changes between **MariaDB 5.5** and **MariaDB 10.3** are as follows:

- **MariaDB Galera Cluster**, a synchronous multi-source cluster, is a standard part of **MariaDB** since 10.1.

- The ARCHIVE storage engine is no longer enabled by default, and the plug-in needs to be specifically enabled.

- The BLACKHOLE storage engine is no longer enabled by default, and the plug-in needs to be specifically enabled.

- InnoDB is used as the default storage engine instead of XtraDB, which was used in **MariaDB 10.1** and earlier versions.
  For more details, see Why does MariaDB 10.2 use InnoDB instead of XtraDB? .

- The new **mariadb-connector-c** packages provide a common client library for MySQL and MariaDB. This library is usable with any version of the **MySQL** and **MariaDB** database servers. As a result, the user is able to connect one build of an application to any of the MySQL and **MariaDB** servers distributed with Red Hat Enterprise Linux 8.

To migrate from **MariaDB 5.5** to **MariaDB 10.3**, you need to perform multiple configuration changes as described in Migrating to MariaDB .

9.2.7.2. Configuration changes

The recommended migration path from **MariaDB 5.5** to **MariaDB 10.3** is to upgrade to **MariaDB 10.0** first, and then upgrade by one version successively.

The main advantage of upgrading one minor version at a time is better adaptation of the database, including both data and configuration, to the changes. The upgrade ends on the same major version as is available in RHEL 8 (MariaDB 10.3), which significantly reduces configuration changes or other issues.

For more information about configuration changes when migrating from **MariaDB 5.5** to **MariaDB 10.0**, see Migrating to MariaDB 10.0 in Red Hat Software Collections documentation.

The migration to following successive versions of **MariaDB** and the required configuration changes is described in these documents:

- Migrating to MariaDB 10.1 in Red Hat Software Collections documentation.

- Migrating to MariaDB 10.2 in Red Hat Software Collections documentation.

- Migrating to MariaDB 10.3 in Red Hat Software Collections documentation.

> **NOTE**
>
> Migration directly from **MariaDB 5.5** to **MariaDB 10.3** is also possible, but you must perform all configuration changes that are required by differences described in the migration documents above.

9.2.7.3. In-place upgrade using the mysql_upgrade utility

To migrate the database files to RHEL 8, users of **MariaDB** on RHEL 7 must perform the in-place upgrade using the **mysql_upgrade** utility. The **mysql_upgrade** utility is provided by the **mariadb-server-utils** subpackage, which is installed as a dependency of the **mariadb-server** package.

To perform an in-place upgrade, you must copy binary data files to the **/var/lib/mysql/** data directory on the RHEL 8 system and use the **mysql_upgrade** utility.

You can use this method for migrating data from:

- The Red Hat Enterprise Linux 7 version of **MariaDB 5.5**

- The Red Hat Software Collections versions of:

  - **MariaDB 5.5** (no longer supported)

  - **MariaDB 10.0** (no longer supported)

  - **MariaDB 10.1** (no longer supported)

  - **MariaDB 10.2** (no longer supported)

  - MariaDB 10.3
    Note that it is recommended to upgrade to **MariaDB 10.3** by one version successively. See the respective Migration chapters in the Release Notes for Red Hat Software Collections .

> **NOTE**
>
> If you are upgrading from the RHEL 7 version of **MariaDB**, the source data is stored in the **/var/lib/mysql/** directory. In case of Red Hat Software Collections versions of **MariaDB**, the source data directory is **/var/opt/rh/<collection_name>/lib/mysql/** (with the exception of the **mariadb55**, which uses the **/opt/rh/mariadb55/root/var/lib/mysql/** data directory).

To perform an upgrade using the **mysql_upgrade** utility, use the following procedure.

Prerequisites

- Before performing the upgrade, back up all your data stored in the **MariaDB** databases.

Procedure

1. Ensure that the **mariadb-server** package is installed on the RHEL 8 system:

   # **yum install mariadb-server**

2. Ensure that the **mariadb** service is not running on either of the source and target systems at the time of copying data:

   # **systemctl stop mariadb.service**

3. Copy the data from the source location to the **/var/lib/mysql/** directory on the RHEL 8 target system.

4. Set the appropriate permissions and SELinux context for copied files on the target system:

   # **restorecon -vr /var/lib/mysql**

5. Start the **MariaDB** server on the target system:

   # **systemctl start mariadb.service**

6. Run the **mysql_upgrade** command to check and repair internal tables:

   $ **mysql_upgrade**

7. When the upgrade is complete, verify that all configuration files within the **/etc/my.cnf.d/** directory include only options valid for **MariaDB 10.3**.

> **IMPORTANT**
>
> There are certain risks and known problems related to an in-place upgrade. For example, some queries might not work or they will be run in different order than before the upgrade. For more information about these risks and problems, and for general information about an in-place upgrade, see MariaDB 10.3 Release Notes .

9.2.8. Upgrading from MariaDB 10.3 to MariaDB 10.5

This part describes migration from **MariaDB 10.3** to **MariaDB 10.5** within RHEL 8.

9.2.8.1. Notable differences between MariaDB 10.3 and MariaDB 10.5

Significant changes between **MariaDB 10.3** and **MariaDB 10.5** include:

- **MariaDB** now uses the **unix_socket** authentication plug-in by default. The plug-in enables users to use operating system credentials when connecting to **MariaDB** through the local UNIX socket file.

- **MariaDB** adds **mariadb-*** named binaries and **mysql*** symbolic links pointing to the **mariadb-*** binaires. For example, the **mysqladmin**, **mysqlaccess**, and **mysqlshow** symlinks point to the **mariadb-admin**, **mariadb-access**, and **mariadb-show** binaries, respectively.

- The **SUPER** privilege has been split into several privileges to better align with each user role. As a result, certain statements have changed required privileges.

- In parallel replication, the **slave_parallel_mode** now defaults to **optimistic**.

- In the InnoDB storage engine, defaults of the following variables have been changed: **innodb_adaptive_hash_index** to **OFF** and **innodb_checksum_algorithm** to **full_crc32**.

- **MariaDB** now uses the **libedit** implementation of the underlying software managing the **MariaDB** command history (the **.mysql_history** file) instead of the previously used **readline** library. This change impacts users working directly with the **.mysql_history** file. Note that **.mysql_history** is a file managed by the **MariaDB** or **MySQL** applications, and users should not work with the file directly. The human-readable appearance is coincidental.

> **NOTE**
>
> To increase security, you can consider not maintaining a history file. To disable the command history recording:
>
> 1. Remove the **.mysql_history** file if it exists.
>
> 2. Use either of the following approaches:
>
>   - Set the **MYSQL_HISTFILE** variable to **/dev/null** and include this setting in any of your shell's startup files.
>
>   - Change the **.mysql_history** file to a symbolic link to **/dev/null**:
>
>      ```
>      $ ln -s /dev/null $HOME/.mysql_history
>      ```

MariaDB Galera Cluster has been upgraded to version 4 with the following notable changes:

- **Galera** adds a new streaming replication feature, which supports replicating transactions of unlimited size. During an execution of streaming replication, a cluster replicates a transaction in small fragments.

- **Galera** now fully supports Global Transaction ID (GTID).

- The default value for the **wsrep_on** option in the **/etc/my.cnf.d/galera.cnf** file has changed from **1** to **0** to prevent end users from starting **wsrep** replication without configuring required additional options.

Changes to the PAM plug-in in **MariaDB 10.5** include:

- **MariaDB 10.5** adds a new version of the Pluggable Authentication Modules (PAM) plug-in. The PAM plug-in version 2.0 performs PAM authentication using a separate **setuid root** helper binary, which enables **MariaDB** to use additional PAM modules.

- The helper binary can be executed only by users in the **mysql** group. By default, the group contains only the **mysql** user. Red Hat recommends that administrators do not add more users to the **mysql** group to prevent password-guessing attacks without throttling or logging through this helper utility.

- In **MariaDB 10.5**, the Pluggable Authentication Modules (PAM) plug-in and its related files have been moved to a new package, **mariadb-pam**. As a result, no new **setuid root** binary is introduced on systems that do not use PAM authentication for **MariaDB**.

- The **mariadb-pam** package contains both PAM plug-in versions: version 2.0 is the default, and version 1.0 is available as the **auth_pam_v1** shared object library.

- The **mariadb-pam** package is not installed by default with the MariaDB server. To make the PAM authentication plug-in available in **MariaDB 10.5**, install the **mariadb-pam** package manually.

9.2.8.2. Upgrading from a RHEL 8 version of MariaDB 10.3 to MariaDB 10.5

This procedure describes upgrading from the **mariadb:10.3** module stream to the **mariadb:10.5** module stream using the **yum** and **mariadb-upgrade** utilities.

The **mariadb-upgrade** utility is provided by the **mariadb-server-utils** subpackage, which is installed as a dependency of the **mariadb-server** package.

Prerequisites

- Before performing the upgrade, back up all your data stored in the **MariaDB** databases.

Procedure

1. Stop the **MariaDB** server:

   ```
   # systemctl stop mariadb.service
   ```

2. Execute the following command to determine if your system is prepared for switching to a later stream:

   ```
   # yum distro-sync
   ```

   This command must finish with the message *Nothing to do. Complete!* For more information, see Switching to a later stream .

3. Reset the **mariadb** module on your system:

   ```
   # yum module reset mariadb
   ```

4. Enable the new **mariadb:10.5** module stream:

   ```
   # yum module enable mariadb:10.5
   ```

5. Synchronize installed packages to perform the change between streams:

```
# yum distro-sync
```

This will update all installed **MariaDB** packages.

6. Adjust configuration so that option files located in **/etc/my.cnf.d/** include only options valid for **MariaDB 10.5**. For details, see upstream documentation for MariaDB 10.4 and MariaDB 10.5.

7. Start the **MariaDB** server.

- When upgrading a database running standalone:

```
# systemctl start mariadb.service
```

- When upgrading a **Galera** cluster node:

```
# galera_new_cluster
```

The **mariadb** service will be started automatically.

8. Execute the **mariadb-upgrade** utility to check and repair internal tables.

- When upgrading a database running standalone:

```
# mariadb-upgrade
```

- When upgrading a **Galera** cluster node:

```
# mariadb-upgrade --skip-write-binlog
```

> **IMPORTANT**
>
> There are certain risks and known problems related to an in-place upgrade. For example, some queries might not work or they will be run in different order than before the upgrade. For more information about these risks and problems, and for general information about an in-place upgrade, see MariaDB 10.5 Release Notes .

9.2.9. Replicating MariaDB with Galera

This section describes how to replicate a **MariaDB** database using the Galera solution on Red Hat Enterprise Linux 8.

9.2.9.1. Introduction to MariaDB Galera Cluster

Galera replication is based on the creation of a synchronous multi-source **MariaDB Galera Cluster** consisting of multiple MariaDB servers. Unlike the traditional primary/replica setup where replicas are usually read-only, nodes in the MariaDB Galera Cluster can be all writable.

The interface between Galera replication and a **MariaDB** database is defined by the write set replication API (**wsrep API**).

The main features of **MariaDB Galera Cluster** are:

- Synchronous replication

- Active-active multi-source topology

- Read and write to any cluster node

- Automatic membership control, failed nodes drop from the cluster

- Automatic node joining

- Parallel replication on row level

- Direct client connections: users can log on to the cluster nodes, and work with the nodes directly while the replication runs

Synchronous replication means that a server replicates a transaction at commit time by broadcasting the write set associated with the transaction to every node in the cluster. The client (user application) connects directly to the Database Management System (DBMS), and experiences behavior that is similar to native **MariaDB**.

Synchronous replication guarantees that a change that happened on one node in the cluster happens on other nodes in the cluster at the same time.

Therefore, synchronous replication has the following advantages over asynchronous replication:

- No delay in propagation of the changes between particular cluster nodes

- All cluster nodes are always consistent

- The latest changes are not lost if one of the cluster nodes crashes

- Transactions on all cluster nodes are executed in parallel

- Causality across the whole cluster

Additional resources

- About Galera replication

- What is MariaDB Galera Cluster

- Getting started with MariaDB Galera Cluster

9.2.9.2. Components to build MariaDB Galera Cluster

To build **MariaDB Galera Cluster**, you must install the following packages on your system:

- **mariadb-server-galera** - contains support files and scripts for **MariaDB Galera Cluster**.

- **mariadb-server** - is patched by **MariaDB** upstream to include the write set replication API (**wsrep API**). This API provides the interface between **Galera** replication and **MariaDB**.

- **galera** - is patched by **MariaDB** upstream to add full support for **MariaDB**. The **galera** package contains the following:

  - **Galera Replication Library** provides the whole replication functionality.

- The **Galera Arbitrator** utility can be used as a cluster member that participates in voting in split-brain scenarios. However, **Galera Arbitrator** cannot participate in the actual replication.

- **Galera Systemd service** and **Galera wrapper script** which are used for deploying the Galera Arbitrator utility. **MariaDB** 10.3 and **MariaDB** 10.5 in RHEL 8 include a Red Hat version of the **garbd** systemd service and a wrapper script for the **galera** package in the **/usr/lib/systemd/system/garbd.service** and **/usr/sbin/garbd-wrapper** files, respectively. Since RHEL 8.6, **MariaDB** 10.3 and **MariaDB** 10.5 now also provide an upstream version of these files located at **/usr/share/doc/galera/garb-systemd** and **/usr/share/doc/galera/garbd.service**.

Additional resources

- Galera Replication Library

- Galera Arbitrator

- mysql-wsrep project

9.2.9.3. Deploying MariaDB Galera Cluster

Prerequisites

- Install **MariaDB Galera Cluster** packages by selecting a stream (version) from the **mariadb** module and specifying the **galera** profile. For example:

  ```
  # {PackageManagerCommand} module install mariadb:10.3/galera
  ```

  As a result, the following packages are installed:

  - **mariadb-server-galera**

  - **mariadb-server**

  - **galera**
    The **mariadb-server-galera** package pulls the **mariadb-server** and **galera** packages as its dependency.

    For more information about which packages you need to install to build **MariaDB Galera Cluster**, see Components to build MariaDB Cluster .

- The **MariaDB** server replication configuration must be updated before the system is added to a cluster for the first time.
  The default configuration is distributed in the **/etc/my.cnf.d/galera.cnf** file.

  Before deploying **MariaDB Galera Cluster**, set the **wsrep_cluster_address** option in the **/etc/my.cnf.d/galera.cnf** file on all nodes to start with the following string:

  ```
  gcomm://
  ```

  - For the initial node, it is possible to set **wsrep_cluster_address** as an empty list:

    ```
    wsrep_cluster_address="gcomm://"
    ```

- For all other nodes, set **wsrep_cluster_address** to include an address to any node which is already a part of the running cluster. For example:

  ```
  wsrep_cluster_address="gcomm://10.0.0.10"
  ```

  For more information about how to set Galera Cluster address, see Galera Cluster Address .

Procedure

1. Bootstrap a first node of a new cluster by running the following wrapper on that node:

   ```
   # *galera_new_cluster*
   ```

   This wrapper ensures that the **MariaDB** server daemon (**mysqld**) runs with the **--wsrep-new-cluster** option. This option provides the information that there is no existing cluster to connect to. Therefore, the node creates a new UUID to identify the new cluster.

   > **NOTE**
   >
   > The **mariadb** service supports a systemd method for interacting with multiple MariaDB server processes. Therefore, in cases with multiple running MariaDB servers, you can bootstrap a specific instance by specifying the instance name as a suffix:
   >
   > ```
   > # galera_new_cluster mariadb@node1
   > ```

2. Connect other nodes to the cluster by running the following command on each of the nodes:

   ```
   # systemctl start mariadb
   ```

   As a result, the node connects to the cluster, and synchronizes itself with the state of the cluster.

Additional resources

- Getting started with MariaDB Galera Cluster .

9.2.9.4. Adding a new node to MariaDB Galera Cluster

To add a new node to **MariaDB Galera Cluster**, use the following procedure.

Note that you can also use this procedure to reconnect an already existing node.

Procedure

- On the particular node, provide an address to one or more existing cluster members in the **wsrep_cluster_address** option within the **[mariadb]** section of the **/etc/my.cnf.d/galera.cnf** configuration file :

  ```
  [mariadb]
  wsrep_cluster_address="gcomm://192.168.0.1"
  ```

When a new node connects to one of the existing cluster nodes, it is able to see all nodes in the cluster.

However, preferably list all nodes of the cluster in **wsrep_cluster_address**.

As a result, any node can join a cluster by connecting to any other cluster node, even if one or more cluster nodes are down. When all members agree on the membership, the cluster's state is changed. If the new node's state is different from the state of the cluster, the new node requests either an Incremental State Transfer (IST) or a State Snapshot Transfer (SST) to ensure consistency with the other nodes.

Additional resources

- [Getting started with MariaDB Galera Cluster](#)

- [Introduction to State Snapshot Transfers](#)

9.2.9.5. Restarting MariaDB Galera Cluster

If you shut down all nodes at the same time, you stop the cluster, and the running cluster no longer exists. However, the cluster's data still exist.

To restart the cluster, bootstrap a first node as described in [Configuring MariaDB Galera Cluster](#) .

> **WARNING**
>
> If the cluster is not bootstrapped, and **mysqld** on the first node is started with only the **systemctl start mariadb** command, the node tries to connect to at least one of the nodes listed in the **wsrep_cluster_address** option in the **/etc/my.cnf.d/galera.cnf** file. If no nodes are currently running, the restart fails.

Additional resources

- [Getting started with MariaDB Galera Cluster](#) .

9.3. USING MYSQL

The **MySQL** server is an open source fast and robust database server. This part describes how to install and configure **MySQL** on a RHEL system, how to back up **MySQL** data, how to migrate from an earlier **MySQL** version, and how to replicate a **MySQL**.

9.3.1. Getting started with MySQL

MySQL is a relational database that converts data into structured information and provides an SQL interface for accessing data. It includes multiple storage engines and plug-ins, as well as geographic information system (GIS) and JavaScript Object Notation (JSON) features.

This part describes:

- How to install the **MySQL** server in [Installing MySQL](#).

- How to adjust **MySQL** configuration in Configuring MySQL.

- How to back up **MySQL** data in Backing up MySQL data.

- How to migrate **MySQL** data in Migrating to a RHEL 8 version of MySQL 8.0 .

- How to replicate a **MySQL** database in Replicating MySQL.

9.3.2. Installing MySQL

In RHEL 8, the **MySQL 8.0** server is available as the **mysql:8.0** module stream.

To install **MySQL**, use the following procedure.

Procedure

1. Install **MySQL** server packages by selecting the **8.0** stream (version) from the **mysql** module and specifying the **server** profile:

   ```
   # yum module install mysql:8.0/server
   ```

2. Start the **mysqld** service:

   ```
   # systemctl start mysqld.service
   ```

3. Enable the **mysqld** service to start at boot:

   ```
   # systemctl enable mysqld.service
   ```

4. *Recommended:* To improve security when installing **MySQL**, run the following command:

   ```
   $ mysql_secure_installation
   ```

   The command launches a fully interactive script, which prompts for each step in the process. The script enables you to improve security in the following ways:

   - Setting a password for root accounts

   - Removing anonymous users

   - Disallowing remote root logins (outside the local host)

> **NOTE**
>
> The **MySQL** and **MariaDB** database servers cannot be installed in parallel in RHEL 8 due to conflicting RPM packages. Parallel installation of components is possible in Red Hat Software Collections for RHEL 7. In RHEL 8, different versions of database servers can be used in containers.

9.3.3. Configuring MySQL

To configure the **MySQL** server for networking, use the following procedure.

Procedure

1. Edit the **[mysqld]** section of the **/etc/my.cnf.d/mysql-server.cnf** file. You can set the following configuration directives:

   - **bind-address** – is the address on which the server listens. Possible options are:

       - a host name

       - an IPv4 address

       - an IPv6 address

   - **skip-networking** – controls whether the server listens for TCP/IP connections. Possible values are:

       - 0 – to listen for all clients

       - 1 – to listen for local clients only

   - **port** – the port on which MySQL listens for TCP/IP connections.

2. Restart the **mysqld** service:

   ```
   # systemctl restart mysqld.service
   ```

9.3.4. Backing up MySQL data

There are two main ways to back up data from a **MySQL** database in Red Hat Enterprise Linux 8:

- Logical backup

- Physical backup

Logical backup consists of the SQL statements necessary to restore the data. This type of backup exports information and records in plain text files.

The main advantage of logical backup over physical backup is portability and flexibility. The data can be restored on other hardware configurations, **MySQL** versions or Database Management System (DBMS), which is not possible with physical backups.

Note that logical backup can be performed if the **mysqld.service** is running. Logical backup does not include log and configuration files.

Physical backup consists of copies of files and directories that store the content.

Physical backup has the following advantages compared to logical backup:

- Output is more compact.

- Backup is smaller in size.

- Backup and restore are faster.

- Backup includes log and configuration files.

Note that physical backup must be performed when the **mysqld.service** is not running or all tables in the database are locked to prevent changes during the backup.

You can use one of the following **MySQL** backup approaches to back up data from a **MySQL** database:

- Logical backup with **mysqldump**

- File system backup

- Replication as a backup solution

9.3.4.1. Performing logical backup with mysqldump

The **mysqldump** client is a backup utility, which can be used to dump a database or a collection of databases for the purpose of a backup or transfer to another database server. The output of **mysqldump** typically consists of SQL statements to re-create the server table structure, populate it with data, or both. **mysqldump** can also generate files in other formats, including XML and delimited text formats, such as CSV.

To perform the **mysqldump** backup, you can use one of the following options:

- Back up one or more selected databases

- Back up all databases

- Back up a subset of tables from one database

Procedure

- To dump a single database, run:

  ```
  # mysqldump [options] --databases db_name > backup-file.sql
  ```

- To dump multiple databases at once, run:

  ```
  # mysqldump [options] --databases db_name1 [db_name2 ...] > backup-file.sql
  ```

- To dump all databases, run:

  ```
  # mysqldump [options] --all-databases > backup-file.sql
  ```

- To load one or more dumped full databases back into a server, run:

  ```
  # mysql < backup-file.sql
  ```

- To load a database to a remote **MySQL** server, run:

  ```
  # mysql --host=remote_host < backup-file.sql
  ```

- To dump a subset of tables from one database, add a list of the chosen tables at the end of the **mysqldump** command:

  ```
  # mysqldump [options] db_name [tbl_name ...] > backup-file.sql
  ```

- To load a subset of tables dumped from one database, run:

```
# mysql db_name < backup-file.sql
```

> **NOTE**
>
> The *db_name* database must exist at this point.

- To see a list of the options that **mysqldump** supports, run:

```
$ mysqldump --help
```

Additional resources

- Logical backup with mysqldump

9.3.4.2. Performing file system backup

To create a file system backup of **MySQL** data files, copy the content of the **MySQL** data directory to your backup location.

To back up also your current configuration or the log files, use the optional steps of the following procedure.

Procedure

1. Stop the **mysqld** service:

```
# systemctl stop mysqld.service
```

2. Copy the data files to the required location:

```
# cp -r /var/lib/mysql /backup-location
```

3. Optionally, copy the configuration files to the required location:

```
# cp -r /etc/my.cnf /etc/my.cnf.d /backup-location/configuration
```

4. Optionally, copy the log files to the required location:

```
# cp /var/log/mysql/* /backup-location/logs
```

5. Start the **mysqld** service:

```
# systemctl start mysqld.service
```

6. When loading the backed up data from the backup location to the **/var/lib/mysql** directory, ensure that **mysql:mysql** is an owner of all data in **/var/lib/mysql**:

```
# chown -R mysql:mysql /var/lib/mysql
```

9.3.4.3. Replication as a backup solution

Replication is an alternative backup solution for source servers. If a source server replicates to a replica server, backups can be run on the replica without any impact on the source. The source can still run while you shut down the replica and back the data up from the replica.

For instructions on how to replicate a **MySQL** database, see Replicating MySQL.

> **WARNING**
>
> Replication itself is not a sufficient backup solution. Replication protects source servers against hardware failures, but it does not ensure protection against data loss. It is recommended that you use any other backup solution on the replica together with this method.

Additional resources

- MySQL replication documentation

9.3.5. Migrating to a RHEL 8 version of MySQL 8.0

RHEL 7 contains **MariaDB 5.5** as the default implementation of a server from the **MySQL** databases family. The Red Hat Software Collections offering for RHEL 7 provides **MySQL 8.0** and several versions of **MariaDB**. RHEL 8 provides **MySQL 8.0**, **MariaDB 10.3**, and **MariaDB 10.5**.

This procedure describes migration from a Red Hat Software Collections version of **MySQL 8.0** to a RHEL 8 version of **MySQL 8.0** using the **mysql_upgrade** utility. The **mysql_upgrade** utility is provided by the **mysql-server** package.

> **NOTE**
>
> In the Red Hat Software Collections version of **MySQL**, the source data directory is **/var/opt/rh/rh-mysql80/lib/mysql/**. In RHEL 8, **MySQL** data is stored in the **/var/lib/mysql/** directory.

Prerequisites

- Before performing the upgrade, back up all your data stored in the **MySQL** databases.

Procedure

1. Ensure that the **mysql-server** package is installed on the RHEL 8 system:

   ```
   # yum install mysql-server
   ```

2. Ensure that the **mysqld** service is not running on either of the source and target systems at the time of copying data:

   ```
   # systemctl stop mysqld.service
   ```

3. Copy the data from the **/var/opt/rh/rh-mysql80/lib/mysql/** directory on the RHEL 7 source system to the **/var/lib/mysql/** directory on the RHEL 8 target system.

4. Set the appropriate permissions and SELinux context for copied files on the target system:

```
# restorecon -vr /var/lib/mysql
```

5. Ensure that **mysql:mysql** is an owner of all data in the **/var/lib/mysql** directory:

```
# chown -R mysql:mysql /var/lib/mysql
```

6. Start the **MySQL** server on the target system:

```
# systemctl start mysqld.service
```

Note: In earlier versions of **MySQL**, the **mysql_upgrade** command was needed to check and repair internal tables. This is now done automatically when you start the server.

9.3.6. Replicating MySQL

MySQL provides various configuration options for replication, ranging from basic to advanced. This section describes a transaction-based way to replicate in **MySQL** on freshly installed **MySQL** servers using global transaction identifiers (GTIDs). Using GTIDs simplifies transaction identification and consistency verification.

To set up replication in **MySQL**, you must:

- Configure a source server

- Configure a replica server

- Create a replication user on the source server

- Connect the replica server to the source server
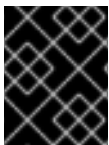
> **IMPORTANT**
>
> If you want to use existing **MySQL** servers for replication, you must first synchronize data. See the upstream documentation for more information.

9.3.6.1. Configuring a MySQL source server

This section describes configuration options required for a **MySQL** source server to properly run and replicate all changes made on the database server.

Prerequisites

- The source server is installed.

Procedure

1. Include the following options in the **/etc/my.cnf.d/mysql-server.cnf** file under the **[mysqld]** section:

- **bind-address=*source_ip_adress***
  This option is required for connections made from replicas to the source.

- **server-id=*id***
  The *id* must be unique.

- **log_bin=*path_to_source_server_log***
  This option defines a path to the binary log file of the **MySQL** source server. For example:
  log_bin=/var/log/mysql/mysql-bin.log.

- **gtid_mode=ON**
  This option enables global transaction identifiers (GTIDs) on the server.

- **enforce-gtid-consistency=ON**
  The server enforces GTID consistency by allowing execution of only statements that can be safely logged using a GTID.

- *Optional:* **binlog_do_db=*db_name***
  Use this option if you want to replicate only selected databases. To replicate more than one selected database, specify each of the databases separately:

  ```
  binlog_do_db=db_name1
  binlog_do_db=db_name2
  binlog_do_db=db_name3
  ```

- *Optional:* **binlog_ignore_db=*db_name***
  Use this option to exclude a specific database from replication.

2. Restart the **mysqld** service:

   ```
   # systemctl restart mysqld.service
   ```

9.3.6.2. Configuring a MySQL replica server

This section describes configuration options required for a **MySQL** replica server to ensure a successful replication.

Prerequisites

- The replica server is installed.

Procedure

1. Include the following options in the **/etc/my.cnf.d/mysql-server.cnf** file under the **[mysqld]** section:

   - **server-id=*id***
     The *id* must be unique.

   - **relay-log=*path_to_replica_server_log***
     The relay log is a set of log files created by the **MySQL** replica server during replication.

   - **log_bin=*path_to_replica_sever_log***

This option defines a path to the binary log file of the **MySQL** replica server. For example: **log_bin=/var/log/mysql/mysql-bin.log**.

This option is not required in a replica but strongly recommended.

- **gtid_mode=ON**
  This option enables global transaction identifiers (GTIDs) on the server.

- **enforce-gtid-consistency=ON**
  The server enforces GTID consistency by allowing execution of only statements that can be safely logged using a GTID.

- **log-replica-updates=ON**
  This option ensures that updates received from the source server are logged in the replica's binary log.

- **skip-replica-start=ON**
  This option ensures that the replica server does not start the replication threads when the replica server starts.

- *Optional:* **binlog_do_db=*db_name***
  Use this option if you want to replicate only certain databases. To replicate more than one database, specify each of the databases separately:

  ```
  binlog_do_db=db_name1
  binlog_do_db=db_name2
  binlog_do_db=db_name3
  ```

- *Optional:* **binlog_ignore_db=*db_name***
  Use this option to exclude a specific database from replication.

2. Restart the **mysqld** service:

   ```
   # systemctl restart mysqld.service
   ```

9.3.6.3. Creating a replication user on the MySQL source server

You must create a replication user and grant this user permissions required for replication traffic. This procedure shows how to create a replication user with appropriate permissions. Execute these steps only on the source server.

Prerequisites

- The source server is installed and configured as described in Configuring a MySQL source server.

Procedure

1. Create a replication user:

   ```
   mysql> CREATE USER 'replication_user'@'replica_server_ip' IDENTIFIED WITH
   mysql_native_password BY 'password';
   ```

2. Grant the user replication permissions:

```
mysql> GRANT REPLICATION SLAVE ON *.* TO 'replication_user'@'replica_server_ip';
```

3. Reload the grant tables in the **MySQL** database:

```
mysql> FLUSH PRIVILEGES;
```

4. Set the source server to read-only state:

```
mysql> SET @@GLOBAL.read_only = ON;
```

9.3.6.4. Connecting the replica server to the source server

On the **MySQL** replica server, you must configure credentials and the address of the source server. Use the following procedure to implement the replica server.

Prerequisites

- The source server is installed and configured as described in Configuring a MySQL source server.

- The replica server is installed and configured as described in Configuring a MySQL replica server.

- You have created a replication user. See Creating a replication user on the MySQL source server.

Procedure

1. Set the replica server to read-only state:

```
mysql> SET @@GLOBAL.read_only = ON;
```

2. Configure the replication source:

```
mysql> CHANGE REPLICATION SOURCE TO
    -> SOURCE_HOST='source_ip_address',
    -> SOURCE_USER='replication_user',
    -> SOURCE_PASSWORD='password',
    -> SOURCE_AUTO_POSITION=1;
```

3. Start the replica thread in the **MySQL** replica server:

```
mysql> START REPLICA;
```

4. Unset the read-only state on both the source and replica servers:

```
mysql> SET @@GLOBAL.read_only = OFF;
```

5. *Optional:* Inspect the status of the replica server for debugging purposes:

```
mysql> SHOW REPLICA STATUS\G;
```

NOTE

If the replica server fails to start or connect, you can skip a certain number of events following the binary log file position displayed in the output of the **SHOW MASTER STATUS** command. For example, skip the first event from the defined position:

> **mysql> SET GLOBAL SQL_SLAVE_SKIP_COUNTER=1;**

and try to start the replica server again.

6. *Optional:* Stop the replica thread in the replica server:

> mysql> STOP REPLICA;

9.3.6.5. Verification steps

1. Create an example database on the source server:

   > mysql> CREATE DATABASE *test_db_name*;

2. Verify that the ***test_db_name*** database replicates on the replica server.

3. Display status information about the binary log files of the MySQL server by executing the following command on either of the source or replica servers:

   > mysql> SHOW MASTER STATUS;

   The **Executed_Gtid_Set** column, which shows a set of GTIDs for transactions executed on the source, must not be empty.

NOTE

The same set of GTIDs is displayed in the **Executed_Gtid_Set** row when you use the **SHOW SLAVE STATUS** on the replica server.

9.3.6.6. Additional resources

- MySQL Replication documentation
- How To Set Up Replication in MySQL
- Replication with Global Transaction Identifiers

9.4. USING POSTGRESQL

The **PostgreSQL** server is an open source robust and highly-extensible database server based on the SQL language. This part describes how to install and configure **PostgreSQL** on a RHEL system, how to back up **PostgreSQL** data, and how to migrate from an earlier **PostgreSQL** version

9.4.1. Getting started with PostgreSQL

The **PostgreSQL** server provides an object-relational database system, which allows you to manage extensive datasets and a high number of concurrent users. For these reasons, the **PostgreSQL** servers can be used in clusters to manage high amounts of data.

The **PostgreSQL** server includes features for ensuring data integrity, building fault-tolerant environments, and building applications. It allows users to extend a database with users' own data types, custom functions, or code from different programming languages without the need to recompile the database.

This part describes:

- How to install **PostgreSQL** in Installing PostgreSQL.

- Users, roles, and privileges in Creating PostgreSQL users.

- How to adjust **PostgreSQL** configuration in Configuring PostgreSQL.

- How to enable Transport Layer Security (TLS) support on the **PostgreSQL** server in Configuring TLS encryption on a PostgreSQL server.

- How to back up your databases in Backing up PostgreSQL data.

- How to migrate to a later version of **PostgreSQL** in Migrating to a RHEL 8 version of PostgreSQL. One of the prerequisites of migration is performing a data backup.

9.4.2. Installing PostgreSQL

In RHEL 8, the **PostgreSQL** server is available in several versions, each provided by a separate stream:

- **PostgreSQL 10** - the default stream

- **PostgreSQL 9.6**

- **PostgreSQL 12** - available since RHEL 8.1.1

- **PostgreSQL 13** - available since RHEL 8.4
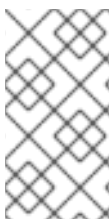
> **NOTE**
>
> By design, it is impossible to install more than one version (stream) of the same module in parallel. Thus you must choose only one of the available streams from the **postgresql** module. Parallel installation of components is possible in Red Hat Software Collections for RHEL 7. In RHEL 8, different versions of database servers can be used in containers.

To install **PostgreSQL**, use the following procedure.

Procedure

1. Install the **PostgreSQL** server packages by selecting a stream (version) from the **postgresql** module and specifying the server profile. For example:

   ```
   # yum module install postgresql:13/server
   ```

   The **postgres** superuser is created automatically.

2. Initialize the database cluster:

```
# postgresql-setup --initdb
```

Red Hat recommends storing the data in the default **/var/lib/pgsql/data** directory.

3. Start the **postgresql** service:

```
# systemctl start postgresql.service
```

4. Enable the **postgresql** service to start at boot:

```
# systemctl enable postgresql.service
```

For information about using module streams, see Installing, managing, and removing user-space components.

> **IMPORTANT**
>
> If you want to upgrade from an earlier **postgresql** stream within RHEL 8, follow both procedures described in Switching to a later stream and in Section 9.4.7, "Migrating to a RHEL 8 version of PostgreSQL".

9.4.3. Creating PostgreSQL users

PostgreSQL users are of the following types:

- The **postgres** UNIX system user – should be used only to run the PostgreSQL server and client applications, such as **pg_dump**. Do not use the **postgres** system user for any interactive work on **PostgreSQL** administration, such as database creation and user management.

- A database superuser – the default **postgres PostgreSQL** superuser is not related to the **postgres** system user. You can limit access of the **postgres** superuser in the **pg_hba.conf** file, otherwise no other permission limitations exist. You can also create other database superusers.

- A role with specific database access permissions:

  - A database user – has a permission to log in by default

  - A group of users – enables managing permissions for the group as a whole

Roles can own database objects (for example, tables and functions) and can assign object privileges to other roles using SQL commands.

Standard database management privileges include **SELECT**, **INSERT**, **UPDATE**, **DELETE**, **TRUNCATE**, **REFERENCES**, **TRIGGER**, **CREATE**, **CONNECT**, **TEMPORARY**, **EXECUTE**, and **USAGE**.

Role attributes are special privileges, such as **LOGIN**, **SUPERUSER**, **CREATEDB**, and **CREATEROLE**.
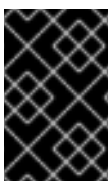
> **IMPORTANT**
>
> Red Hat recommends performing most tasks as a role that is not a superuser. A common practice is to create a role that has the **CREATEDB** and **CREATEROLE** privileges and use this role for all routine management of databases and roles.

Prerequisites

- The **PostgreSQL** server is installed.

- The database cluster is initialized.

Procedure

- To create a user, set a password for the user, and assign the user the **CREATEROLE** and **CREATEDB** permissions:

  > postgres=# CREATE USER *mydbuser* WITH PASSWORD '*mypasswd*' CREATEROLE CREATEDB;

  Replace *mydbuser* with the username and *mypasswd* with the user's password.

Additional resources

- PostgreSQL database roles

- PostgreSQL privileges

- Configuring PostgreSQL

Example 9.1. Initializing, creating, and connecting to a PostgreSQL database

This example demonstrates how to initialize a PostgreSQL database, create a database user with routine database management privileges, and how to create a database that is accessible from any system account through the database user with management privileges.

1. Install the PosgreSQL server:

   > # yum module install postgresql:13/server

2. Initialize the database cluster:

   > # postgresql-setup --initdb
   > * Initializing database in '/var/lib/pgsql/data'
   > * Initialized, logs are in /var/lib/pgsql/initdb_postgresql.log

3. Set the password hashing algorithm to **scram-sha-256**.

   a. In the **/var/lib/pgsql/data/postgresql.conf** file, change the following line:

      > #password_encryption = md5   # md5 or scram-sha-256

      to:

      > password_encryption = scram-sha-256

   b. In the **/var/lib/pgsql/data/pg_hba.conf** file, change the following line for the IPv4 local connections:

      > host   all   all   127.0.0.1/32   ident

      to:

```
host    all         all        127.0.0.1/32        scram-sha-256
```

4. Start the postgresql service:

```
# systemctl start postgresql.service
```

5. Log in as the system user named **postgres**:

```
# su - postgres
```

6. Start the **PostgreSQL** interactive terminal:

```
$ psql
psql (13.7)
Type "help" for help.

postgres=#
```

7. Optional: Obtain information about the current database connection:

```
postgres=# \conninfo
You are connected to database "postgres" as user "postgres" via socket in
"/var/run/postgresql" at port "5432".
```

8. Create a user named **mydbuser**, set a password for **mydbuser**, and assign **mydbuser** the **CREATEROLE** and **CREATEDB** permissions:

```
postgres=# CREATE USER mydbuser WITH PASSWORD 'mypasswd' CREATEROLE
CREATEDB;
CREATE ROLE
```

The **mydbuser** user now can perform routine database management operations: create databases and manage user indexes.

9. Log out of the interactive terminal by using the **\q** meta command:

```
postgres=# \q
```

10. Log out of the **postgres** user session:

```
$ logout
```

11. Log in to the **PostgreSQL** terminal as **mydbuser**, specify the hostname, and connect to the default **postgres** database, which was created during initialization:

```
# psql -U mydbuser -h 127.0.0.1 -d postgres
Password for user mydbuser:
Type the password.
psql (13.7)
Type "help" for help.

postgres=>
```

12. Create a database named **mydatabase**:

    ```
    postgres=> CREATE DATABASE mydatabase;
    CREATE DATABASE
    postgres=>
    ```

13. Log out of the session:

    ```
    postgres=# \q
    ```

14. Connect to mydatabase as **mydbuser**:

    ```
    # psql -U mydbuser -h 127.0.0.1 -d mydatabase
    Password for user mydbuser:
    psql (13.7)
    Type "help" for help.
    mydatabase=>
    ```

15. Optional: Obtain information about the current database connection:

    ```
    mydatabase=> \conninfo
    You are connected to database "mydatabase" as user "mydbuser" on host "127.0.0.1" at
    port "5432".
    ```

9.4.4. Configuring PostgreSQL

In a **PostgreSQL** database, all data and configuration files are stored in a single directory called a database cluster. Red Hat recommends storing all data, including configuration files, in the default **/var/lib/pgsql/data/** directory.

PostgreSQL configuration consists of the following files:

- **postgresql.conf** – is used for setting the database cluster parameters.

- **postgresql.auto.conf** – holds basic **PostgreSQL** settings similarly to **postgresql.conf**. However, this file is under the server control. It is edited by the **ALTER SYSTEM** queries, and cannot be edited manually.

- **pg_ident.conf** – is used for mapping user identities from external authentication mechanisms into the **PostgreSQL** user identities.

- **pg_hba.conf** – is used for configuring client authentication for **PostgreSQL** databases.

To change the **PostgreSQL** configuration, use the following procedure.

Procedure

1. Edit the respective configuration file, for example, **/var/lib/pgsql/data/postgresql.conf**.

2. Restart the **postgresql** service so that the changes become effective:

    ```
    # systemctl restart postgresql.service
    ```

> **Example 9.2. Configuring PostgreSQL database cluster parameters**
>
> This example shows basic settings of the database cluster parameters in the **/var/lib/pgsql/data/postgresql.conf** file.
>
> ```
> # This is a comment
> log_connections = yes
> log_destination = 'syslog'
> search_path = '"$user", public'
> shared_buffers = 128MB
> password_encryption = scram-sha-256
> ```

> **Example 9.3. Setting client authentication in PostgreSQL**
>
> This example demonstrates how to set client authentication in the **/var/lib/pgsql/data/pg_hba.conf** file.
>
> ```
> # TYPE    DATABASE      USER       ADDRESS            METHOD
> local    all          all                            trust
> host     postgres     all        192.168.93.0/24     ident
> host     all          all        .example.com        scram-sha-256
> ```

9.4.5. Configuring TLS encryption on a PostgreSQL server

By default, **PostgreSQL** uses unencrypted connections. For more secure connections, you can enable Transport Layer Security (TLS) support on the **PostgreSQL** server and configure your clients to establish encrypted connections.

Prerequisites

- The **PostgreSQL** server is installed.

- The database cluster is initialized.

Procedure

1. Install the OpenSSL library:

   ```
   # yum install openssl
   ```

2. Generate a TLS certificate and a key:

   ```
   # openssl req -new -x509 -days 365 -nodes -text -out server.crt \
       -keyout server.key -subj "/CN=dbhost.yourdomain.com"
   ```

   Replace *dbhost.yourdomain.com with* your database host and domain name.

3. Copy your signed certificate and your private key to the required locations on the database server:

   ```
   # cp server.{key,crt} /var/lib/pgsql/data/.
   ```

4. Change the owner and group ownership of the signed certificate and your private key to the **postgres** user:

```
# chown postgres:postgres /var/lib/pgsql/data/server.{key,crt}
```

5. Restrict the permissions for your private key so that it is readable only by the owner:

```
# chmod 0400 /var/lib/pgsql/data/server.key
```

6. Set the password hashing algorithm to **scram-sha-256** by changing the following line in the **/var/lib/pgsql/data/postgresql.conf** file:

```
#password_encryption = md5          # md5 or scram-sha-256
```

to:

```
password_encryption = scram-sha-256
```

7. Configure PostgreSQL to use SSL/TLS by changing the following line in the **/var/lib/pgsql/data/postgresql.conf** file:

```
#ssl = off
```

to:

```
ssl=on
```

8. Restrict access to all databases to accept only connections from clients using TLS by changing the following line for the IPv4 local connections in the **/var/lib/pgsql/data/pg_hba.conf** file:

```
host  all  all  127.0.0.1/32  ident
```

to:

```
hostssl  all  all  127.0.0.1/32  scram-sha-256
```

Alternatively, you can restrict access for a single database and a user by adding the following new line:

```
hostssl mydatabase mydbuser 127.0.0.1/32  scram-sha-256
```

Replace *mydatabase* with the database name and *mydbuser* with the username.

9. Make the changes effective by restarting the **postgresql** service:

```
# systemctl restart postgresql.service
```

Verification

- To manually verify that the connection is encrypted:

  1. Connect to the **PostgreSQL** database as the *mydbuser* user, specify the host name and the

1. Connect to the **PostgreSQL** database as the *mydbuser* user, specify the host name and the database name:

```
$ psql -U mydbuser -h 127.0.0.1 -d mydatabase
Password for user mydbuser:
```

Replace *mydatabase* with the database name and *mydbuser* with the username.

2. Obtain information about the current database connection:

```
mydbuser=> \conninfo
You are connected to database "mydatabase" as user "mydbuser" on host "127.0.0.1" at
port "5432".
SSL connection (protocol: TLSv1.3, cipher: TLS_AES_256_GCM_SHA384, bits: 256,
compression: off)
```

- You can write a simple application that verifies whether a connection to **PostgreSQL** is encrypted. This example demonstrates such an application written in C that uses the **libpq** client library, which is provided by the **libpq-devel** package:

```
#include <stdio.h>
#include <stdlib.h>
#include <libpq-fe.h>

int main(int argc, char* argv[])
{
//Create connection
PGconn* connection = PQconnectdb("hostaddr=127.0.0.1 password=mypassword port=5432
dbname=mydatabase user=mydbuser");

if (PQstatus(connection) ==CONNECTION_BAD)
   {
   printf("Connection error\n");
   PQfinish(connection);
   return -1; //Execution of the program will stop here
   }
   printf("Connection ok\n");
   //Verify TLS
   if (PQsslInUse(connection)){
    printf("TLS in use\n");
    printf("%s\n", PQsslAttribute(connection,"protocol"));
   }
   //End connection
   PQfinish(connection);
   printf("Disconnected\n");
   return 0;
}
```

Replace *mypassword* with the password, *mydatabase* with the database name, and *mydbuser* with the username.

NOTE

You must load the **pq** libraries for compilation by using the **-lpq** option. For example, to compile the application by using the GCC compiler:

```
$ gcc source_file.c -lpq -o myapplication
```

where the *source_file.c* contains the example code above, and *myapplication* is the name of your application for verifying secured **PostgreSQL** connection.

Example 9.4. Initializing, creating, and connecting to a PostgreSQL database using TLS encryption

This example demonstrates how to initialize a PostgreSQL database, create a database user and a database, and how to connect to the database using a secured connection.

1. Install the PosgreSQL server:

   ```
   # yum module install postgresql:13/server
   ```

2. Initialize the database cluster:

   ```
   # postgresql-setup --initdb
   * Initializing database in '/var/lib/pgsql/data'
   * Initialized, logs are in /var/lib/pgsql/initdb_postgresql.log
   ```

3. Install the OpenSSL library:

   ```
   # yum install openssl
   ```

4. Generate a TLS certificate and a key:

   ```
   # openssl req -new -x509 -days 365 -nodes -text -out server.crt \
     -keyout server.key -subj "/CN=dbhost.yourdomain.com"
   ```

   Replace *dbhost.yourdomain.com with* your database host and domain name.

5. Copy your signed certificate and your private key to the required locations on the database server:

   ```
   # cp server.{key,crt} /var/lib/pgsql/data/.
   ```

6. Change the owner and group ownership of the signed certificate and your private key to the **postgres** user:

   ```
   # chown postgres:postgres /var/lib/pgsql/data/server.{key,crt}
   ```

7. Restrict the permissions for your private key so that it is readable only by the owner:

   ```
   # chmod 0400 /var/lib/pgsql/data/server.key
   ```

8. Set the password hashing algorithm to **scram-sha-256**. In the **/var/lib/pgsql/data/postgresql.conf** file, change the following line:

> #password_encryption = md5   # md5 or scram-sha-256

to:

> password_encryption = scram-sha-256

9. Configure PostgreSQL to use SSL/TLS. In the **/var/lib/pgsql/data/postgresql.conf** file, change the following line:

> #ssl = off

to:

> ssl=on

10. Start the **postgresql** service:

> # systemctl start postgresql.service

11. Log in as the system user named **postgres**:

> # su - postgres

12. Start the **PostgreSQL** interactive terminal as the **postgres** user:

```
$ psql -U postgres
psql (13.7)
Type "help" for help.

postgres=#
```

13. Create a user named **mydbuser** and set a password for **mydbuser**:

```
postgres=# CREATE USER mydbuser WITH PASSWORD 'mypasswd';
CREATE ROLE
postgres=#
```

14. Create a database named **mydatabase**:

```
postgres=# CREATE DATABASE mydatabase;
CREATE DATABASE
postgres=#
```

15. Grant all permissions to the **mydbuser** user:

```
postgres=# GRANT ALL PRIVILEGES ON DATABASE mydatabase TO mydbuser;
GRANT
postgres=#
```

16. Log out of the interactive terminal:

> postgres=# \q

17. Log out of the **postgres** user session:

    ```
    $ logout
    ```

18. Restrict access to all databases to accept only connections from clients using TLS by changing the following line for the IPv4 local connections in the **/var/lib/pgsql/data/pg_hba.conf** file:

    ```
    host  all  all  127.0.0.1/32  ident
    ```

    to:

    ```
    hostssl  all  all  127.0.0.1/32  scram-sha-256
    ```

19. Make the changes effective by restarting the **postgresql** service:

    ```
    # systemctl restart postgresql.service
    ```

20. Connect to the **PostgreSQL** database as the **mydbuser** user, specify the host name and the database name:

    ```
    $ psql -U mydbuser -h 127.0.0.1 -d mydatabase
    Password for user mydbuser:
    psql (13.7)
    SSL connection (protocol: TLSv1.3, cipher: TLS_AES_256_GCM_SHA384, bits: 256,
    compression: off)
    Type "help" for help.

    mydatabase=>
    ```

9.4.6. Backing up PostgreSQL data

To back up **PostgreSQL** data, use one of the following approaches:

SQL dump

See Backing up with SQL dump.

File system level backup

See File system level backup .

Continuous archiving

See Continuous archiving.

9.4.6.1. Backing up PostgreSQL data with an SQL dump

The SQL dump method is based on generating a dump file with SQL commands. When a dump is uploaded back to the database server, it recreates the database in the same state as it was at the time of the dump.

The SQL dump is ensured by the following **PostgreSQL** client applications:

- **pg_dump** dumps a single database without cluster-wide information about roles or tablespaces

- **pg_dumpall** dumps each database in a given cluster and preserves cluster-wide data, such as role and tablespace definitions.

By default, the **pg_dump** and **pg_dumpall** commands write their results into the standard output. To store the dump in a file, redirect the output to an SQL file. The resulting SQL file can be either in a text format or in other formats that allow for parallelism and for more detailed control of object restoration.

You can perform the SQL dump from any remote host that has access to the database.

9.4.6.1.1. Advantages and disadvantages of an SQL dump

An SQL dump has the following advantages compared to other **PostgreSQL** backup methods:

- An SQL dump is the only **PostgreSQL** backup method that is not server version-specific. The output of the **pg_dump** utility can be reloaded into later versions of **PostgreSQL**, which is not possible for file system level backups or continuous archiving.

- An SQL dump is the only method that works when transferring a database to a different machine architecture, such as going from a 32-bit to a 64-bit server.

- An SQL dump provides internally consistent dumps. A dump represents a snapshot of the database at the time **pg_dump** began running.

- The **pg_dump** utility does not block other operations on the database when it is running.

A disadvantage of an SQL dump is that it takes more time compared to file system level backup.

9.4.6.1.2. Performing an SQL dump using pg_dump

To dump a single database without cluster-wide information, use the **pg_dump** utility.

Prerequisites

- You must have read access to all tables that you want to dump. To dump the entire database, you must run the commands as the **postgres** superuser or a user with database administrator privileges.

Procedure

- Dump a database without cluster-wide information:

  ```
  $ pg_dump dbname > dumpfile
  ```

To specify which database server **pg_dump** will contact, use the following command-line options:

- The **-h** option to define the host.
  The default host is either the local host or what is specified by the **PGHOST** environment variable.

- The **-p** option to define the port.
  The default port is indicated by the **PGPORT** environment variable or the compiled-in default.

9.4.6.1.3. Performing an SQL dump using pg_dumpall

To dump each database in a given database cluster and to preserve cluster-wide data, use the **pg_dumpall** utility.

Prerequisites

- You must run the commands as the **postgres** superuser or a user with database administrator privileges.

Procedure

- Dump all databases in the database cluster and preserve cluster-wide data:
  $ **pg_dumpall > dumpfile**

To specify which database server **pg_dumpall** will contact, use the following command-line options:

- The **-h** option to define the host.
  The default host is either the local host or what is specified by the **PGHOST** environment variable.

- The **-p** option to define the port.
  The default port is indicated by the **PGPORT** environment variable or the compiled-in default.

- The **-l** option to define the default database.
  This option enables you to choose a default database different from the **postgres** database created automatically during initialization.

9.4.6.1.4. Restoring a database dumped using pg_dump

To restore a database from an SQL dump that you dumped using the **pg_dump** utility, follow the steps below.

Prerequisites

- You must run the commands as the **postgres** superuser or a user with database administrator privileges.

Procedure

1. Create a new database:

   ```
   $ createdb dbname
   ```

2. Verify that all users who own objects or were granted permissions on objects in the dumped database already exist. If such users do not exist, the restore fails to recreate the objects with the original ownership and permissions.

3. Run the **psql** utility to restore a text file dump created by the **pg_dump** utility:

   ```
   $ psql dbname < dumpfile
   ```

   where **dumpfile** is the output of the **pg_dump** command. To restore a non-text file dump, use the **pg_restore** utility instead:

   ```
   $ pg_restore non-plain-text-file
   ```

9.4.6.1.5. Restoring databases dumped using pg_dumpall

To restore data from a database cluster that you dumped using the **pg_dumpall** utility, follow the steps below.

Prerequisites

- You must run the commands as the **postgres** superuser or a user with database administrator privileges.

Procedure

1. Ensure that all users who own objects or were granted permissions on objects in the dumped databases already exist. If such users do not exist, the restore fails to recreate the objects with the original ownership and permissions.

2. Run the **psql** utility to restore a text file dump created by the **pg_dumpall** utility:

   ```
   $ psql < dumpfile
   ```

   where **dumpfile** is the output of the **pg_dumpall** command.

9.4.6.1.6. Performing an SQL dump of a database on another server

Dumping a database directly from one server to another is possible because **pg_dump** and **psql** can write to and read from pipes.

Procedure

- To dump a database from one server to another, run:

  ```
  $ pg_dump -h host1 dbname | psql -h host2 dbname
  ```

9.4.6.1.7. Handling SQL errors during restore

By default, **psql** continues to execute if an SQL error occurs, causing the database to restore only partially.

To change the default behavior, use one of the following approaches when restoring a dump.

Prerequisites

- You must run the commands as the **postgres** superuser or a user with database administrator privileges.

Procedure

- Make **psql** exit with an exit status of 3 if an SQL error occurs by setting the **ON_ERROR_STOP** variable:
  ```
  $ psql --set ON_ERROR_STOP=on dbname < dumpfile
  ```

- Specify that the whole dump is restored as a single transaction so that the restore is either fully completed or canceled.

- When restoring a text file dump using the **psql** utility:

  ```
  $ psql -1
  ```

- When restoring a non-text file dump using the **pg_restore** utility:

  ```
  $ pg_restore -e
  ```

  Note that when using this approach, even a minor error can cancel a restore operation that has already run for many hours.

9.4.6.1.8. Additional resources

- [PostgreSQL Documentation – SQL dump](#)

9.4.6.2. Backing up PostgreSQL data with a file system level backup

To create a file system level backup, copy **PostgreSQL** database files to another location. For example, you can use any of the following approaches:

- Create an archive file using the **tar** utility.

- Copy the files to a different location using the **rsync** utility.

- Create a consistent snapshot of the data directory.

9.4.6.2.1. Advantages and limitations of file system backing up

File system level backing up has the following advantage compared to other **PostgreSQL** backup methods:

- File system level backing up is usually faster than an SQL dump.

File system level backing up has the following limitations compared to other **PostgreSQL** backup methods:

- This backing up method is not suitable when you want to upgrade from RHEL 7 to RHEL 8 and migrate your data to the upgraded system. File system level backup is specific to an architecture and a RHEL major version. You can restore your data on your RHEL 7 system if the upgrade is not successful but you cannot restore the data on a RHEL 8 system.

- The database server must be shut down before backing up and restoring data.

- Backing up and restoring certain individual files or tables is impossible. Backing up a file system works only for complete backing up and restoring of an entire database cluster.

9.4.6.2.2. Performing file system level backing up

To perform file system level backing up, use the following procedure.

Procedure

1. Choose the location of a database cluster and initialize this cluster:

   ```
   # postgresql-setup --initdb
   ```

2. Stop the postgresql service:

```
# systemctl stop postgresql.service
```

3. Use any method to create a file system backup, for example a **tar** archive:

```
$ tar -cf backup.tar /var/lib/pgsql/data
```

4. Start the postgresql service:

```
# systemctl start postgresql.service
```

Additional resources

- [PostgreSQL Documentation – file system level backup](#)

9.4.6.3. Backing up PostgreSQL data by continuous archiving

9.4.6.3.1. Introduction to continuous archiving

PostgreSQL records every change made to the database's data files into a write ahead log (WAL) file that is available in the **pg_wal/** subdirectory of the cluster's data directory. This log is intended primarily for a crash recovery. After a crash, the log entries made since the last checkpoint can be used for restoring the database to a consistency.

The continuous archiving method, also known as an online backup, combines the WAL files with a copy of the database cluster in the form of a base backup performed on a running server or a file system level backup.

If a database recovery is needed, you can restore the database from the copy of the database cluster and then replay log from the backed up WAL files to bring the system to the current state.

With the continuous archiving method, you must keep a continuous sequence of all archived WAL files that extends at minimum back to the start time of your last base backup. Therefore the ideal frequency of base backups depends on:

- The storage volume available for archived WAL files.

- The maximum possible duration of data recovery in situations when recovery is necessary. In cases with a long period since the last backup, the system replays more WAL segments, and the recovery therefore takes more time.

> **NOTE**
>
> You cannot use **pg_dump** and **pg_dumpall** SQL dumps as a part of a continuous archiving backup solution. SQL dumps produce logical backups and do not contain enough information to be used by a WAL replay.

To perform a database backup and restore using the continuous archiving method, follow these instructions:

1. Set up and test your procedure for archiving WAL files – see WAL archiving.

2. Perform a base backup - see base backup.

To restore your data, follow instructions in Restoring database with continuous archiving.

9.4.6.3.2. Advantages and disadvantages of continuous archiving

Continuous archiving has the following advantages compared to other **PostgreSQL** backup methods:

- With the continuous backup method, it is possible to use a base backup that is not entirely consistent because any internal inconsistency in the backup is corrected by the log replay. Therefore you can perform a base backup on a running **PostgreSQL** server.

- A file system snapshot is not needed; **tar** or a similar archiving utility is sufficient.

- Continuous backup can be achieved by continuing to archive the WAL files because the sequence of WAL files for the log replay can be indefinitely long. This is particularly valuable for large databases.

- Continuous backup supports point-in-time recovery. It is not necessary to replay the WAL entries to the end. The replay can be stopped at any point and the database can be restored to its state at any time since the base backup was taken.

- If the series of WAL files are continuously available to another machine that has been loaded with the same base backup file, it is possible to restore the other machine with a nearly-current copy of the database at any point.

Continuous archiving has the following disadvantages compared to other **PostgreSQL** backup methods:

- Continuous backup method supports only restoration of an entire database cluster, not a subset.

- Continuous backup requires extensive archival storage.

9.4.6.3.3. Setting up WAL archiving

A running **PostgreSQL** server produces a sequence of write ahead log (WAL) records. The server physically divides this sequence into WAL segment files, which are given numeric names that reflect their position in the WAL sequence. Without WAL archiving, the segment files are reused and renamed to higher segment numbers.

When archiving WAL data, the contents of each segment file are captured and saved at a new location before the segment file is reused. You have multiple options where to save the content, such as an NFS-mounted directory on another machine, a tape drive, or a CD.

Note that WAL records do not include changes to configuration files.

To enable WAL archiving, use the following procedure.

Procedure

1. In the **/var/lib/pgsql/data/postgresql.conf** file:

   a. Set the **wal_level** configuration parameter to **replica** or higher.

   b. Set the **archive_mode** parameter to **on**.

c. Specify the shell command in the **archive_command** configuration parameter. You can use the **cp** command, another command, or a shell script.

2. Restart the **postgresql** service to enable the changes:

```
# systemctl restart postgresql.service
```

3. Test your archive command and ensure it does not overwrite an existing file and that it returns a non-zero exit status if it fails.

4. To protect your data, ensure that the segment files are archived into a directory that does not have group or world read access.

> **NOTE**
>
> The archive command is executed only on completed WAL segments. A server that generates little WAL traffic can have a substantial delay between the completion of a transaction and its safe recording in archive storage. To limit how old unarchived data can be, you can:
>
> - Set the **archive_timeout** parameter to force the server to switch to a new WAL segment file with a given frequency.
>
> - Use the **pg_switch_wal** parameter to force a segment switch to ensure that a transaction is archived immediately after it finishes.

Example 9.5. Shell command for archiving WAL segments

This example shows a simple shell command you can set in the **archive_command** configuration parameter.

The following command copies a completed segment file to the required location:

```
archive_command = 'test ! -f /mnt/server/archivedir/%f && cp %p /mnt/server/archivedir/%f'
```

where the **%p** parameter is replaced by the relative path to the file to archive and the **%f** parameter is replaced by the file name.

This command copies archivable WAL segments to the **/mnt/server/archivedir/** directory. After replacing the **%p** and **%f** parameters, the executed command looks as follows:

```
test ! -f /mnt/server/archivedir/00000001000000A900000065 && cp
pg_wal/00000001000000A900000065 /mnt/server/archivedir/00000001000000A900000065
```

A similar command is generated for each new file that is archived.

Additional resources

- [PostgreSQL 10 Documentation](#).

9.4.6.3.4. Making a base backup

You can create a base backup in several ways. This section describes the simplest way of performing a base backup using the **pg_basebackup** utility on a running **PostgreSQL** server.

The base backup process creates a backup history file that is stored into the WAL archive area and is named after the first WAL segment file that you need for the base backup.

The backup history file is a small text file containing the starting and ending times, and WAL segments of the backup. If you used the label string to identify the associated dump file, you can use the backup history file to determine which dump file to restore.
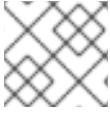
> **NOTE**
>
> Consider keeping several backup sets to be certain that you can recover your data.

To perform a base backup, use the following procedure.

Prerequisites

- You must run the commands as the **postgres** superuser, a user with database administrator privileges, or another user with at least **REPLICATION** permissions.

- You must keep all the WAL segment files generated during and after the base backup.

Procedure

1. Use the **pg_basebackup** utility to perform the base backup.

   - To create a base backup as individual files (plain format):

     ```
     $ pg_basebackup -D backup_directory -Fp
     ```

     Replace *backup_directory* with your desired backup location.

     If you use tablespaces and perform the base backup on the same host as the server, you must also use the **--tablespace-mapping** option, otherwise the backup will fail upon an attempt to write the backup to the same location.

   - To create a base backup as a **tar** archive (**tar** and compressed format):

     ```
     $ pg_basebackup -D backup_directory -Ft -z
     ```

     Replace *backup_directory* with your desired backup location.

     To restore such data, you must manually extract the files in the correct locations.

2. After the base backup process is complete, safely archive the copy of the database cluster and the WAL segment files used during the backup, which are specified in the backup history file.

3. Delete WAL segments numerically lower than the WAL segment files used in the base backup because these are older than the base backup and no longer needed for a restore.

To specify which database server **pg_basebackup** will contact, use the following command-line options:

- The **-h** option to define the host.

  The default host is either the local host or a host specified by the **PGHOST** environment

The default host is either the local host or a host specified by the **PGHOST** environment variable.

- The **-p** option to define the port.
  The default port is indicated by the **PGPORT** environment variable or the compiled-in default.

Additional resources

- PostgreSQL Documentation – base backup

- PostgreSQL Documentation – pg_basebackup utility

9.4.6.3.5. Restoring the database using a continuous archive backup

To restore a database using a continuous backup, use the following procedure.

Procedure

1. Stop the server:

   ```
   # systemctl stop postgresql.service
   ```

2. Copy the necessary data to a temporary location.
   Preferably, copy the whole cluster data directory and any tablespaces. Note that this requires enough free space on your system to hold two copies of your existing database.

   If you do not have enough space, save the contents of the cluster's **pg_wal** directory, which can contain logs that were not archived before the system went down.

3. Remove all existing files and subdirectories under the cluster data directory and under the root directories of any tablespaces you are using.

4. Restore the database files from your base backup.
   Ensure that:

   - The files are restored with the correct ownership (the database system user, not **root**).

   - The files are restored with the correct permissions.

   - The symbolic links in the **pg_tblspc**/ subdirectory are restored correctly.

5. Remove any files present in the **pg_wal**/ subdirectory.
   These files resulted from the base backup and are therefore obsolete. If you did not archive **pg_wal**/, recreate it with proper permissions.

6. Copy any unarchived WAL segment files that you saved in step 2 into **pg_wal**/.

7. Create the **recovery.conf** recovery command file in the cluster data directory and specify the shell command in the **restore_command** configuration parameter. You can use the **cp** command, another command, or a shell script. For example:

   ```
   restore_command = 'cp /mnt/server/archivedir/%f "%p"'
   ```

8. Start the server:

> # **systemctl start postgresql.service**

The server will enter the recovery mode and proceed to read through the archived WAL files that it needs.

If the recovery is terminated due to an external error, the server can be restarted and it will continue the recovery. When the recovery process is completed, the server renames **recovery.conf** to **recovery.done**. This prevents the server from accidental re-entering the recovery mode after it starts normal database operations.

9. Check the contents of the database to verify that the database has recovered into the required state.
    If the database has not recovered into the required state, return to step 1. If the database has recovered into the required state, allow the users to connect by restoring the client authentication configuration in the **pg_hba.conf** file.

For more information about restoring using the continuous backup, see PostgreSQL Documentation.

9.4.6.3.6. Additional resources

- Continuous archiving method

9.4.7. Migrating to a RHEL 8 version of PostgreSQL

Red Hat Enterprise Linux 7 contains **PostgreSQL 9.2** as the default version of the **PostgreSQL** server. In addition, several versions of **PostgreSQL** are provided as Software Collections for RHEL 7.

Red Hat Enterprise Linux 8 provides **PostgreSQL 10** (the default **postgresql** stream), **PostgreSQL 9.6**, **PostgreSQL 12**, and **PostgreSQL 13**.

Users of **PostgreSQL** on Red Hat Enterprise Linux can use two migration paths for the database files:

- Fast upgrade using the pg_upgrade utility

- Dump and restore upgrade

The fast upgrade method is quicker than the dump and restore process. However, in certain cases, the fast upgrade does not work, and you can only use the dump and restore process. Such cases include:

- Cross-architecture upgrades

- Systems using the **plpython** or **plpython2** extensions. Note that RHEL 8 AppStream repository includes only the **postgresql-plpython3** package, not the **postgresql-plpython2** package.

- Fast upgrade is not supported for migration from Red Hat Software Collections versions of **PostgreSQL**.

As a prerequisite for migration to a later version of **PostgreSQL**, back up all your **PostgreSQL** databases.

Dumping the databases and performing backup of the SQL files is required for the dump and restore process and recommended for the fast upgrade method.

Before migrating to a later version of **PostgreSQL**, see the upstream compatibility notes for the version of **PostgreSQL** to which you want to migrate, as well as for all skipped **PostgreSQL** versions between the one you are migrating from and the target version.

9.4.7.1. Fast upgrade using the pg_upgrade utility

During a fast upgrade, you must copy binary data files to the **/var/lib/pgsql/data/** directory and use the **pg_upgrade** utility.

You can use this method for migrating data:

- From the RHEL 7 system version of **PostgreSQL 9.2** to the RHEL 8 version of **PostgreSQL 10**

- From the RHEL 8 version of **PostgreSQL 10** to the RHEL 8 version of **PostgreSQL 12** or **PostgreSQL 13**

If you want to upgrade from an earlier **postgresql** stream within RHEL 8, follow the procedure described in Switching to a later stream and then migrate your **PostgreSQL** data.

For migrating between other combinations of **PostgreSQL** versions within RHEL, and for migration from the Red Hat Software Collections versions of **PostgreSQL** to RHEL, use Dump and restore upgrade.

The following procedure describes migration from the RHEL 7 system version of **Postgreql 9.2** to a RHEL 8 version of **PostgreSQL** using the fast upgrade method.

Prerequisites

- Before performing the upgrade, back up all your data stored in the **PostgreSQL** databases. By default, all data is stored in the **/var/lib/pgsql/data/** directory on both the RHEL 7 and RHEL 8 systems.

Procedure

1. On the RHEL 8 system, enable the stream (version) to which you wish to migrate:

   ```
   # yum module enable postgresql:stream
   ```

   Replace *stream* with the selected version of the **PostgreSQL** server.

   You can omit this step if you want to use the default stream, which provides **PostgreSQL 10**.

2. On the RHEL 8 system, install the **postgresql-server** and **postgresql-upgrade** packages:

   ```
   # yum install postgresql-server postgresql-upgrade
   ```

   Optionally, if you used any **PostgreSQL** server modules on RHEL 7, install them also on the RHEL 8 system in two versions, compiled both against **PostgreSQL 9.2** (installed as the **postgresql-upgrade** package) and the target version of **PostgreSQL** (installed as the **postgresql-server** package). If you need to compile a third-party **PostgreSQL** server module, build it both against the **postgresql-devel** and **postgresql-upgrade-devel** packages.

3. Check the following items:

   - Basic configuration: On the RHEL 8 system, check whether your server uses the default **/var/lib/pgsql/data** directory and the database is correctly initialized and enabled. In addition, the data files must be stored in the same path as mentioned in the **/usr/lib/systemd/system/postgresql.service** file.

- **PostgreSQL** servers: Your system can run multiple **PostgreSQL** servers. Ensure that the data directories for all these servers are handled independently.

- **PostgreSQL** server modules: Ensure that the **PostgreSQL** server modules that you used on RHEL 7 are installed on your RHEL 8 system as well. Note that plug-ins are installed in the **/usr/lib64/pgsql/** directory (or in the **/usr/lib/pgsql/** directory on 32-bit systems).

4. Ensure that the **postgresql** service is not running on either of the source and target systems at the time of copying data.

   ```
   # systemctl stop postgresql.service
   ```

5. Copy the database files from the source location to the **/var/lib/pgsql/data/** directory on the RHEL 8 system.

6. Perform the upgrade process by running the following command as the **PostgreSQL** user:

   ```
   # postgresql-setup --upgrade
   ```

   This launches the **pg_upgrade** process in the background.

   In case of failure, **postgresql-setup** provides an informative error message.

7. Copy the prior configuration from **/var/lib/pgsql/data-old** to the new cluster.
   Note that the fast upgrade does not reuse the prior configuration in the newer data stack and the configuration is generated from scratch. If you want to combine the old and new configurations manually, use the *.conf files in the data directories.

8. Start the new **PostgreSQL** server:

   ```
   # systemctl start postgresql.service
   ```

9. Run the **analyze_new_cluster.sh** script located in the **PostgreSQL** home directory:

   ```
   su postgres -c '~/analyze_new_cluster.sh'
   ```

10. If you want the new **PostgreSQL** server to be automatically started on boot, run:

    ```
    # systemctl enable postgresql.service
    ```

9.4.7.2. Dump and restore upgrade

When using the dump and restore upgrade, you must dump all databases contents into an SQL file dump file.

Note that the dump and restore upgrade is slower than the fast upgrade method and it may require some manual fixing in the generated SQL file.

You can use this method for migrating data from:

- The Red Hat Enterprise Linux 7 system version of **PostgreSQL 9.2**

- Any earlier Red Hat Enterprise Linux 8 version of **PostgreSQL**

- An earlier or equal version of **PostgreSQL** from Red Hat Software Collections:

- PostgreSQL 9.2 (no longer supported)

- PostgreSQL 9.4 (no longer supported)

- PostgreSQL 9.6 (no longer supported)

- PostgreSQL 10

- PostgreSQL 12

On RHEL 7 and RHEL 8 systems, **PostgreSQL** data is stored in the **/var/lib/pgsql/data/** directory by default. In case of Red Hat Software Collections versions of **PostgreSQL**, the default data directory is **/var/opt/rh/collection_name/lib/pgsql/data/** (with the exception of **postgresql92**, which uses the **/opt/rh/postgresql92/root/var/lib/pgsql/data/** directory).

If you want to upgrade from an earlier **postgresql** stream within RHEL 8, follow the procedure described in Switching to a later stream and then migrate your **PostgreSQL** data.

To perform the dump and restore upgrade, change the user to **root**.

The following procedure describes migration from the RHEL 7 system version of **Postgreql 9.2** to a RHEL 8 version of **PostgreSQL**.

Procedure

1. On your RHEL 7 system, start the **PostgreSQL 9.2** server:

   ```
   # systemctl start postgresql.service
   ```

2. On the RHEL 7 system, dump all databases contents into the **pgdump_file.sql** file:

   ```
   su - postgres -c "pg_dumpall > ~/pgdump_file.sql"
   ```

3. Verify that the databases were dumped correctly:

   ```
   su - postgres -c 'less "$HOME/pgdump_file.sql"'
   ```

   As a result, the path to the dumped sql file is displayed: **/var/lib/pgsql/pgdump_file.sql**.

4. On the RHEL 8 system, enable the stream (version) to which you wish to migrate:

   ```
   # yum module enable postgresql:stream
   ```

   Replace *stream* with the selected version of the **PostgreSQL** server.

   You can omit this step if you want to use the default stream, which provides **PostgreSQL 10**.

5. On the RHEL 8 system, install the **postgresql-server** package:

   ```
   # yum install postgresql-server
   ```

   Optionally, if you used any **PostgreSQL** server modules on RHEL 7, install them also on the RHEL 8 system. If you need to compile a third-party **PostgreSQL** server module, build it against the **postgresql-devel** package.

6. On the RHEL 8 system, initialize the data directory for the new **PostgreSQL** server:

   ```
   # postgresql-setup --initdb
   ```

7. On the RHEL 8 system, copy the **pgdump_file.sql** into the **PostgreSQL** home directory, and check that the file was copied correctly:

   ```
   su - postgres -c 'test -e "$HOME/pgdump_file.sql" && echo exists'
   ```

8. Copy the configuration files from the RHEL 7 system:

   ```
   su - postgres -c 'ls -1 $PGDATA/*.conf'
   ```

   The configuration files to be copied are:

   - **/var/lib/pgsql/data/pg_hba.conf**

   - **/var/lib/pgsql/data/pg_ident.conf**

   - **/var/lib/pgsql/data/postgresql.conf**

9. On the RHEL 8 system, start the new **PostgreSQL** server:

   ```
   # systemctl start postgresql.service
   ```

10. On the RHEL 8 system, import data from the dumped sql file:

    ```
    su - postgres -c 'psql -f ~/pgdump_file.sql postgres'
    ```

> **NOTE**
>
> When upgrading from a Red Hat Software Collections version of **PostgreSQL**, adjust the commands to include **scl enable collection_name.** For example, to dump data from the **rh-postgresql96** Software Collection, use the following command:
>
> ```
> su - postgres -c 'scl enable rh-postgresql96 "pg_dumpall > ~/pgdump_file.sql"'
> ```

CHAPTER 10. DEPLOYING AND CONFIGURING MAIL TRANSPORT AGENTS

Mail transport agent (MTA) transports email messages between hosts using SMTP. The delivery of email messages can include several MTAs while reaching their destination. During the process of message delivery, the usage of MTA depends on the MTAs or the access configuration of the network.

RHEL provides two MTAs: Sendmail and Postfix. Postfix replaced Sendmail as the preferred MTA in RHEL 7.

10.1. POSTFIX

A Postfix transfer agent handles all processes related to mail delivery. Postfix is available through the **BaseOS** repository. However, **postfix-mysql** or **postfix-ldap** subpackages are built from the postfix source RPM that is available through the **Appstream** repository.

Postfix configuration files are stored in the **/etc/postfix/** directory.

Following is a list of the commonly used files:

- access — it is used for access control and specifies hosts that are allowed to connect to Postfix.

- main.cf — it is the global Postfix configuration file, and specifies many configuration options.

- master.cf — it specifies Postfix interaction with various processes to accomplish mail delivery.

- transport — it maps email addresses to relay hosts. The aliases file is a configurable list required by the mail protocol that describes user ID aliases. It is in the **/etc** directory and is shared between Postfix and Sendmail.

10.2. INSTALLING POSTFIX

If the mail server package is not selected during the system installation, Postfix will not be available by default. Perform the following steps to install Postfix:

Prerequisites

- Registering your system

- Disabling and removing Sendmail

  - To disable and remove Sendmail

    ```
    # yum remove sendmail
    ```

- Configuring firewall for sending and receiving emails

Procedure

1. To install postfix

   ```
   # yum install postfix
   ```

2. To start and enable the postfix service

```
# systemctl start postfix

# systemctl enable postfix
```

Verification

- Verify if the postfix service is running

  ```
  $ rpm -qa | grep postfix
  ```

- Restart the postfix service, in the following scenarios:

  - if the output is stopped/waiting or not running

  - after changing any options in the configuration files under the /etc/postfix/ directory in order for those changes to take effect

    ```
    # service postfix restart
    ```

Additional resources

- **/etc/postfix/main.cf** configuration file.

10.3. CONFIGURING POSTFIX

The **main.cf** is the global Postfix configuration file and specifies many configuration options.

Procedure

Following are few options you can add in the **/etc/postfix/main.cf** file to configure Postfix:

1. myhostname:- replace **host.domain.tld** with the mail server's hostname.

   For example:

   myhostname = *mail.example.com*

2. mydomain:- replace **domain.tld** with the domain mail server.

   For example:

   mydomain = *example.com*

   > **NOTE**
   >
   > With correct DNS setups, the myhostname and mydomain options should be autodetected and set automatically without user intervention.

3. mail_spool_directory:- allows specifying the location of the mailbox files.

> For example:
>
> mail_spool_directory = /var/mail

4. mynetworks:– add the list of valid and trusted remote SMTP servers. This option should include only local network IP addresses or networks separated by commas or whitespaces. Adding local network addresses avoids unauthorized access to mail servers.

5. Uncomment the inet_interfaces = all line. This option allows Postfix to be accessible from the internet. With its default configuration, it will only receive emails from the local machine.

6. Comment the inet_interfaces = localhost line. This option allows Postfix to be accessible from the internet. With its default configuration, it will only receive emails from the local machine.

7. Restart the postfix service.

> # systemctl reload postfix

Verification

- To verify an email communication between local users on the system

> # echo "This is a test message" | mail -s <SUBJECT> <name@mydomain.com>

Press Ctrl+D to send the message.

- To verify the open relay is not active on your new mail server, send an email from your mail server to a domain that your new mail server doesn't accept mail:

> hello *mydomain.com*
> mail from: *name@mydomain.com*
> rcpt to: *name@notmydomain.com*
>
> 554 Relay access denied - the server is not going to relay.
>
> 250 OK or similar- the server is going to relay.
>
> The 'rcpt to' option should only accept mail addressed to addresses @mydomain.com
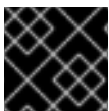
> **IMPORTANT**
>
> In case of errors, check the `**/var/log/maillog** file.

Additional resources

- **/etc/postfix/main.cf** configuration file.

10.4. SECURING THE POSTFIX SERVICE

Postfix is a mail transfer agent (MTA) that uses the Simple Mail Transfer Protocol (SMTP) to deliver electronic messages between other MTAs and to email clients or delivery agents. Although MTAs can encrypt traffic between one another, they might not do so by default. You can also mitigate risks to various attacks by changing setting to more secure values.

10.4.1. Reducing Postfix network-related security risks

To reduce the risk of attackers invading your system through the network, perform as many of the following tasks as possible.

- Do not share the **/var/spool/postfix/** mail spool directory on a Network File System (NFS) shared volume. NFSv2 and NFSv3 do not maintain control over user and group IDs. Therefore, if two or more users have the same UID, they can receive and read each other's mail, which is a security risk.

  > **NOTE**
  >
  > This rule does not apply to NFSv4 using Kerberos, because the **SECRPC_GSS** kernel module does not use UID-based authentication. However, to reduce the security risks, you should not put the mail spool directory on NFS shared volumes.

- To reduce the probability of Postfix server exploits, mail users must access the Postfix server using an email program. Do not allow shell accounts on the mail server, and set all user shells in the **/etc/passwd** file to **/sbin/nologin** (with the possible exception of the **root** user).

- To protect Postfix from a network attack, it is set up to only listen to the local loopback address by default. You can verify this by viewing the **inet_interfaces = localhost** line in the **/etc/postfix/main.cf** file. This ensures that Postfix only accepts mail messages (such as **cron** job reports) from the local system and not from the network. This is the default setting and protects Postfix from a network attack. To remove the localhost restriction and allow Postfix to listen on all interfaces, set the **inet_interfaces** parameter to **all** in **/etc/postfix/main.cf**.

10.4.2. Postfix configuration options for limiting DoS attacks

An attacker can flood the server with traffic, or send information that triggers a crash, causing a denial of service (DoS) attack. You can configure your system to reduce the risk of such attacks by setting limits in the **/etc/postfix/main.cf** file. You can change the value of the existing directives or you can add new directives with custom values in the *<directive>* = *<value>* format.

Use the following list of directives for limiting a DoS attack:

smtpd_client_connection_rate_limit

   This directive limits the maximum number of connection attempts any client can make to this service per time unit. The default value is **0**, which means a client can make as many connections per time unit as Postfix can accept. By default, the directive excludes clients in trusted networks.

anvil_rate_time_unit

   This directive is a time unit to calculate the rate limit. The default value is **60** seconds.

smtpd_client_event_limit_exceptions

   This directive excludes clients from the connection and rate limit commands. By default, the directive excludes clients in trusted networks.

smtpd_client_message_rate_limit

   This directive defines the maximum number of message deliveries from client to request per time unit (regardless of whether or not Postfix actually accepts those messages).

default_process_limit

   This directive defines the default maximum number of Postfix child processes that provide a given service. You can ignore this rule for specific services in the **master.cf** file. By default, the value is **100**.

queue_minfree

This directive defines the minimum amount of free space required to receive mail in the queue file system. The directive is currently used by the Postfix SMTP server to decide if it accepts any mail at all. By default, the Postfix SMTP server rejects **MAIL FROM** commands when the amount of free space is less than 1.5 times the **message_size_limit**. To specify a higher minimum free space limit, specify a **queue_minfree** value that is at least 1.5 times the **message_size_limit**. By default, the **queue_minfree** value is **0**.

header_size_limit

This directive defines the maximum amount of memory in bytes for storing a message header. If a header is large, it discards the excess header. By default, the value is **102400** bytes.

message_size_limit

This directive defines the maximum size of a message including the envelope information in bytes. By default, the value is **10240000** bytes.

10.4.3. Configuring Postfix to use SASL

Postfix supports Simple Authentication and Security Layer (SASL) based SMTP Authentication (AUTH). SMTP AUTH is an extension of the Simple Mail Transfer Protocol. Currently, the Postfix SMTP server supports the SASL implementations in the following ways:

Dovecot SASL

The Postfix SMTP server can communicate with the Dovecot SASL implementation using either a UNIX-domain socket or a TCP socket. Use this method if Postfix and Dovecot applications are running on separate machines.

Cyrus SASL

When enabled, SMTP clients must authenticate with the SMTP server using an authentication method supported and accepted by both the server and the client.

Prerequisites

- The **dovecot** package is installed on the system

Procedure

1. Set up Dovecot:

   a. Include the following lines in the **/etc/dovecot/conf.d/10-master.conf** file:

   ```
   service auth {
     unix_listener /var/spool/postfix/private/auth {
       mode = 0660
       user = postfix
       group = postfix
     }
   }
   ```

   The previous example uses UNIX-domain sockets for communication between Postfix and Dovecot. The example also assumes default Postfix SMTP server settings, which include the mail queue located in the **/var/spool/postfix/** directory, and the application running under the **postfix** user and group.

   b. Optional: Set up Dovecot to listen for Postfix authentication requests through TCP:

```
service auth {
  inet_listener {
     port = port-number
  }
}
```

c. Specify the method that the email client uses to authenticate with Dovecot by editing the **auth_mechanisms** parameter in **/etc/dovecot/conf.d/10-auth.conf** file:

```
auth_mechanisms = plain login
```

The **auth_mechanisms** parameter supports different plaintext and non-plaintext authentication methods.

2. Set up Postfix by modifying the **/etc/postfix/main.cf** file:

a. Enable SMTP Authentication on the Postfix SMTP server:

```
smtpd_sasl_auth_enable = yes
```

b. Enable the use of Dovecot SASL implementation for SMTP Authentication:

```
smtpd_sasl_type = dovecot
```

c. Provide the authentication path relative to the Postfix queue directory. Note that the use of a relative path ensures that the configuration works regardless of whether the Postfix server runs in **chroot** or not:

```
smtpd_sasl_path = private/auth
```

This step uses UNIX-domain sockets for communication between Postfix and Dovecot.

To configure Postfix to look for Dovecot on a different machine in case you use TCP sockets for communication, use configuration values similar to the following:

```
smtpd_sasl_path = inet: ip-address : port-number
```

In the previous example, replace the *ip-address* with the IP address of the Dovecot machine and *port-number* with the port number specified in Dovecot's **/etc/dovecot/conf.d/10-master.conf** file.

d. Specify SASL mechanisms that the Postfix SMTP server makes available to clients. Note that you can specify different mechanisms for encrypted and unencrypted sessions.

```
smtpd_sasl_security_options = noanonymous, noplaintext
smtpd_sasl_tls_security_options = noanonymous
```

The previous directives specify that during unencrypted sessions, no anonymous authentication is allowed and no mechanisms that transmit unencrypted user names or passwords are allowed. For encrypted sessions that use TLS, only non-anonymous authentication mechanisms are allowed.

Additional resources

- Postfix SMTP server policy – SASL mechanism properties

- Postfix and Dovecot SASL

- Configuring SASL authentication in the Postfix SMTP server

- Postfix SMTP server policy – SASL mechanism properties

- Postfix and Dovecot SASL

- Configuring SASL authentication in the Postfix SMTP server

CHAPTER 11. INSTALLING AND CONFIGURING DOVECOT FOR IMAP AND POP3

The **imap-login** and **pop3-login** processes that implement the IMAP and POP3 protocols are spawned by the master **dovecot** daemon included in the Dovecot package. The use of IMAP and POP is configured through the **/etc/dovecot/dovecot.conf**; by default, dovecot runs IMAP and POP3 together with their secure versions using SSL.

Procedure

To install and configure dovecot to use IMAP, complete the following steps:

1. To install Dovecot

   ```
   # yum install dovecot
   ```

2. To enable and start the Dovecot service.

   ```
   # systemctl enable dovecot

   # systemctl start dovecot
   ```

   The **dovecot. conf** is the main configuration file. Perform the following steps to add few options in the **/etc/dovecot/dovecot.conf** file:

   - listen:- Allows to set the IP addresses to listen to the services. For IPv4 addresses use the asterisk (*) and for IPV6 addresses use a colon (: :)

     ```
     For example:

     listen = *, : :
     ```

   - protocols:- Allows specifying the type of protocol.

     ```
     For example:

     protocols = imap, pop3
     ```

   - mail_location:- Sets the sender's mail location. By default, this option is empty, as Dovecot locates the mail automatically. Following is the format of the mail_location option:

   - mailbox-format:- *<path>* [: key = *<value>* ...]

   - Make the changes operational for the current session

     ```
     # systemctl restart dovecot
     ```

Verification

- Verify the Dovecot status

  ```
  # doveadm instance list
  ```

IMPORTANT

To check the logs, run the following command:

```
# journalctl -u dovecot -b
```

CHAPTER 12. SECURING DOVECOT

Dovecot contains self-signed SSL certificates in the **/etc/dovecot/conf.d/10-ssl.conf** file. Because Dovecot does not have CA certificates, you receive a warning message when you connect to the service. Ensure that you open the default SMTP, IMAP, SSL/TLS IMAP, and POP3 ports, using the following the command:

```
# firewall-cmd --permanent --add-port=110/tcp --add-port=995/tcp
# firewall-cmd --permanent --add-port=143/tcp --add-port=993/tcp
# firewall-cmd --reload
```
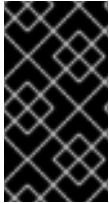
> **IMPORTANT**
>
> To check logs for problems related to Dovecot, use the following **journalctl** command:
>
> ```
> # journalctl -u dovecot -b
> ```

Additional resources

- Configuring the Dovecot mail server in the Securing networks document.

CHAPTER 13. CONFIGURING FIREWALL FOR SENDING AND RECEIVING EMAILS

Configure the firewall for sending and receiving emails using the following steps:

Procedure

1. To add the service

   # firewall-cmd --permanent --add-service=*servicename*

   Replace the *servicename* with any of the services in the **/etc/services**. For example, smtp, submission.

2. Reload the service for the change to take effect

   # systemctl reload firewalld

CHAPTER 14. SECURING EMAIL COMMUNICATION USING SSL

You can secure email communication using self signed certification. SSL certification can be done in 2 ways:

- by applying to a Certificate Authority (CA) for an SSL certificate

- by creating a self-signed certificate.

Procedure

Perform the following steps to create a self-signed SSL certificate for IMAP or POP

1. Edit the certificate parameters in the **/etc/pki/dovecot/dovecot-openssl.cnf** as you prefer, and type the following command:

   ```
   # rm -f certs/dovecot.pem private/dovecot.pem
   # /usr/libexec/dovecot/mkcert.sh
   ```

2. Ensure you have the following configurations in your /etc/dovecot/conf.d/10-ssl.conf file:

   ```
   ssl_cert = </etc/pki/dovecot/certs/dovecot.pem
   ssl_key = </etc/pki/dovecot/private/dovecot.pem
   ```

3. Execute the following command to restart the dovecot daemon:

   ```
   # systemctl restart dovecot
   ```

CHAPTER 15. CONFIGURING PRINTING

Printing on Red Hat Enterprise Linux 8 is based on the Common Unix Printing System (CUPS).

This documentation describes how to configure your system to be able to operate as a CUPS server.

15.1. ACTIVATING THE CUPS SERVICE

This section describes how to activate the **cups** service on your system.

Prerequisites

- The **cups** package, which is available in the Appstream repository, must be installed on your system:

  ```
  # yum install cups
  ```

Procedure

1. Start the **cups** service:

   ```
   # systemctl start cups
   ```

2. Configure the **cups** service to be automatically started at boot time:

   ```
   # systemctl enable cups
   ```

3. Optionally, check the status of the **cups** service:

   ```
   $ systemctl status cups
   ```

15.2. PRINT SETTINGS TOOLS

To achieve various tasks related to printing, you can choose one of the following tools:

- **CUPS web user interface (UI)**

- **GNOME Control center**

> **WARNING**
>
> The **Print Settings** configuration tool, which was used in Red Hat Enterprise Linux 7, is no longer available.

Tasks that you can achieve by using tools mentioned above include:

- Adding and configuring a new printer

- Maintaining printer configuration

- Managing printer classes

Note that this documentation covers only printing in **CUPS web user interface (UI)** If you want to print using **GNOME Control center**, you need to have a GUI available. For more information about printing using **GNOME Control center**, see Handling printing using GNOME .

15.3. ACCESSING AND CONFIGURING THE CUPS WEB UI

This section describes accessing the **CUPS web user interface**(web UI) and configuring it to be able to manage printing through this interface.

Procedure

1. Allow the CUPS server to listen to connections from the network by setting **Port 631** in the **/etc/cups/cupsd.conf** file:

   #Listen localhost:631
   Port 631

   > **WARNING**
   >
   > Enabling the CUPS server to listen on port 631 opens this port for any address accessible by the server. Therefore, use this setting only in local networks that are unreachable from an external network. Red Hat doesn't recommend configuring CUPS server on a publicly accessible server.

2. Allow your system to access the CUPS server by including the following directive in the **/etc/cups/cupsd.conf** file:

   <Location />
   Allow from <your_ip_address>
   Order allow,deny
   </Location>

   where *<your_ip_address>* is the real IP address of your system. You can also use regular expressions for subnets.

> **WARNING**
>
> The CUPS configuration offers the **Allow from all** directive in the
> *<Location>* tags, but Red Hat recommends using this directive only in
> trusted networks. The setup **Allow from all** enables access for all users who
> can connect to the server through port 631. If you set the **Port** directive to
> 631, and the server is accessible from an outside network, anyone on the
> Internet can access the CUPS service on your system.

3. Restart the cups.service:

   ```
   # systemctl restart cups
   ```

4. Open your browser, and go to **http://<IP_address_of_the_CUPS_server>:631/**.
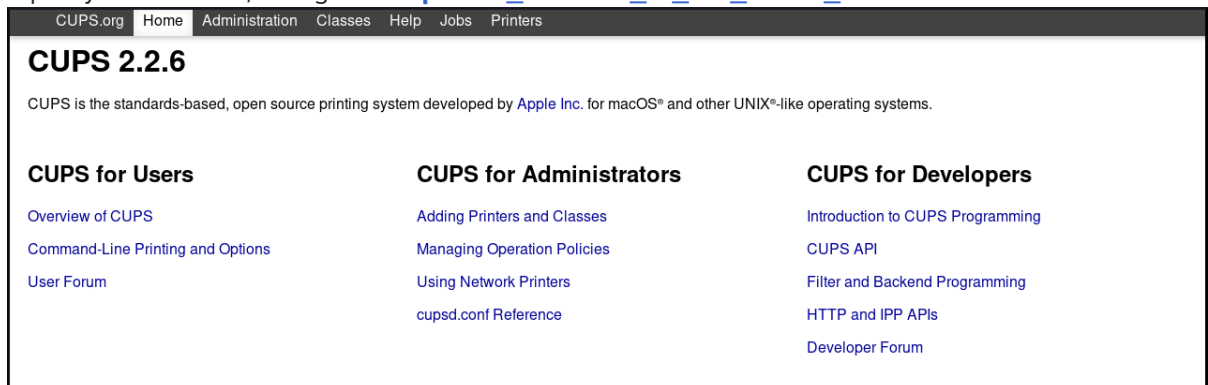
All menus except for the **Administration** menu are now available.

If you click on the **Administration** menu, you receive the **Forbidden** message:
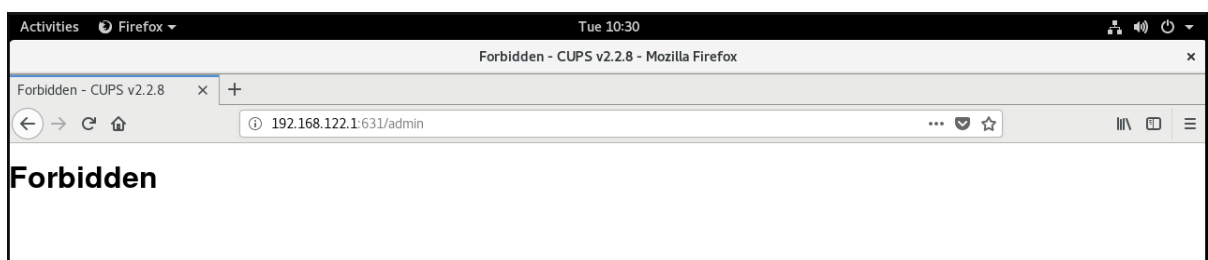
To acquire the access to the **Administration** menu, follow the instructions in Acquiring administration access to the CUPS web UI.

15.4. ACQUIRING ADMINISTRATION ACCESS TO THE CUPS WEB UI

This section describes how to acquire administration access to the **CUPS web UI**

Procedure

1. To be able to access the **Administration** menu in the **CUPS web UI**, include the following lines
   in the **/etc/cups/cupsd.conf** file:

```
<Location /admin>
Allow from <your_ip_address>
Order allow,deny
</Location>
```
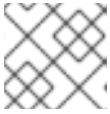
> **NOTE**
>
> Replace **<your_ip_address>** with the real IP address of your system.

2. To be able to access configuration files in the **CUPS web UI**, include the following in the **/etc/cups/cupsd.conf** file:

```
<Location /admin/conf>
AuthType Default
Require user @SYSTEM
Allow from <your_ip_address>
Order allow,deny
</Location>
```

> **NOTE**
>
> Replace **<your_ip_address>** with the real IP address of your system.

3. To be able to access log files in the **CUPS web UI**, include the following in the **/etc/cups/cupsd.conf** file:

```
<Location /admin/log>
AuthType Default
Require user @SYSTEM
Allow from <your_ip_address>
Order allow,deny
</Location>
```
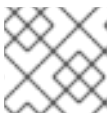
> **NOTE**
>
> Replace **<your_ip_address>** with the real IP address of your system.

4. To specify the use of encryption for authenticated requests in the **CUPS web UI**, include **DefaultEncryption** in the **/etc/cups/cupsd.conf** file:

```
DefaultEncryption IfRequested
```

With this setting, you will receive an authentication window to enter the username of a user allowed to add printers when you attempt to access the **Administration** menu. However, there are also other options on how to set **DefaultEncryption**. For more details, see the **cupsd.conf** man page.

5. Restart the **cups** service:

```
# systemctl restart cups
```

> ⚠️ **WARNING**
>
> If you do not restart the **cups** service, the changes in **/etc/cups/cupsd.conf** will not be applied. Consequently, you will not be able to obtain administration access to the **CUPS web UI**.

Additional resources

- The **cupsd.conf** man page

15.5. CONFIGURING DRIVERLESS PRINTING

As an administrator, you can configure driverless printing to use printers or remote CUPS queues without any special software.

RHEL 8 provides driverless printing support for following driverless standards:

- **IPP Everywhere model** in CUPS supports AirPrint, IPP Everywhere and Wi-Fi Direct standards.

- **Driverless model** in cups-filters supports the same standards as CUPS and in addition also PCLm document format.

These standards use the Internet Printing Protocol (IPP) 2.0 or higher to communicate the printer setup, and eliminate the need to install specific drivers for specific printers. To use the printer without a specific driver, you need to have a printer, which supports one of the driverless standards. To determine if your printer supports a driverless standard, choose one of the following options:

- Refer to the printer specification, and search for a driverless standard support or ask your vendor.

- Search for certified printers.

- Determine the driverless support based on the attributes of a printer with the ipptool command.

To install a print queue on the client with IPP Everywhere model, which points to a queue on the print server, you need to have both your remote print server and your client with RHEL 8.6 installation or newer.

> **NOTE**
>
> You can verify the driverless support based on the attributes of a print server with the ipptool command.

15.5.1. Determining printer attributes using ipptool

To determine if your printer or print server supports a driverless standard, you can inspect your printer attributes using the **ipptool** command available in the **ipptool** package.

Procedure

- Display attributes of a printer or a print server:

> $ **ipptool -tv <URI> get-printer-attributes.test**

NOTE

Replace <URI> with the URI of your printer, for example
ipp://<hostname_or_IP_address>:631/ipp/print for printers or
ipp://<hostname_or_IP_address>:631/printers/<remote_print_queue> for
remote print queues from print servers.

Your printer or print server supports driverless printing if:

- the **ipp-version-supported** attribute contains **2.0** or higher for IPP protocol 2.0, and

- the **document-format-supported** attribute contains one of the supported document formats
  listed in the driverless printing standards.

15.5.2. Adding a driverless printer in CUPS web UI

Since RHEL 8.6, you can add a driverless printer in the CUPS web UI and use it to print directly from an
application to network printers or print servers using CUPS, without installing any specific drivers or
software for specific printers.

Prerequisites

- You have administration access to the **CUPS web UI** as described in Acquiring administration
  access to CUPS web UI.

- Your printer or print server has the IPP Everywhere standard implementation.

- Open IPP port: port **631** for IPP, or port **443** for secure printing with IPPS.

- Enable the **ipp** and **ipp-client** communication in the firewall of the print server.

- If your destination is another CUPS server, allow remote access on the remote server, or if you
  are using a network printer, open the web user interface, search for IPP related settings: IPP or
  AirPrint, and enable those settings.

Procedure

1. Start the **CUPS web UI** as described in Accessing and configuring the CUPS.

2. In your browser, go to **localhost:631** and select the **Administration** tab.

3. Under **Printers** click **Add printer**.
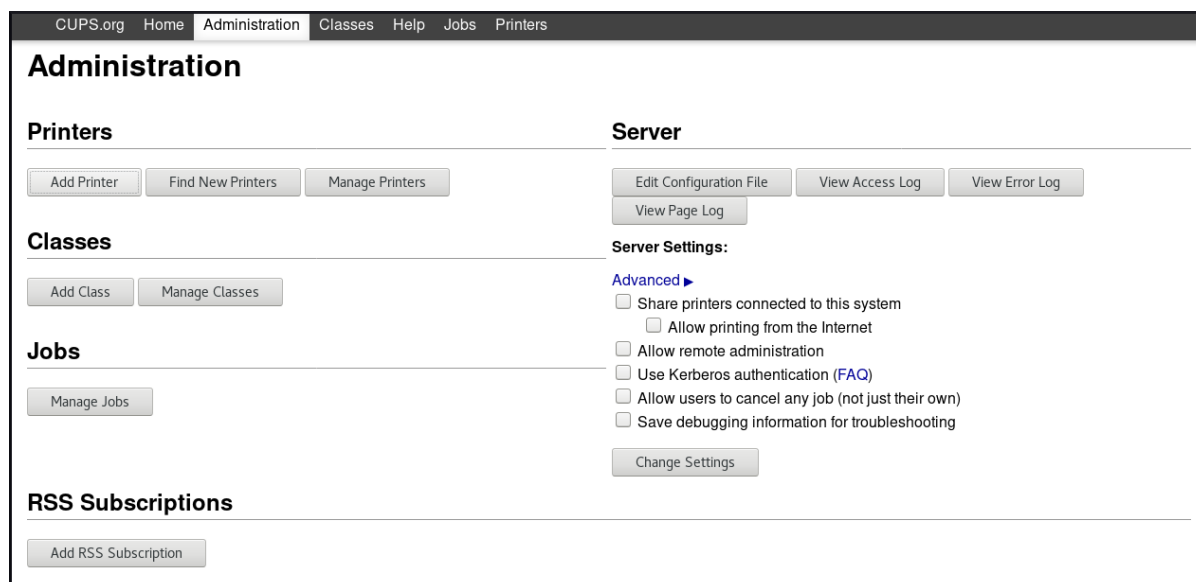
4. Authenticate with your username and password:
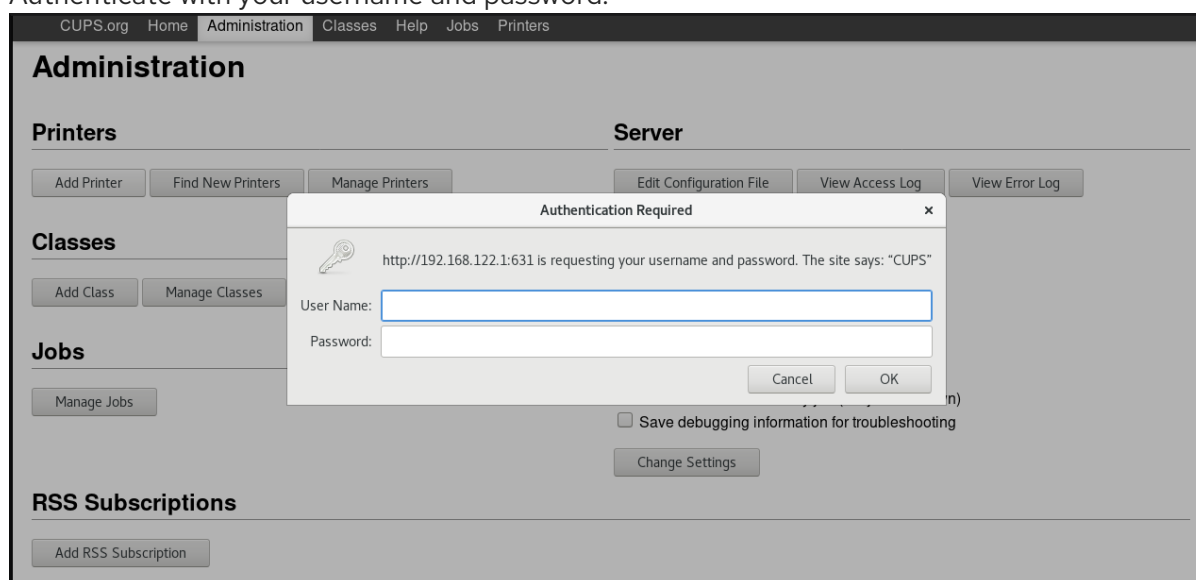
> **IMPORTANT**
>
> To be able to add a new printer by using the **CUPS web UI**, you must authenticate as a user who belongs to a group defined by the **SystemGroup** directive in **/etc/cups/cups-files**. The default groups are:
>
> - root
>
> - sys
>
> - wheel

5. In the **Administrator** tab, under **Add Printer**, select one of the options:

- **Internet Printing Protocol (ipp)** or

- **Internet Printing Protocol (ipps)** option, and click **Continue**.

6. In the **Connection** field, enter the URI of your device and click **Continue**.

NOTE

The URI consists of the following parts:

- protocol **ipp://** or **ipps://** if printer or print server supports encryption,

- hostname or IP address of the printer,

- port,

- resource part **/ipp/print** for printers, or **/printers/<remote_queue_name>** for remote CUPS queues.
  For example: **ipp://myprinter.mydomain:631/ipp/print** or **ipp://myserver.mydomain:631/printers/myqueue**.

7. Add details about the new printer: name, description and location. To set a printer to be shared over the network, check the **Share This Printer** checkbox.

CUPS.org   Home   Administration   Classes   Help   Jobs   Printers

Add Printer

Add Printer

Name: hp_laserjet
(May contain any printable characters except "/", "#", and space)

Description: HP LaserJet color
(Human-readable description such as "HP LaserJet with Duplexer")

Location: Lab 1
(Human-readable location such as "Lab 1")

Connection: ipp://<hostname_or_IP>:631/ipp/print

Sharing: ☑ Share This Printer

Continue

NOTE

'Name' is the only required field, other fields are optional.

8. From the **Make** dropdown menu, select the printer manufacturer, and click **Continue**.

9. To proceed with the installation of a driverless printer, from the dropdown menu, select **IPP Everywhere**, and click **Add Printer**.

10. After adding the new printer, you can set the default print options of your choice.

The last window confirms that you set the driverless printer and it is ready to use.

15.5.3. Configuring permanent print queues using `cups-browsed`

As an administrator, you can use the **cups-browsed** service from the **cups-filters** package to browse queues on remote print server. The **cups-browsed** service creates local queues pointing to these remote queues. A print queue is an abstraction for a printer that represents a connection to the physical device.

15.5.3.1. Configuring permanent print queues for printers installed on remote print servers in a different network

To locally install CUPS queues from a remote server, edit the configuration for the **cups-browsed** service on the machine where you want the permanent queues pointing to the remote CUPS server.

Prerequisites

- Printers in a different network must be installed on a remote server.

- The IPP port in the firewall of the server is enabled.

- The server is configured for remote access from a machine where **cups-browsed** runs, and asks the server for queues.

Procedure

1. Edit the **/etc/cups/cups-browsed.conf** file and add a hostname or an IP address of the designated server to the **BrowsePoll** directive:

   ```
   BrowsePoll <hostname or IP-address>
   ```

2. Restart the **cups-browsed** service to apply the changes:

> # systemctl restart cups-browsed

Verification steps

- Display the list of local printers that contains your local print queues:

  > $ lpstat -v
  > Device for <remote_queue_name>: implicitclass:<remote_queue_name>

> **NOTE**
>
> Your printer can take several minutes to appear on the local printers list, depending on how many queues the remote server contains.

15.6. ADDING A PRINTER WITH A CLASSIC DRIVER IN THE CUPS WEB UI

This section describes how to add a new printer using the **CUPS web user interface**

Prerequisites

- You have administration access to the **CUPS web UI** as described in Acquiring administration access to CUPS web UI.

Procedure

1. Start the **CUPS web UI** as described in Starting CUPS web UI

2. In your browser, go to **localhost:631** and select the **Administration** tab.

3. Under **Printers** click **Add printer**.
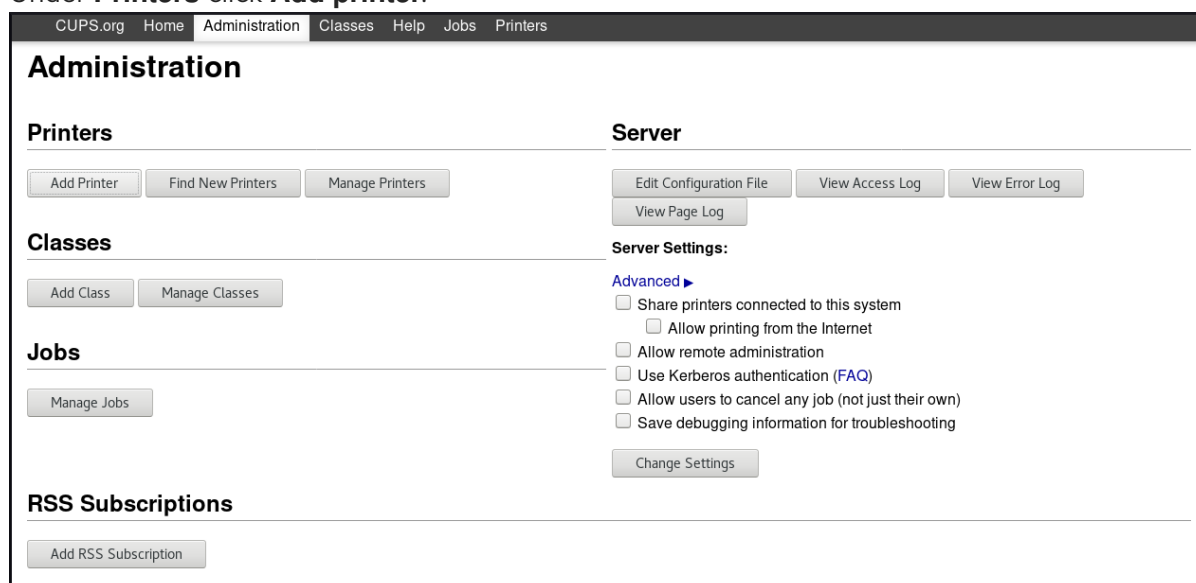
4. Authenticate by username and password:
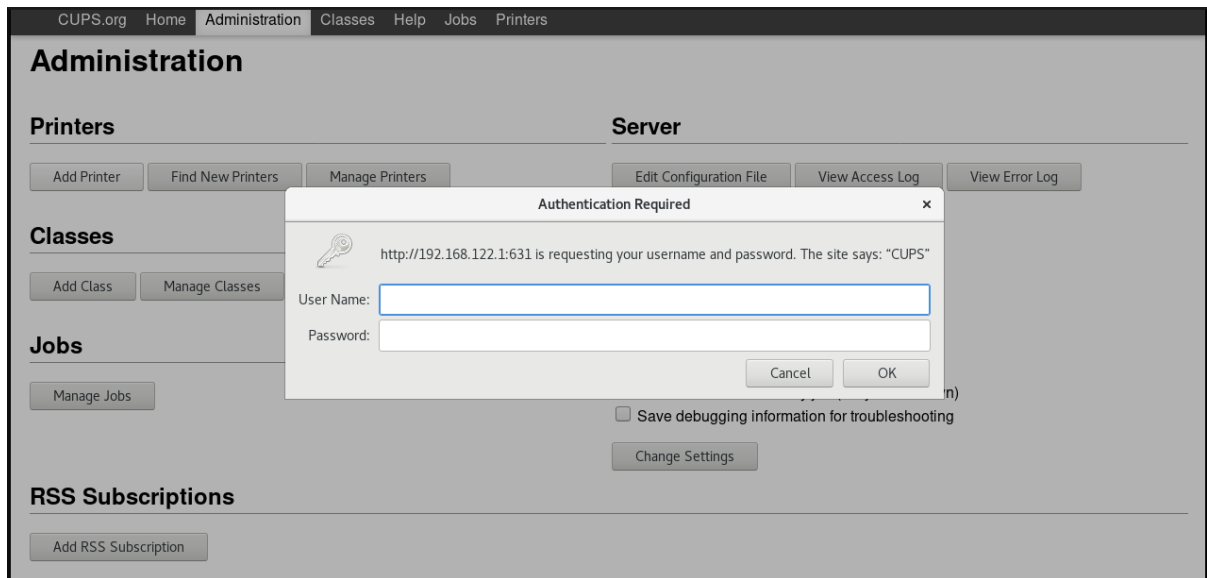
IMPORTANT

To be able to add a new printer by using the **CUPS web UI**, you must authenticate as a user who belongs to groups defined by **SystemGroup** directive in **/etc/cups/cups-files**.

Default groups:

- root

- sys

- wheel

5. If a local printer is connected, or CUPS finds a network printer available, select the printer. If neither local printer nor network printer is available, select one of the printer types from **Other Network Printers**, for example **APP Socket/HP Jet direct**, enter the IP address of the printer, and then click **Continue**.

6. If you have selected for example **APP Socket/HP Jet direct** as shown above, enter the IP address of the printer, and then click **Continue**.

7. You can add more details about the new printer, such as the name, description and location. To set a printer to be shared over the network, check the **Share This Printer** checkbox.

8. Select the printer manufacturer, and then click **Continue**.

Alternatively, you can provide a **postscript printer description**(PPD) file to be used as a driver for the printer, by clicking the **Browse...** button at the bottom.

9. Select the model of the printer, and then click **Add Printer**.

10. After the printer has been added, the next window allows you to set the default print options.

After clicking **Set Default Options**, you will receive a confirmation that the new printer has been added successfully.

Verification steps

- Print a test page especially if you have set up a printer:

- Go to **Printers** menu, and click **Maintenance → Print Test Page**.
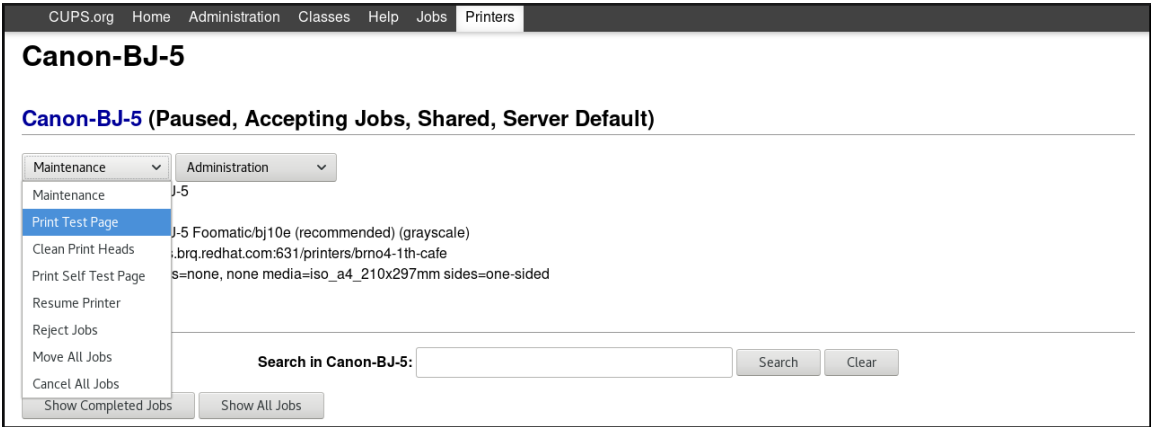
15.7. CONFIGURING A PRINTER IN THE CUPS WEB UI

This section describes how to configure a new printer, and how to maintain a configuration of a printer using the **CUPS web UI**

Prerequisites

- You have administration access to the **CUPS web UI** as described in Acquiring administration access to CUPS web UI.

Procedure

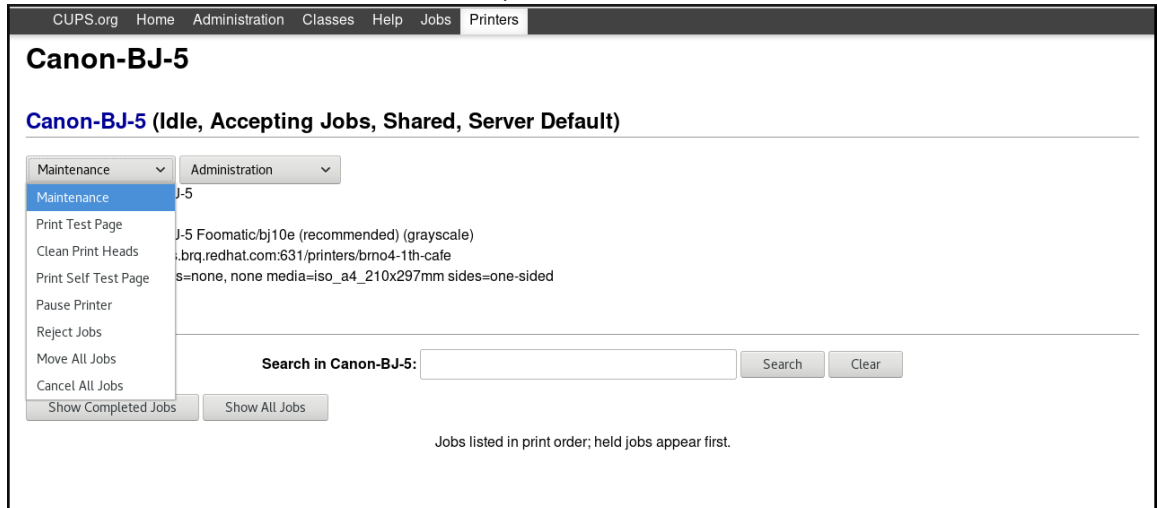1. Click the **Printers** menu to see available printers that you can configure.
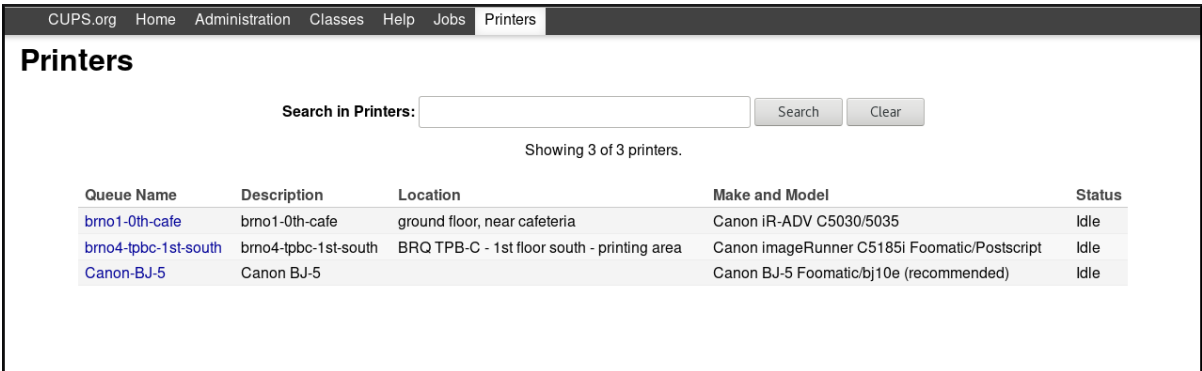
2. Choose one printer that you want to configure.
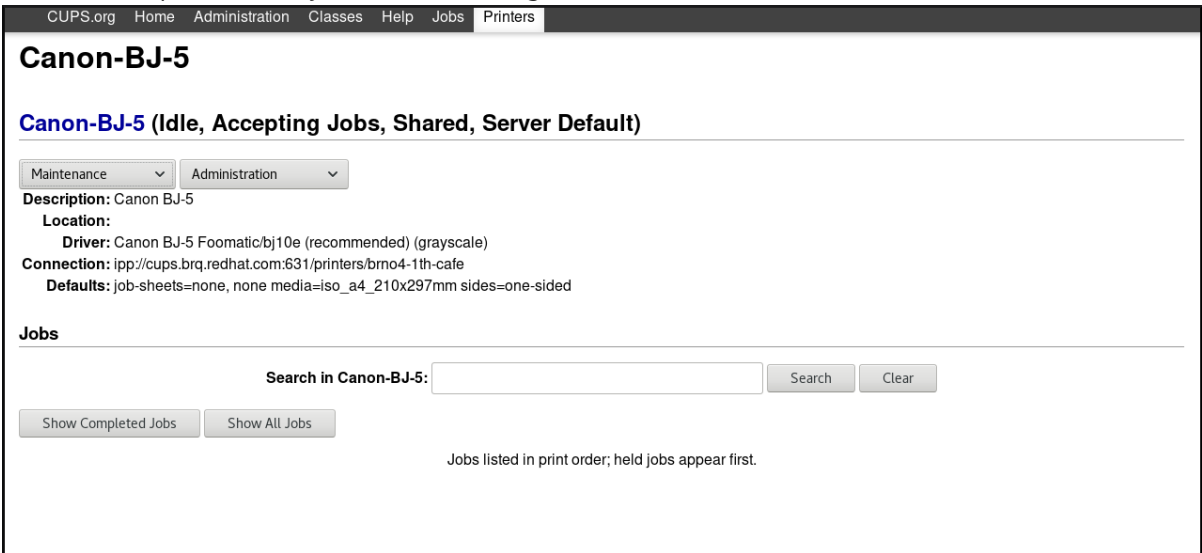
3. Perform your selected task by using one of the available menus:

- Choose **Maintenance** from the first drop-down menu.

- Choose **Administration** from the second drop-down menu.
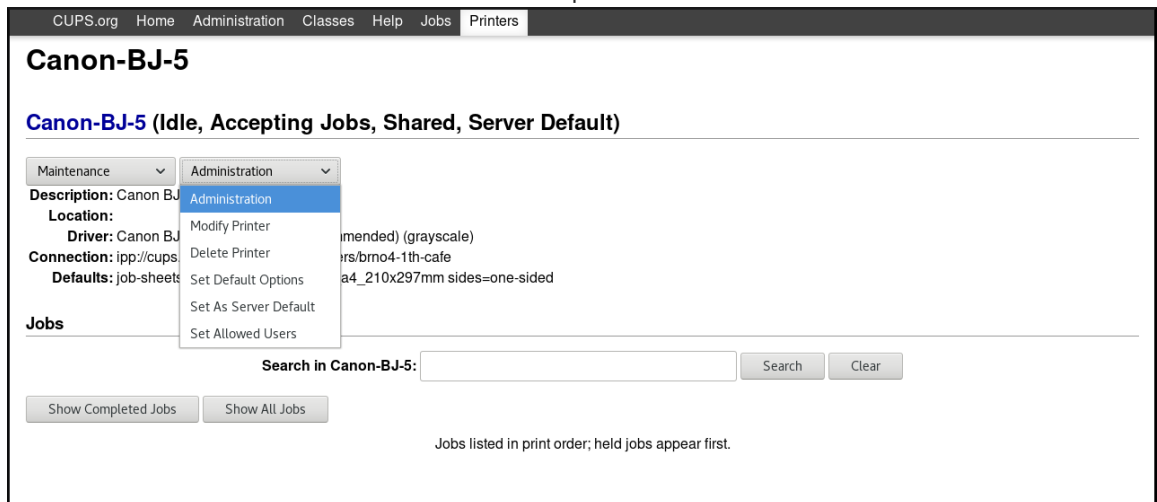
- You can also check completed print jobs or all active print jobs by clicking the **Show Completed Jobs** or **Show All Jobs** buttons.

Verification steps

- Print a test page especially if you have changed a printer configuration:
  - Go to **Printers** menu, and click **Maintenance → Print Test Page**.
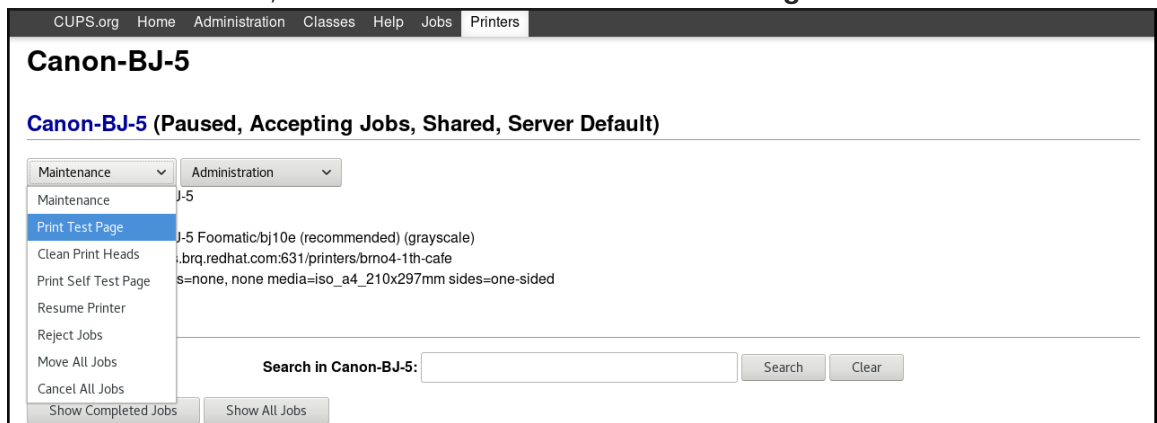
15.8. SETTING PRINT OPTIONS USING THE CUPS WEB UI

This section describes how to set common print options, such as the media size and type, print quality or the color mode, in the **CUPS web UI**.
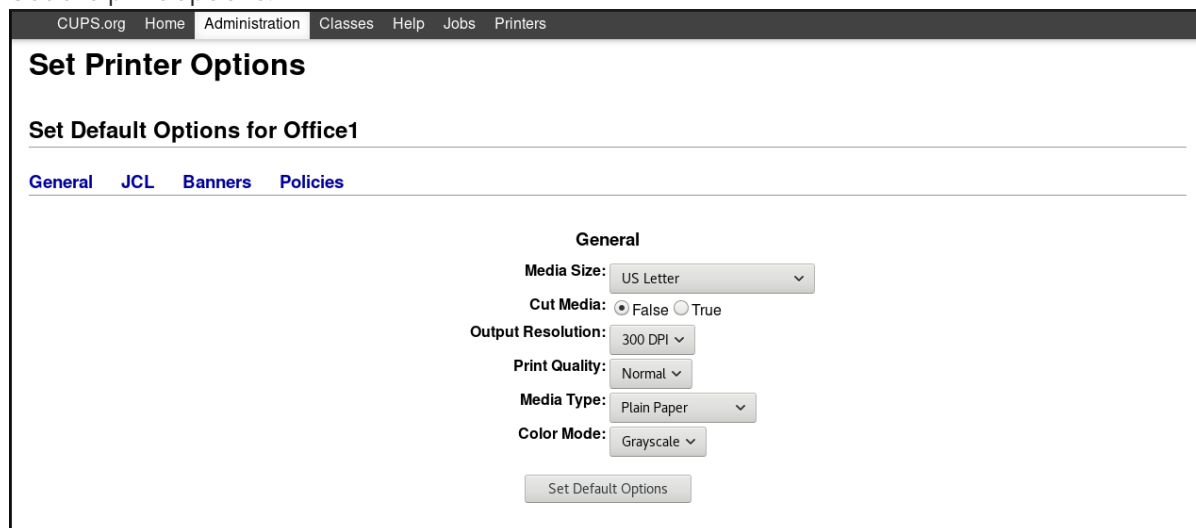
Prerequisites

You have administration access to the **CUPS web UI** as described in Acquiring administration access to CUPS web UI.

Procedure

1. Go to **Administration** menu, and click **Maintenance → Set Default Options**.
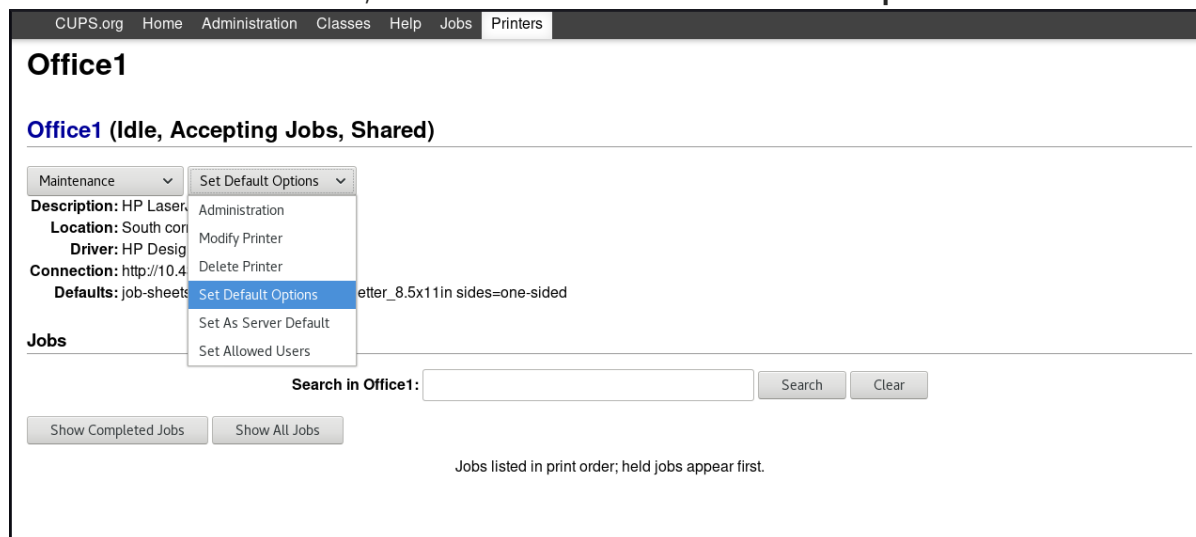
2. Set the print options.

15.9. INSTALLING CERTIFICATES FOR A PRINT SERVER

To install certificates for a print server, you can choose one of the following options:

- Automatic installation using a self-signed certificate

- Manual installation using a certificate and a private key generated by a Certification Authority
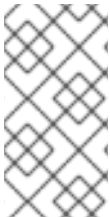
Prerequisites

For the **cupsd** daemon on the server:

1. Set the following directive in the **/etc/cups/cupsd.conf** file to:
   Encryption Required

2. Restart the cups service:

   ```
   $ sudo systemctl restart cups
   ```

Automatic installation using a self-signed certificate

With this option, CUPS generates the certificate and the key automatically.

> **NOTE**
>
> The self-signed certificate does not provide as strong security as certificates generated by Identity Management (IdM), Active Directory (AD), or Red Hat Certificate System (RHCS) Certification Authorities, but it can be used for print servers located within a secure local network.
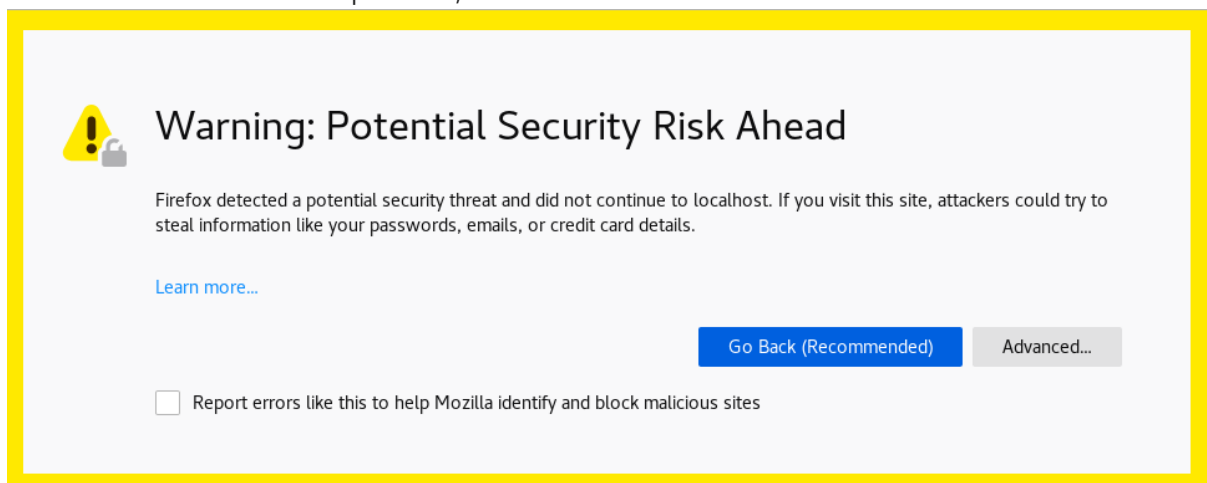
Procedure

1. To access the CUPS Web UI, open your browser and go to **https://<server>:631**
   where *<server>* is either the server IP address or the server host name.

   > **NOTE**
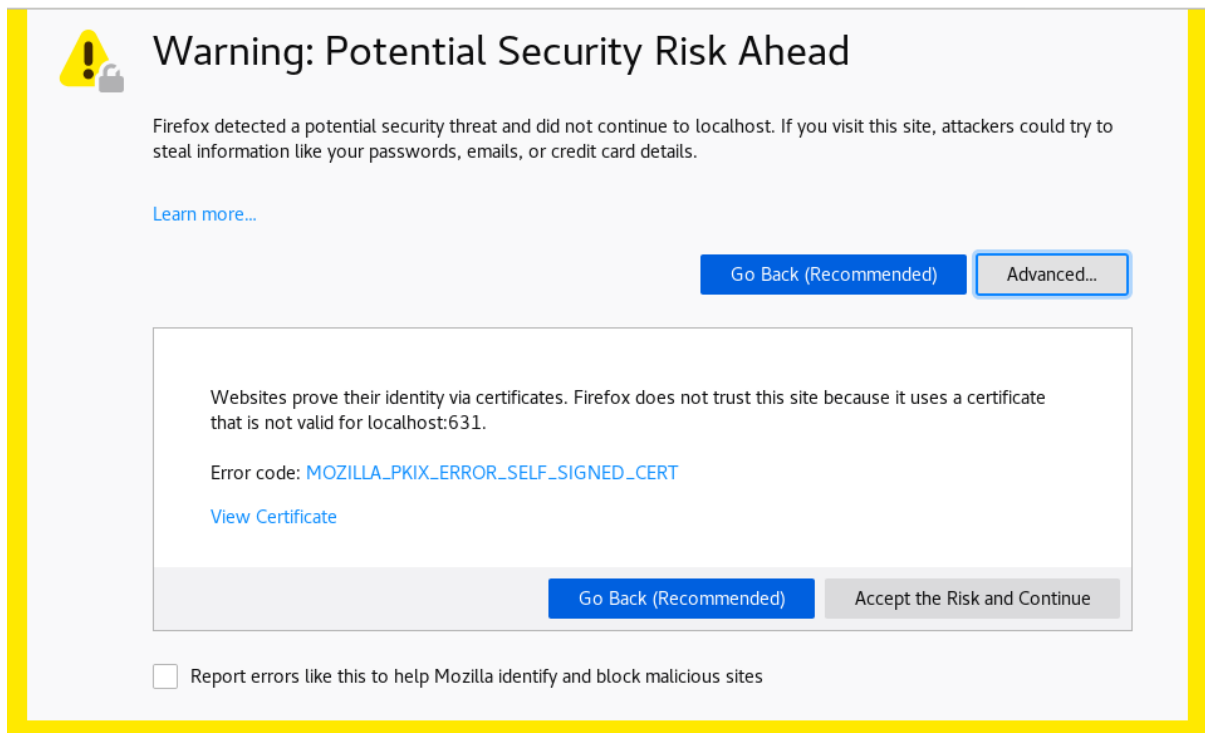   >
   > When CUPS connects to a system for the first time, the browser shows a warning about the self-signed certificate being a potential security risk.

2. To confirm that it is safe to proceed, click the **Advanced...** button.

3. Click **Accept the Risk and Continue** button.

CUPS now starts to use the self-generated certificate and the key.

> **NOTE**
>
> When you access the CUPS Web UI after an automatic installation, the browser displays a warning icon in the address bar. This is because you added the security exception by confirming the security risk warning. If you want to remove this warning icon permanently, perform the manual installation with a certificate and a private key generated by a Certification Authority.

Manual installation using a certificate and a private key generated by a Certification Authority

For print servers located within a public network or to permanently remove the warning in the browser, import the certificate and the key manually.

Prerequisites

- You have certificate and private key files generated by IdM, AD, or by RHCS Certification Authorities.

Procedure

1. Copy the **.crt** and **.key** files into the **/etc/cups/ssl** directory of the system on which you want to use the CUPS Web UI.

2. Rename the copied files to **<hostname>.crt** and **<hostname>.key**.
   Replace *<hostname>* with the host name of the system to which you want to connect the CUPS Web UI.

3. Set the following permissions to the renamed files:

   - **# chmod 644 /etc/cups/ssl/<hostname>.crt**

- **# chmod 644 /etc/cups/ssl/<hostname>.key**

- **# chown root:root /etc/cups/ssl/<hostname>.crt**

- **# chown root:root /etc/cups/ssl/<hostname>.key**

4. Restart the cups service:

   - **# systemctl restart cupsd**

15.10. USING SAMBA TO PRINT TO A WINDOWS PRINT SERVER WITH KERBEROS AUTHENTICATION

With the **samba-krb5-printing** wrapper, Active Directory (AD) users who are logged in to Red Hat Enterprise Linux can authenticate to Active Directory (AD) by using Kerberos and then print to a local CUPS print server that forwards the print job to a Windows print server.

The benefit of this configuration is that the administrator of CUPS on Red Hat Enterprise Linux does not need to store a fixed user name and password in the configuration. CUPS authenticates to AD with the Kerberos ticket of the user that sends the print job.

This section describes how to configure CUPS for this scenario.

> **NOTE**
>
> Red Hat only supports submitting print jobs to CUPS from your local system, and not to re-share a printer on a Samba print server.

Prerequisites

- The printer that you want to add to the local CUPS instance is shared on an AD print server.

- You joined the Red Hat Enterprise Linux host as a member to the AD.

- CUPS is installed on Red Hat Enterprise Linux and the **cups** service is running. For details, see Activating CUPS service .

- The PostScript Printer Description (PPD) file for the printer is stored in the **/usr/share/cups/model/** directory.

Procedure

1. Install the **samba-krb5-printing**, **samba-client**, and **krb5-workstation** packages:

   ```
   # yum install samba-krb5-printing samba-client krb5-workstation
   ```

2. Optional: Authenticate as a domain administrator and display the list of printers that are shared on the Windows print server:

   ```
   # smbclient -L win_print_srv.ad.example.com -U
   administrator@AD_KERBEROS_REALM --use-kerberos=required

   Sharename     Type     Comment
   ---------     ----     -------
   ```

> ...
> Example   Printer   Example
> ...

3. Optional: Display the list of CUPS models to identify the PPD name of your printer:

   > **lpinfo -m**
   > ...
   > *samsung.ppd Samsung M267x 287x Series PXL*
   > ...

   You require the name of the PPD file when you add the printer in the next step.

4. Add the printer to CUPS:

   > # **lpadmin -p "*example_printer*" -v smb://*win_print_srv.ad.example.com*/*Example* -m *samsung.ppd* -o auth-info-required=negotiate -E**

   The command uses the following options:

   - **-p *printer_name*** sets the name of the printer in CUPS.

   - **-v *URI_to_Windows_printer*** sets the URI to the Windows printer. Use the following format: **smb://*host_name*/*printer_share_name***.

   - **-m *PPD_file*** sets the PPD file the printer uses.

   - **-o auth-info-required=negotiate** configures CUPS to use Kerberos authentication when it forwards print jobs to the remote server.

   - **-E** enables the printer and CUPS accepts jobs for the printer.

Verification steps

1. Log into the Red Hat Enterprise Linux host as an AD domain user.

2. Authenticate as an AD domain user:

   > # **kinit *domain_user_name@AD_KERBEROS_REALM***

3. Print a file to the printer you added to the local CUPS print server:

   > # **lp -d *example_printer file***

15.11. WORKING WITH CUPS LOGS

15.11.1. Types of CUPS logs

CUPS provides three different kinds of logs:

- Error log – Stores error messages, warnings, and debugging messages.

- Access log – Stores messages about how many times CUPS clients and web UI have been accessed.

- Page log – Stores messages about the total number of pages printed for each print job.

In Red Hat Enterprise Linux 8, all three types are logged centrally in **systemd-journald** together with logs from other programs.

> **WARNING**
>
> In Red Hat Enterprise Linux 8, the logs are no longer stored in specific files within the **/var/log/cups** directory, which was used in Red Hat Enterprise Linux 7.

15.11.2. Accessing all CUPS logs

You can list all CUPS logs available in **systemd-journald**.

Procedure

- Filter CUPS logs:

  ```
  $ journalctl -u cups
  ```

15.11.3. Accessing CUPS logs for a specific print job

If you need to find a CUPS log for a specific print job, you can do it by filtering the logs using the ID of a print job.

Procedure

- Filter logs for a specific print job:

  ```
  $ journalctl -u cups JID=N
  ```

  Where **N** is the ID of a print job.

15.11.4. Accessing CUPS logs by specific time frame

If you need to access CUPS logs during a certain time period, you can filter the logs in **systemd-journald**.

Procedure

- Filter logs within the specified time frame:

  ```
  $ journalctl -u cups --since=YYYY-MM-DD --until=YYYY-MM-DD
  ```

  Where **YYYY** is year, **MM** is month and **DD** is day.

Additional resources

- The **journalctl(1)** man page

15.11.5. Configuring the CUPS log location

In Red Hat Enterprise Linux 8, CUPS logs are by default logged into **systemd-journald** service. You can find the default configuration in the **/etc/cups/cups-files.conf** file:

> ErrorLog syslog

> **IMPORTANT**
>
> Red Hat recommends keeping the default location of CUPS logs.

Procedure

- To send the logs into a different location, change the **ErrorLog** entry in the **/etc/cups/cups-files.conf** to match your required location:

  > ErrorLog <your_required_location>

> **WARNING**
>
> If you change the default location of CUPS log, you might experience SELinux issues.